

File Name: T-ICML-O\_1\_I1\_introduction

Format: Presenter in Studio

Presenter: Evan Jones



---

## Linear and DNN Models for Image Classification

Evan Jones

# Agenda

---

## **Introduction**

Linear Models

Deep Neural Network Models

DNN Dropout

# Learn how to...

---

Understand how image data is represented as floating point numbers that can be flattened

Compare functions for model confidence in image classification (Softmax)

Train and evaluate a Linear model for image classification using TensorFlow


Train and evaluate a Deep Neural Network (DNN) model for image classification using TensorFlow

Understand how to apply dropout as a regularization technique for DNNs

# Problem: Recognizing Handwritten Digits




# Introducing the MNIST dataset of labeled images



60,000 total images


Training 50K										Test 10K				
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
2	2	2	2	2	2	2	2	2	2	2	2	2	2	2
3	3	3	3	3	3	3	3	3	3	3	3	3	3	3
4	4	4	4	4	4	4	4	4	4	4	4	4	4	4
5	5	5	5	5	5	5	5	5	5	5	5	5	5	5
6	6	6	6	6	6	6	6	6	6	6	6	6	6	6
7	7	7	7	7	7	7	7	7	7	7	7	7	7	7
8	8	8	8	8	8	8	8	8	8	8	8	8	8	8
9	9	9	9	9	9	9	9	9	9	9	9	9	9	9

Each greyscale image is  
28 x 28 pixels

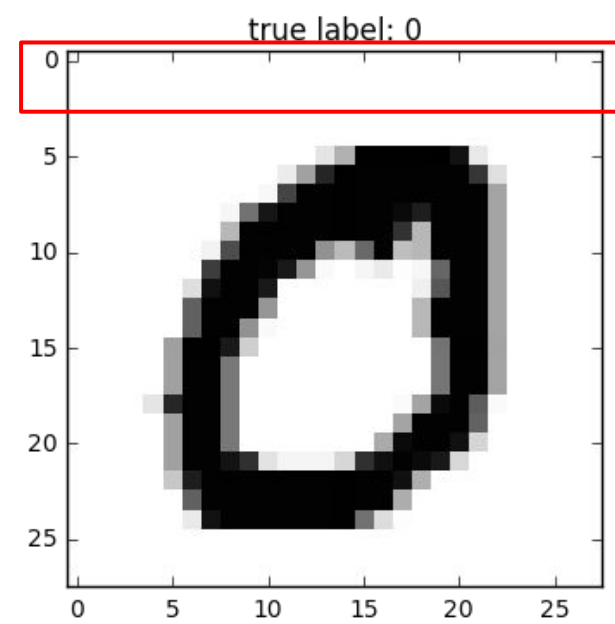






Each greyscale image is  
28 x 28 pixels




Unstack the pixels  
into one long array




The flattened image is represented as an array



How many output classes  
do we have?



How many output classes  
do we have?




0
1
2
3
4
5
6
7
8
9


# What if the model is unsure?

A 10x10 grid of handwritten '2's. Each row contains 10 '2's, and there are 10 rows in total. The '2's are written in a cursive, handwritten style, with variations in slant, loop, and stroke thickness, demonstrating different ways to form the digit '2' by hand.

What if the model is  
unsure?




# Assessing the model's confidence with a function



0	20%
1	70%
2	
3	
4	
5	
6	
7	
8	10%
9	




# Assessing the model's confidence with a function



??


# Assessing the model's confidence with a function

Sigmoid?



# Assessing the model's confidence with a function

Sigmoid? Nope.



# Softmax exponentiating its inputs and then normalizing them

$$\text{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

## Input



$z_1^L = 4.9$



$z_2^L = 1.2$



$z_3^L = 1.2$



$z_4^L = 3.2$

## Softmax



$a_1^L = 0.812$



$a_2^L = 0.020$





$a_3^L = 0.020$



$a_4^L = 0.148$

# Softmax exponentiating its inputs and then normalizing them



$$\text{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$



0	2%
1	
2	10%
3	
4	11%
5	3%
6	<u>70%</u>
7	
8	3%
9	1%

# Softmax exponentiating its inputs and then normalizing them

$$\text{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$




0	2%
1	
2	10%
3	
4	11%
5	3%
6	<u>70%</u>
7	
8	3%
9	1%

X

# Training and evaluating our image classification model

## Minimize cross-entropy error



$$-\frac{1}{N} \times \sum_1^N y_i \times \log(\hat{y}_i) + (1 - y_i) \times \log(1 - \hat{y}_i)$$

Numerically more stable:

`softmax_cross_entropy_with_logits`

## Accuracy for model performance

$$\frac{\text{\# correctly classified images}}{\text{\# number total images}}$$

8 → 8  
0 → 0  
1 → 1  
9 → 9  
4 → 4  
6 → 1 ✖  
7 → 7

$$\frac{7}{8} = 88\%$$

File Name: T-ICML-O\_1\_I2\_linear\_models

Format: Presenter in Studio

Presenter: Evan Jones



# Agenda

---

Introduction

**Linear Models**

Deep Neural Network Models

DNN Dropout

# Training and evaluating our image classification model

## Learning

```
def my_model_fn(features, labels, mode):
    predictions, num_classes = my_model(features)
    loss = ...
    train_op = ...
    return tf.estimator.EstimatorSpec(
        mode=mode,
        predictions=predictions,
        loss=loss,
        train_op=train_op)

estimator = tf.estimator.Estimator(model_fn = my_model_fn, params)

train_input_fn = tf.estimator.inputs.numpy_input_fn(params)
eval_input_fn = tf.estimator.inputs.numpy_input_fn(params)
train_spec = tf.estimator.TrainSpec(input_fn = train_input_fn, params)
eval_spec = tf.estimator.EvalSpec(input_fn = eval_input_fn, params)
tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)`
```

# Training and evaluating our image classification model

```
def my_model_fn(features, labels, mode):  
    predictions, num_classes = my_model(features)  
    loss = ...  
    train_op = ...  
    return tf.estimator.EstimatorSpec(  
        mode=mode,  
        predictions=predictions,  
        loss=loss,  
        train_op=train_op)  
  
estimator = tf.estimator.Estimator(model_fn = my_model_fn, params)  
  
train_input_fn = tf.estimator.inputs.numpy_input_fn(params)  
eval_input_fn = tf.estimator.inputs.numpy_input_fn(params)  
train_spec = tf.estimator.TrainSpec(input_fn = train_input_fn, params)  
eval_spec = tf.estimator.EvalSpec(input_fn = eval_input_fn, params)  
tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)`
```


# Training and evaluating our image classification model

```
def my_model_fn(features, labels, mode):  
    predictions, num_classes = my_model(features)  
    loss = ...  
    train_op = ...  
    return tf.estimator.EstimatorSpec(  
        mode=mode,  
        predictions=predictions,  
        loss=loss,  
        train_op=train_op)  
  
estimator = tf.estimator.Estimator(model_fn = my_model_fn, params)  
  
train_input_fn = tf.estimator.inputs.numpy_input_fn(params)  
eval_input_fn = tf.estimator.inputs.numpy_input_fn(params)  
train_spec = tf.estimator.TrainSpec(input_fn = train_input_fn, params)  
eval_spec = tf.estimator.EvalSpec(input_fn = eval_input_fn, params)  
tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```

# Training and evaluating our image classification model

```
def my_model_fn(features, labels, mode):  
    predictions, num_classes = my_model(features)  
    loss = ...  
    train_op = ...  
    return tf.estimator.EstimatorSpec(  
        mode=mode,  
        predictions=predictions,  
        loss=loss,  
        train_op=train_op)  
  
estimator = tf.estimator.Estimator(model_fn = my_model_fn, params)  
  
train_input_fn = tf.estimator.inputs.numpy_input_fn(params)  
eval_input_fn = tf.estimator.inputs.numpy_input_fn(params)  
train_spec = tf.estimator.TrainSpec(input_fn = train_input_fn, params)  
eval_spec = tf.estimator.EvalSpec(input_fn = eval_input_fn, params)  
tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```

# Making predictions with our image classification model



```
HEIGHT=28
WIDTH=28
NCLASSES=10

def linear_model(img):
    """Uses a linear model to compute a vector representing relative confidence
    that img belongs to one of NCLASSES classes.

    Args:
        img [batchsize, HEIGHT, WIDTH]: A tensor of floats representing a batch
        of images.

    Returns:
        logits: the output of the model
        NCLASSES: the number of classes
    """
    X = tf.reshape(img, [-1, HEIGHT*WIDTH]) # [-1, HEIGHT * WIDTH]
    W = tf.Variable(tf.zeros([HEIGHT*WIDTH, NCLASSES])) # [HEIGHT * WIDTH, CLASSES]
    b = tf.Variable(tf.zeros([NCLASSES])) # [NCLASSES]
    ylogits = tf.matmul(X, W) + b # [-1, NCLASSES]
    return ylogits, NCLASSES
```

# Making predictions with our image classification model

```
HEIGHT=28
WIDTH=28
NCLASSES=10

def linear_model(img):
    """Uses a linear model to compute a vector representing relative confidence
    that img belongs to one of NCLASSES classes.

    Args:
        img [batchsize, HEIGHT, WIDTH]: A tensor of floats representing a batch
        of images.

    Returns:
        logits: the output of the model
        NCLASSES: the number of classes
    """
    X = tf.reshape(img, [-1, HEIGHT*WIDTH]) # [-1, HEIGHT * WIDTH]
    W = tf.Variable(tf.zeros([HEIGHT*WIDTH, NCLASSES])) # [HEIGHT * WIDTH, CLASSES]
    b = tf.Variable(tf.zeros([NCLASSES])) # [NCLASSES]
    ylogits = tf.matmul(X, W) + b # [-1, NCLASSES]
    return ylogits, NCLASSES
```

# Making predictions with our image classification model

```
HEIGHT=28
WIDTH=28
NCLASSES=10

def linear_model(img):
    """Uses a linear model to compute a vector representing relative confidence
    that img belongs to one of NCLASSES classes.

    Args:
        img [batchsize, HEIGHT, WIDTH]: A tensor of floats representing a batch
        of images.

    Returns:
        logits: the output of the model
        NCLASSES: the number of classes
    """
    X = tf.reshape(img, [-1, HEIGHT*WIDTH]) # [-1, HEIGHT * WIDTH]
    W = tf.Variable(tf.zeros([HEIGHT*WIDTH, NCLASSES])) # [HEIGHT * WIDTH, CLASSES]
    b = tf.Variable(tf.zeros([NCLASSES])) # [NCLASSES]
    ylogits = tf.matmul(X, W) + b # [-1, NCLASSES]
    return ylogits, NCLASSES
```



# Making predictions with our image classification model

```
HEIGHT=28
WIDTH=28
NCLASSES=10

def linear_model(img):
    """Uses a linear model to compute a vector representing relative confidence
    that img belongs to one of NCLASSES classes.

    Args:
        img [batchsize, HEIGHT, WIDTH]: A tensor of floats representing a batch
        of images.

    Returns:
        logits: the output of the model
        NCLASSES: the number of classes
    """
    X = tf.reshape(img, [-1, HEIGHT*WIDTH]) # [-1, HEIGHT * WIDTH]
    W = tf.Variable(tf.zeros([HEIGHT*WIDTH, NCLASSES])) # [HEIGHT * WIDTH, CLASSES]
    b = tf.Variable(tf.zeros([NCLASSES])) # [NCLASSES]
    ylogits = tf.matmul(X, W) + b # [-1, NCLASSES]
    return ylogits, NCLASSES
```

# Making predictions with our image classification model

```
HEIGHT=28
WIDTH=28
NCLASSES=10

def linear_model(img):
    """Uses a linear model to compute a vector representing relative confidence
    that img belongs to one of NCLASSES classes.

    Args:
        img [batchsize, HEIGHT, WIDTH]: A tensor of floats representing a batch
        of images.

    Returns:
        logits: the output of the model
        NCLASSES: the number of classes
    """
    X = tf.reshape(img, [-1, HEIGHT*WIDTH]) # [-1, HEIGHT * WIDTH]
    W = tf.Variable(tf.zeros([HEIGHT*WIDTH, NCLASSES])) # [HEIGHT * WIDTH, CLASSES]
    b = tf.Variable(tf.zeros([NCLASSES])) # [NCLASSES]
    ylogits = tf.matmul(X, W) + b # [-1, NCLASSES]
    return ylogits, NCLASSES
```

# Making predictions with our image classification model

```
HEIGHT=28
WIDTH=28
NCLASSES=10

def linear_model(img):
    """Uses a linear model to compute a vector representing relative confidence
    that img belongs to one of NCLASSES classes.

    Args:
        img [batchsize, HEIGHT, WIDTH]: A tensor of floats representing a batch
        of images.

    Returns:
        logits: the output of the model
        NCLASSES: the number of classes
    """
    X = tf.reshape(img, [-1, HEIGHT*WIDTH]) # [-1, HEIGHT * WIDTH]
    W = tf.Variable(tf.zeros([HEIGHT*WIDTH, NCLASSES])) # [HEIGHT * WIDTH, CLASSES]
    b = tf.Variable(tf.zeros([NCLASSES])) # [NCLASSES]
    ylogits = tf.matmul(X, W) + b # [-1, NCLASSES]
    return ylogits, NCLASSES
```



# Training and evaluating our image classification model

```
def my_model_fn(features, labels, mode):  
    logits, num_classes = linear_model(features['image'])  
    probabilities = tf.nn.softmax(logits)  
    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits_v2(  
        logits=logits, labels=labels))  
    train_op = ...  
    return tf.estimator.EstimatorSpec(  
        mode=mode,  
        predictions=probabilities,  
        loss=loss,  
        train_op=train_op)
```

# Training and evaluating our image classification model

```
def my_model_fn(features, labels, mode):  
    logits, num_classes = linear_model(features['image'])  
    probabilities = tf.nn.softmax(logits)  
    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits_v2(  
        logits=logits, labels=labels))  
    train_op = ...  
    return tf.estimator.EstimatorSpec(  
        mode=mode,  
        predictions=probabilities,  
        loss=loss,  
        train_op=train_op)
```

# Training and evaluating our image classification model

```
def my_model_fn(features, labels, mode):  
    logits, num_classes = linear_model(features['image'])  
    probabilities = tf.nn.softmax(logits)  
    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits_v2(  
        logits=logits, labels=labels))  
    train_op = ...  
    return tf.estimator.EstimatorSpec(  
        mode=mode,  
        predictions=probabilities,  
        loss=loss,  
        train_op=train_op)
```

# Training and evaluating our image classification model

```
def my_model_fn(features, labels, mode):  
    logits, num_classes = linear_model(features['image'])  
    probabilities = tf.nn.softmax(logits)  
    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits_v2(  
        logits=logits, labels=labels))  
    train_op = ...  
    return tf.estimator.EstimatorSpec(  
        mode=mode,  
        predictions=probabilities,  
        loss=loss,  
        train_op=train_op)
```

# Training and evaluating our image classification model

```
def my_model_fn(features, labels, mode):  
    logits, num_classes = linear_model(features['image'])  
    probabilities = tf.nn.softmax(logits)  
    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits_v2(  
        logits=logits, labels=labels))  
    train_op = ...  
    return tf.estimator.EstimatorSpec(  
        mode=mode,  
        predictions=probabilities,  
        loss=loss,  
        train_op=train_op)
```



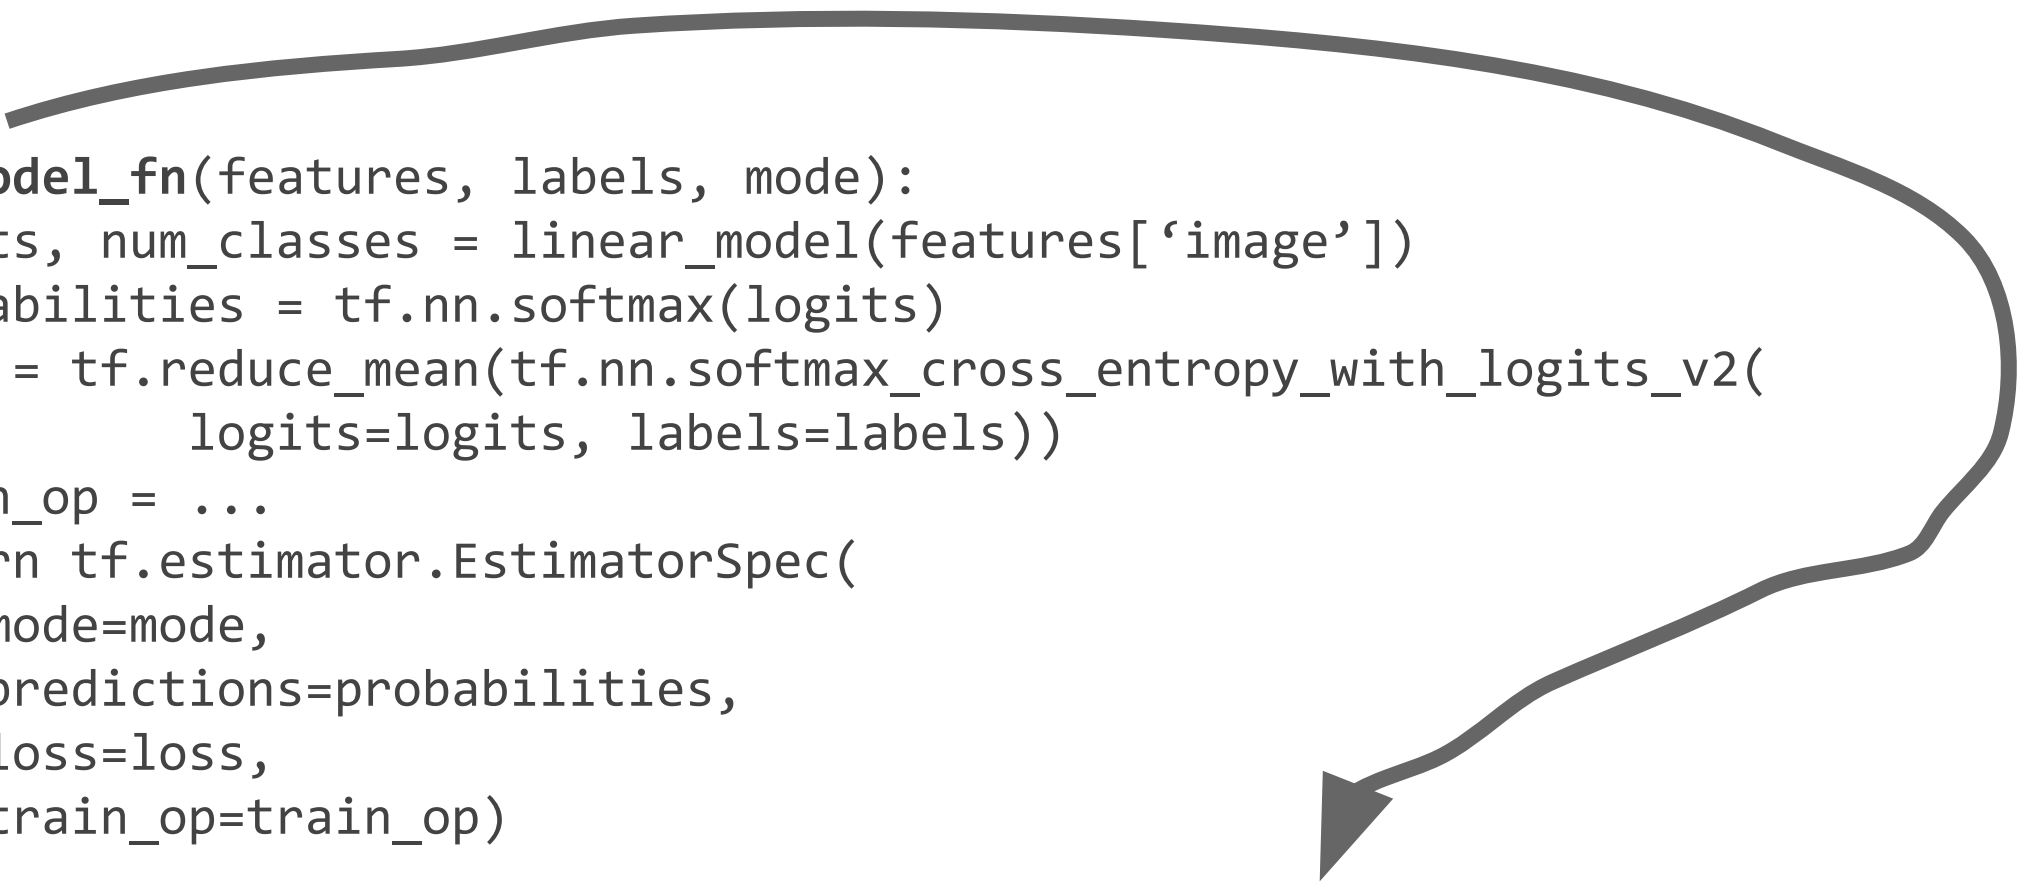
# Training and evaluating our image classification model

```
def my_model_fn(features, labels, mode):  
    logits, num_classes = linear_model(features['image'])  
    probabilities = tf.nn.softmax(logits)  
    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits_v2(  
        logits=logits, labels=labels))  
    train_op = ...  
    return tf.estimator.EstimatorSpec(  
        mode=mode,  
        predictions=probabilities,  
        loss=loss,  
        train_op=train_op)
```

# Training and evaluating our image classification model

```
def my_model_fn(features, labels, mode):  
    logits, num_classes = linear_model(features['image'])  
    probabilities = tf.nn.softmax(logits)  
    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits_v2(  
        logits=logits, labels=labels))  
    train_op = ...  
    return tf.estimator.EstimatorSpec(  
        mode=mode,  
        predictions=probabilities,  
        loss=loss,  
        train_op=train_op)  
  
estimator = tf.estimator.Estimator(model_fn = my_model_fn, params)
```

# Training and evaluating our image classification model



```
def my_model_fn(features, labels, mode):  
    logits, num_classes = linear_model(features['image'])  
    probabilities = tf.nn.softmax(logits)  
    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits_v2(  
        logits=logits, labels=labels))  
    train_op = ...  
    return tf.estimator.EstimatorSpec(  
        mode=mode,  
        predictions=probabilities,  
        loss=loss,  
        train_op=train_op)  
  
estimator = tf.estimator.Estimator(model_fn = my_model_fn, params)
```

# Training and evaluating our image classification model

```
def my_model_fn(features, labels, mode):  
    logits, num_classes = linear_model(features['image'])  
    probabilities = tf.nn.softmax(logits)  
    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits_v2(  
        logits=logits, labels=labels))  
    train_op = ...  
    return tf.estimator.EstimatorSpec(  
        mode=mode,  
        predictions=probabilities,  
        loss=loss,  
        train_op=train_op)  
  
estimator = tf.estimator.Estimator(model_fn = my_model_fn, params)
```



# Training and evaluating our image classification model

```
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets('mnist/data', one_hot=True, reshape=False)

train_input_fn = tf.estimator.inputs.numpy_input_fn(
    x={'image':mnist.train.images},
    y=mnist.train.labels,
    batch_size=100,
    num_epochs=None,
    shuffle=True,
    queue_capacity=5000
)
```

# Training and evaluating our image classification model

```
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets('mnist/data', one_hot=True, reshape=False)

train_input_fn = tf.estimator.inputs.numpy_input_fn(
    x={'image':mnist.train.images},
    y=mnist.train.labels,
    batch_size=100,
    num_epochs=None,
    shuffle=True,
    queue_capacity=5000
)
```

# Training and evaluating our image classification model

```
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets('mnist/data', one_hot=True, reshape=False)

train_input_fn = tf.estimator.inputs.numpy_input_fn(
    x={'image':mnist.train.images},
    y=mnist.train.labels,
    batch_size=100,
    num_epochs=None,
    shuffle=True,
    queue_capacity=5000
)
```



# Training and evaluating our image classification model

[illegible]

# Training and evaluating our image classification model

[illegible]

# Training and evaluating our image classification model

[illegible]

# Training and evaluating our image classification model

[illegible]

# Training and evaluating our image classification model

[illegible]

# Training and evaluating our image classification model

[illegible]

# Training and evaluating our image classification model

[illegible]

# Training and evaluating our image classification model

```
train_spec = tf.estimator.TrainSpec(input_fn = train_input_fn,
                                     max_steps = hparams['train_steps'])

exporter = tf.estimator.LatestExporter('Servo', model.serving_input_fn)

eval_spec = tf.estimator.EvalSpec(input_fn = eval_input_fn,
                                   steps = None,
                                   exporters = exporter,
                                   throttle_secs = EVAL_INTERVAL)

tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```



Video Name: T-ICML-O\_1\_I3\_lab\_intro:\_linear\_models

Format: Studio with Presenter

Presenter: Evan Jones

# Lab

---

## Image Classification with a Linear Model

Evan Jones

Video Name: T-ICML-O\_1\_I4\_lab\_solution:\_linear\_models

Format: Studio with Presenter

Presenter: Evan Jones



File Name: T-ICML-O\_1\_I5\_dnn\_models

Format: Presenter in Studio

Presenter: Evan Jones

# Agenda

---

Introduction


Linear Models

**Deep Neural Network  
Models**


DNN Dropout

## OUTPUT

Training loss 0.002





Colors shows  
data, neuron and  
weight values.



## OUTPUT

Training loss 0.572








## OUTPUT

Training loss 0.504

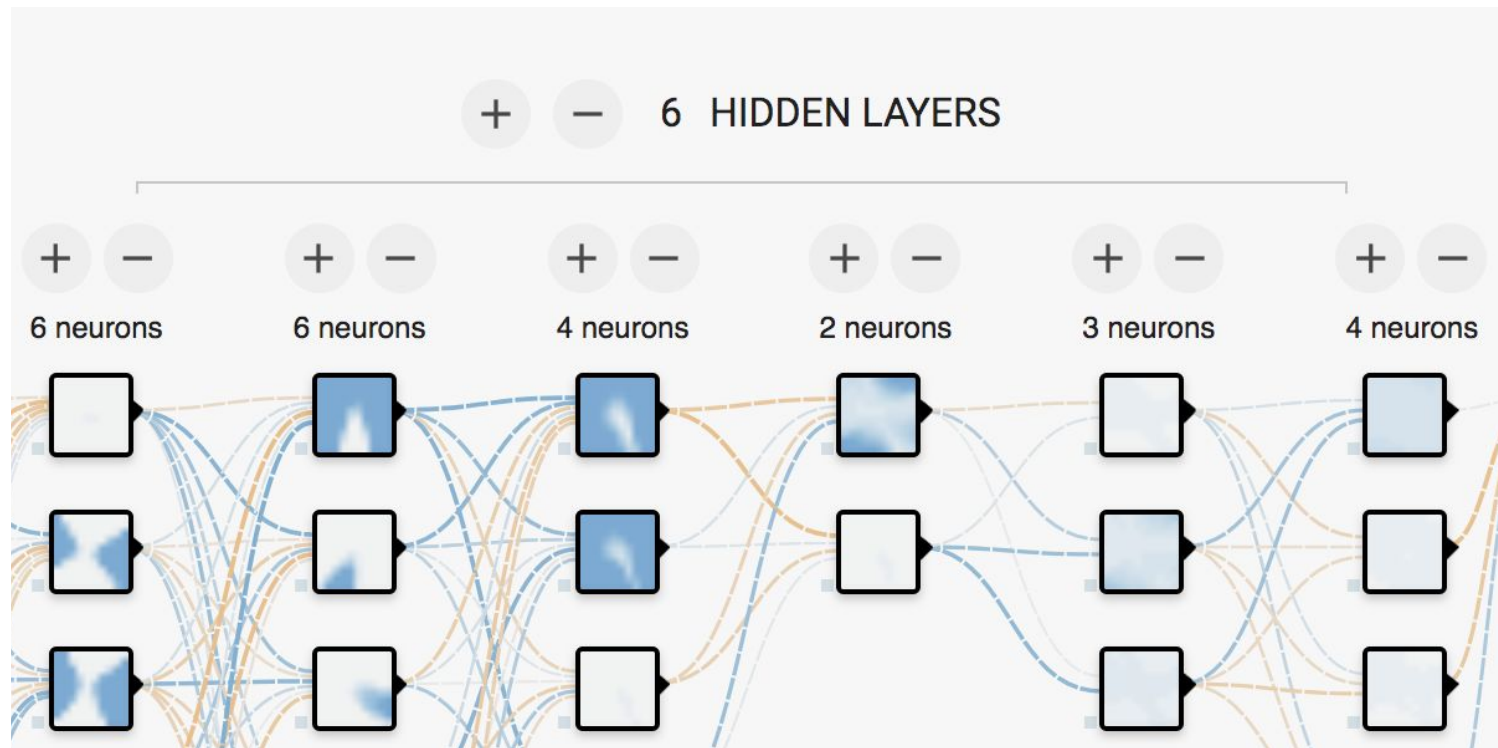


Activation

ReLU

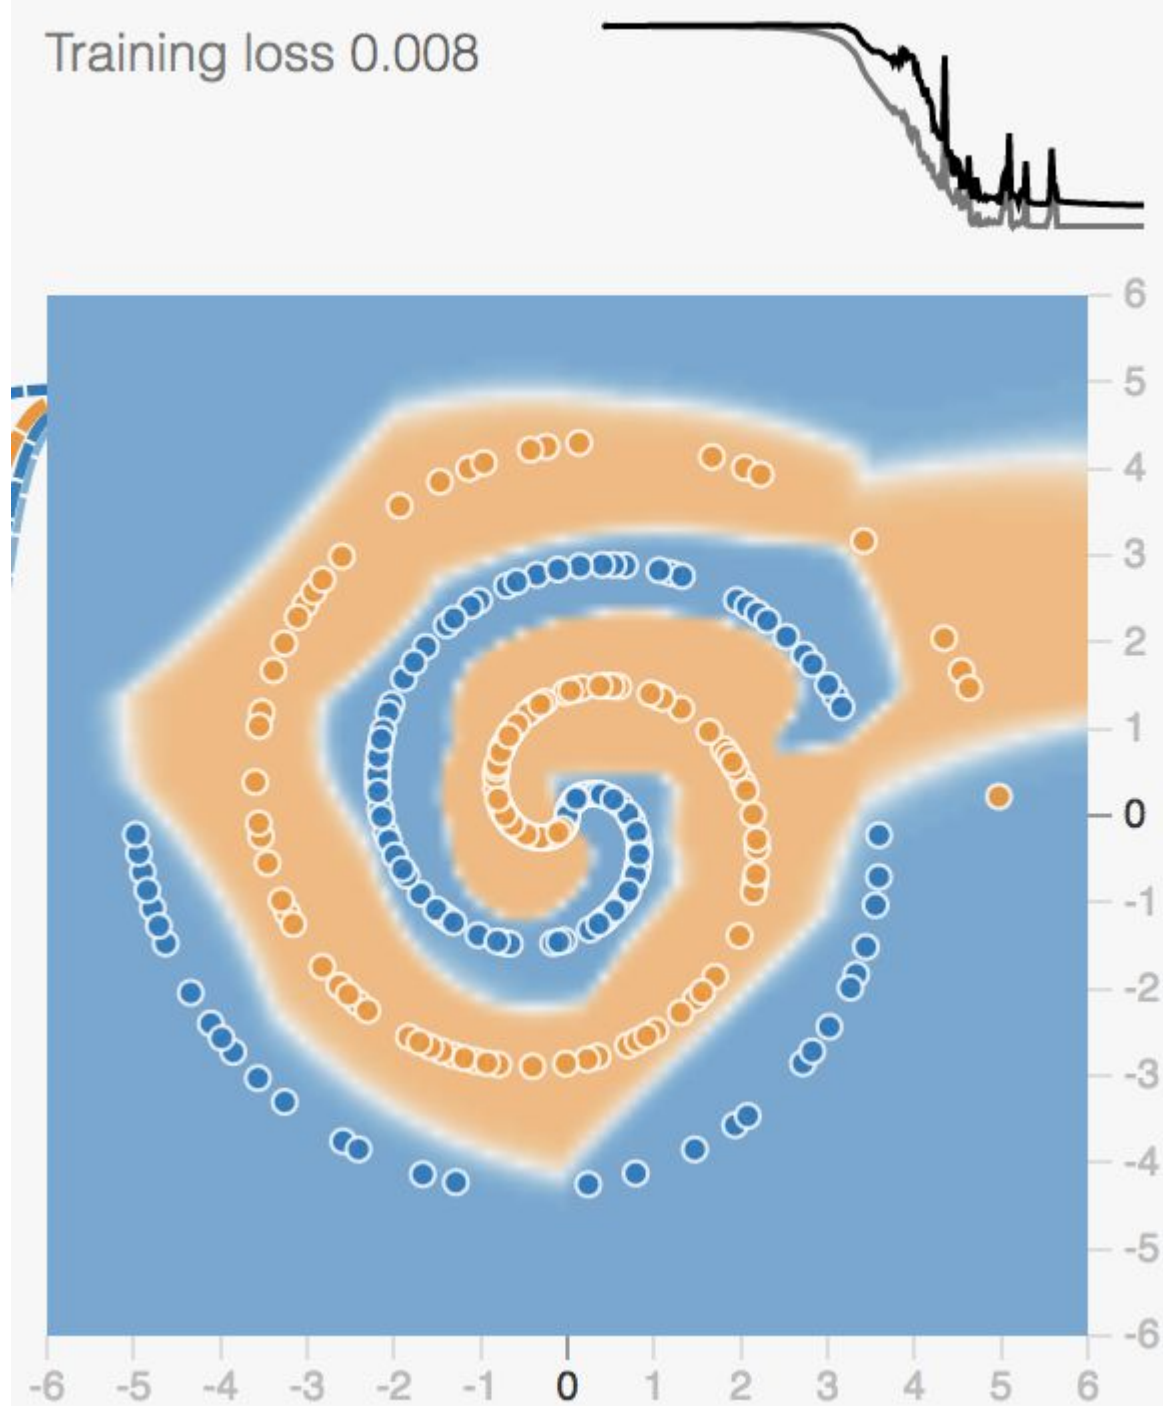
---





## OUTPUT

Training loss 0.008



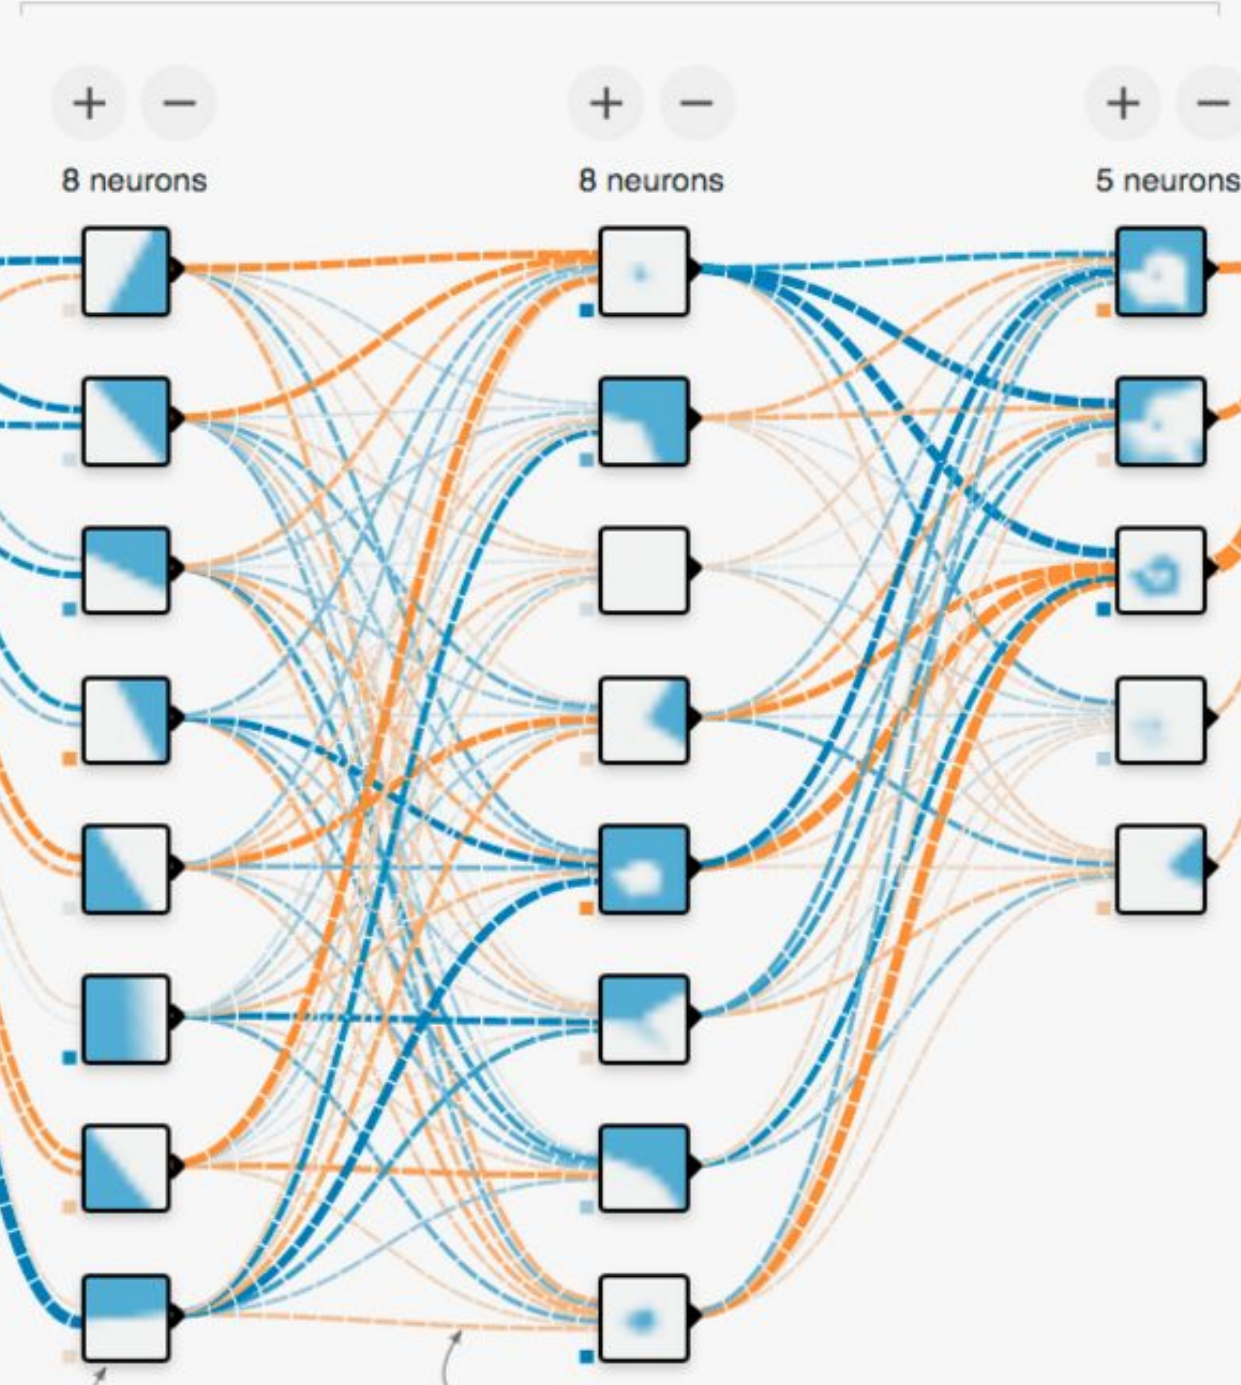


## FEATURES

Which properties  
do you want to  
feed in?




## 3 HIDDEN LAYERS



## OUTPUT


Training loss 0.001






## FEATURES

Which properties  
do you want to  
feed in?



## OUTPUT


Training loss 0.001



Colors shows  
data, neuron and  
weight values.



Models can learn which features to look at and which are most useful





Video Name: T-ICML-O\_1\_I6\_lab\_intro:\_dnn\_models

Format: Studio with Presenter

Presenter: Evan Jones

# Lab

---

## Image Classification with a Deep Neural Network Model

Evan Jones

# Lab Steps

---

- Import the training dataset of MNIST handwritten images
- Reshape and preprocess the image data
- Setup your neural network model with 10 classes (one for each possible digit 0 through 9)
- Define and create your EstimatorSpec in tensorflow to create your custom estimator
- Define and run your train\_and\_evaluate function to train against the input dataset of 60,000 images and evaluate your model's performance

Video Name: T-ICML-O\_1\_I7\_lab\_solution:\_dnn\_models

Format: Studio with Presenter

Presenter: Evan Jones




Video Name: T-ICML-O\_1\_I8\_dropout

Format: Studio with Presenter

Presenter: Evan Jones

An infinitely large DNN  
could classify anything right?

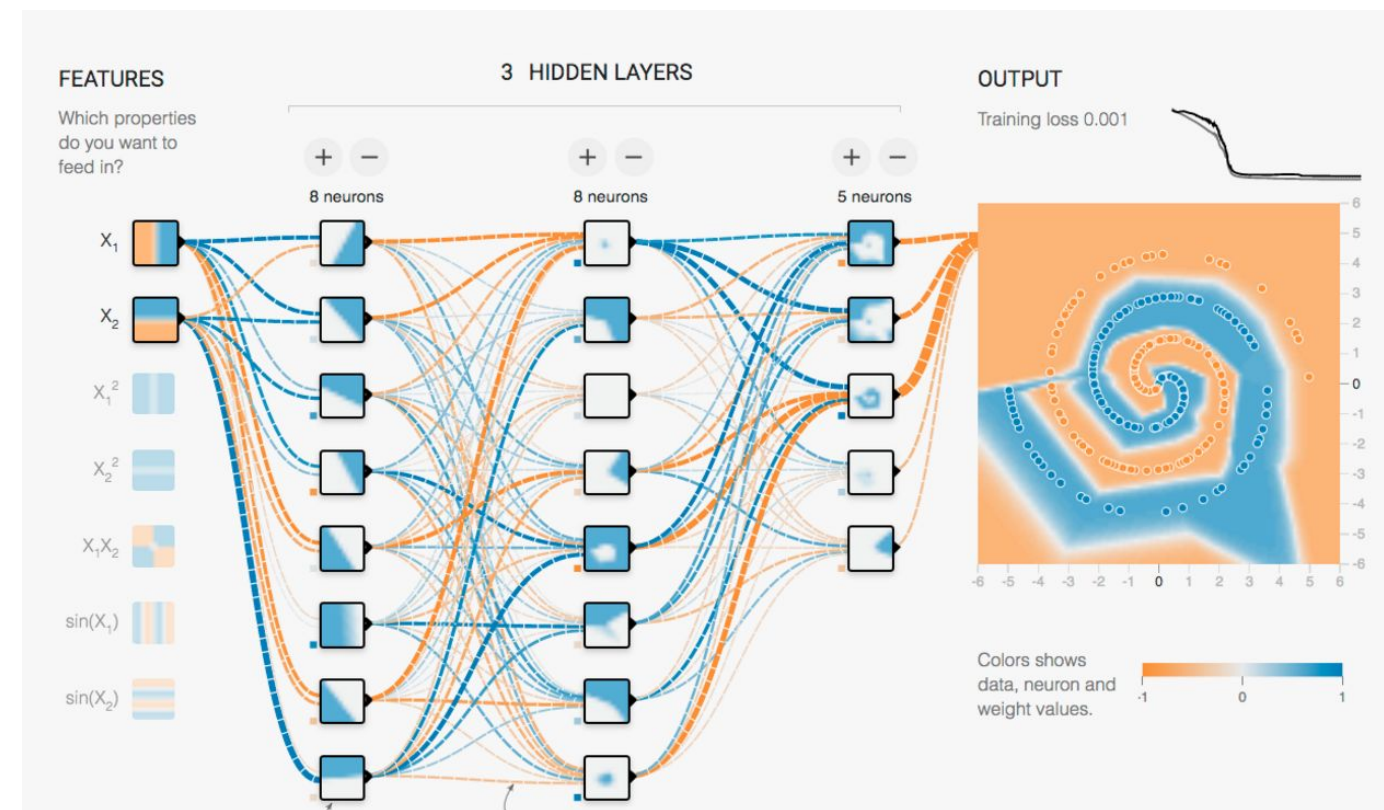


An infinitely large DNN  
could ~~classify~~ *memorize* anything

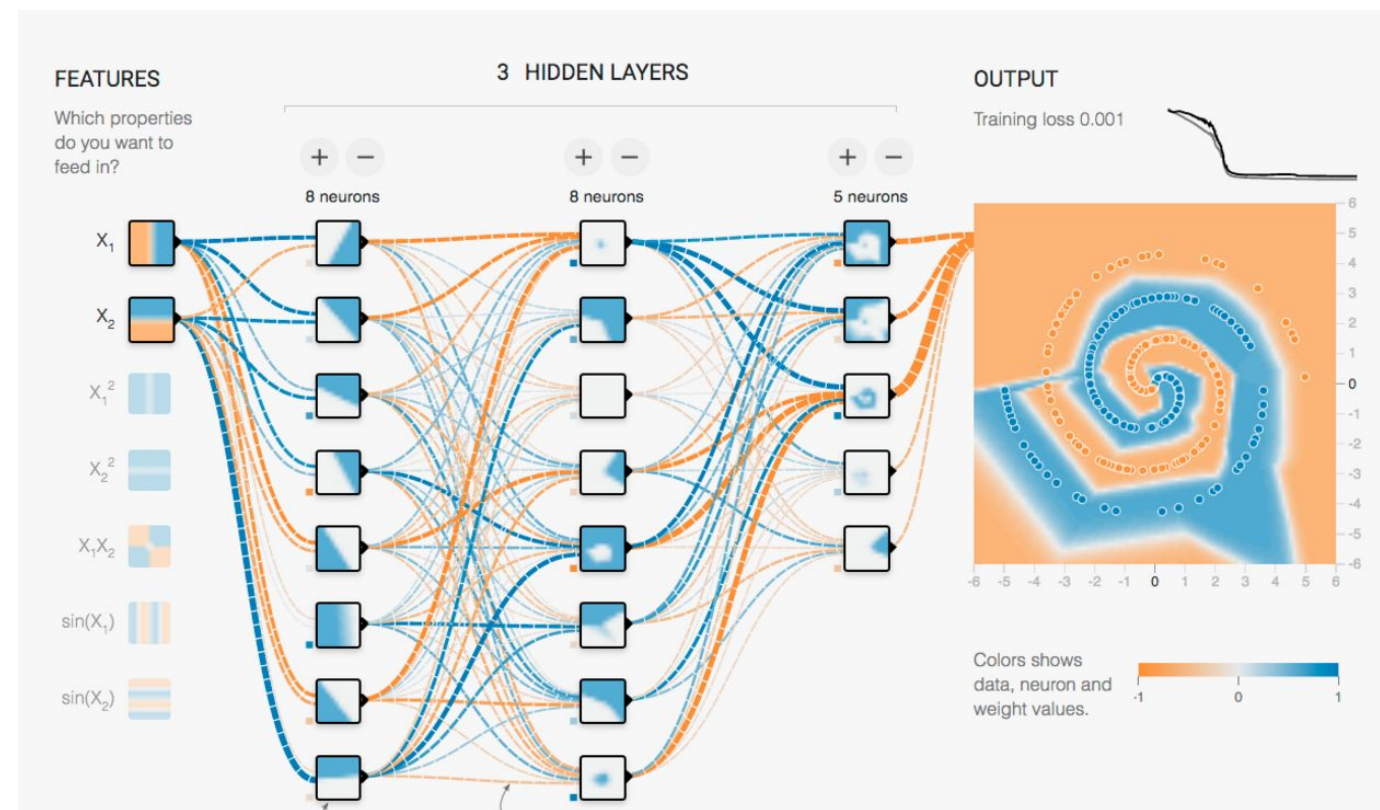




# Why not have a really large DNN?



# Why not have a really large DNN?



1. Computational power
2. Training time
3. Likely to overfit

Recall: How can we combat  
overfitting in a DNN?

Combat overfitting with  
**regularization**


Which form of regularization  
is only used with neural  
networks?

1. Dropout
2. L1
3. L2


Which form of regularization  
is only used with neural  
networks?

1. **Dropout**
2. L1
3. L2

Use dropout to avoid  
overfitting in a DNN



Dropout can be thought of  
as an ensemble of NNs;  
useful in real-world  
problems





$$\lambda \sum_{i=1}^k |\omega_i|$$

L1 regularization term

$$\lambda \sum_{i=1}^k \omega_i^2$$

L2 regularization term

# Adding a dropout layer to your DNN

```
layer_before = tf.layers.dense(previous_layer, 30)

dropout_layer = tf.layers.dropout(layer_before,
    rate=drop_probability,
    # only dropout when training
    training=(mode == tf.estimator.ModeKeys.TRAIN))

layer_after = tf.layers.dense(dropout_layer, 15)
```

Video Name:

T-ICML-O\_1\_I9\_lab\_intro:\_dnn\_models\_with\_dropout

Format: Studio with Presenter

Presenter: Evan Jones

# Lab

---

Image Classification with a  
DNN Model with Dropout

Evan Jones

# Lab Steps

---

- Import the training dataset of MNIST handwritten images
- Reshape and preprocess the image data
- Setup your neural network model with 10 classes (one for each possible digit 0 through 9)
- Add a Dropout layer
- Define and create your EstimatorSpec in tensorflow to create your custom estimator
- Define and run your train\_and\_evaluate function to train against the input dataset of 60,000 images and evaluate your model's performance

Video Name:

T-ICML-O\_1\_I10\_lab\_solution:\_dnn\_models\_with\_dropout

Format: Studio with Presenter

Presenter: Evan Jones

