

File Name: T-ICML-O_2_I1_introduction

Format: Presenter in Studio

Presenter: Carl Osipov



Implementing Convolutional Neural Networks for Image Classification Tasks

Carl Osipov

Agenda

Convolutional Neural Networks

Understanding Convolutions

Parameters (padding, stride)

Pooling layers

Implementing CNNs

Architecture walkthrough

Agenda

Convolutional Neural Networks

Understanding Convolutions

Parameters (padding, stride)

Pooling layers

Implementing CNNs

Architecture walkthrough

Agenda

Convolutional Neural Networks

Understanding Convolutions

Parameters (padding, stride)

Pooling layers

Implementing CNNs

Architecture walkthrough

Agenda

Convolutional Neural Networks

Understanding Convolutions

Parameters (padding, stride)

Pooling layers

Implementing CNNs

Architecture walkthrough

Agenda

Convolutional Neural Networks

Understanding Convolutions

Parameters (padding, stride)

Pooling layers

Implementing CNNs

Architecture walkthrough

Agenda

Convolutional Neural Networks

Understanding Convolutions

Parameters (padding, stride)

Pooling layers

Implementing CNNs

Architecture walkthrough

What are Convolutional
Neural Networks (CNNs)?



8 Megapixel resolution

3264 (w) x 2448 (h) x 3 (RGB) =

23,970,816 per image

← 3264px Width →

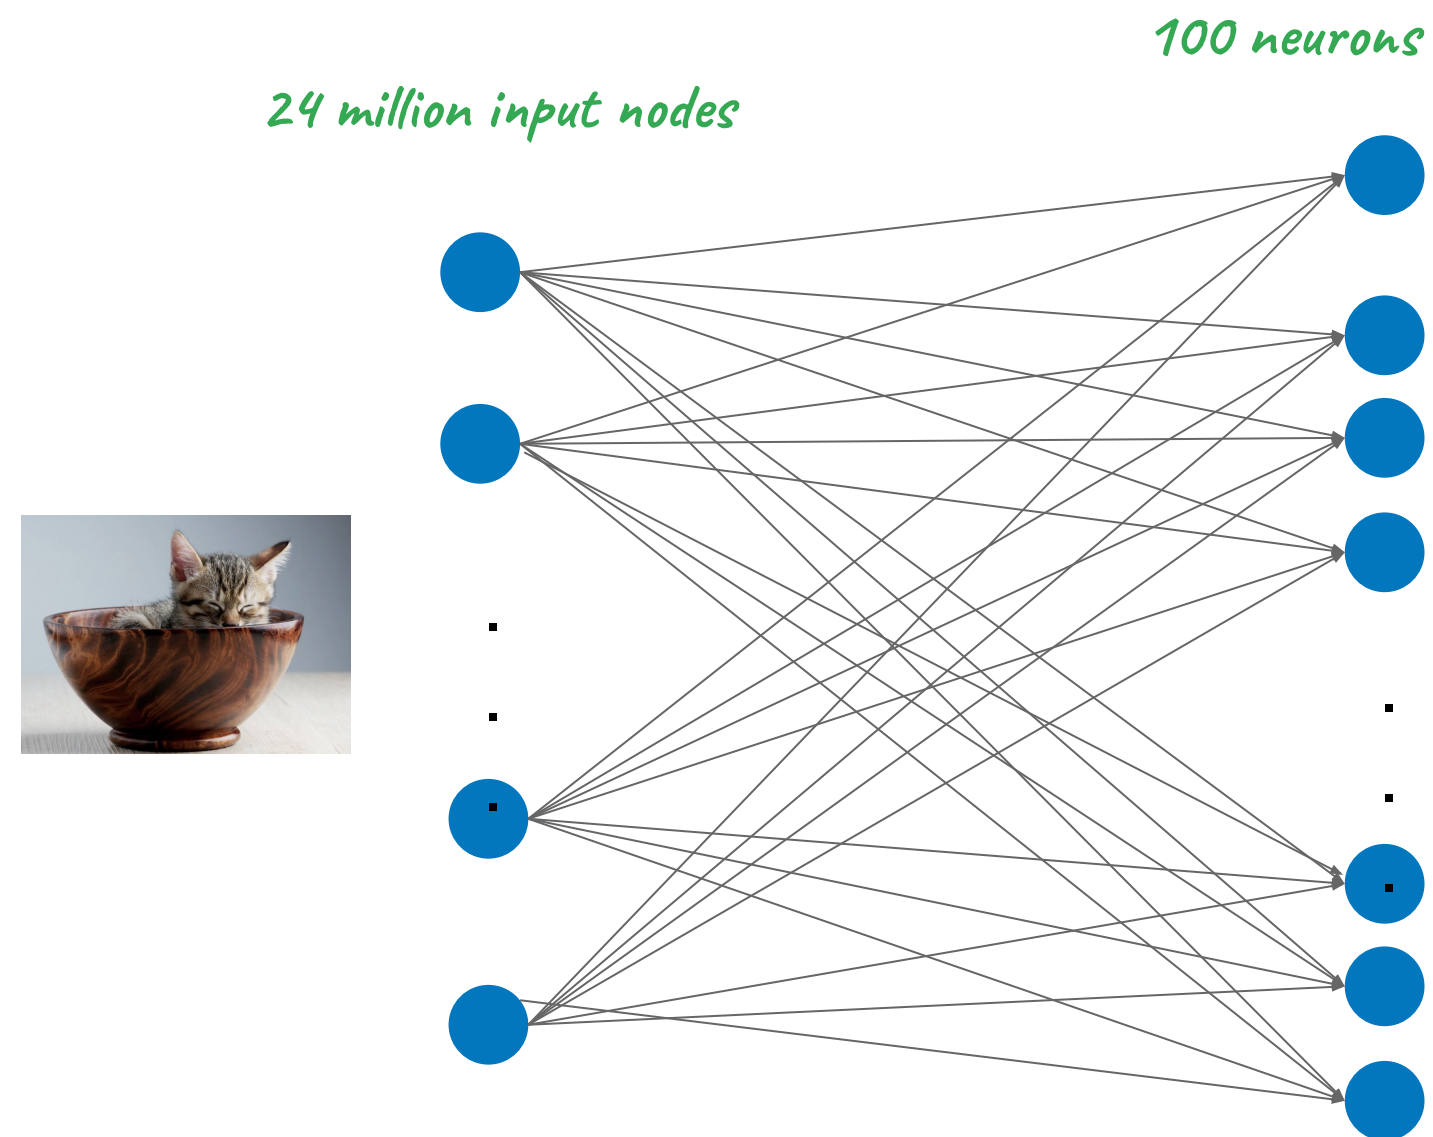
← 2448px Height →




23,970,816 * 100 neurons =

23 Billion weights

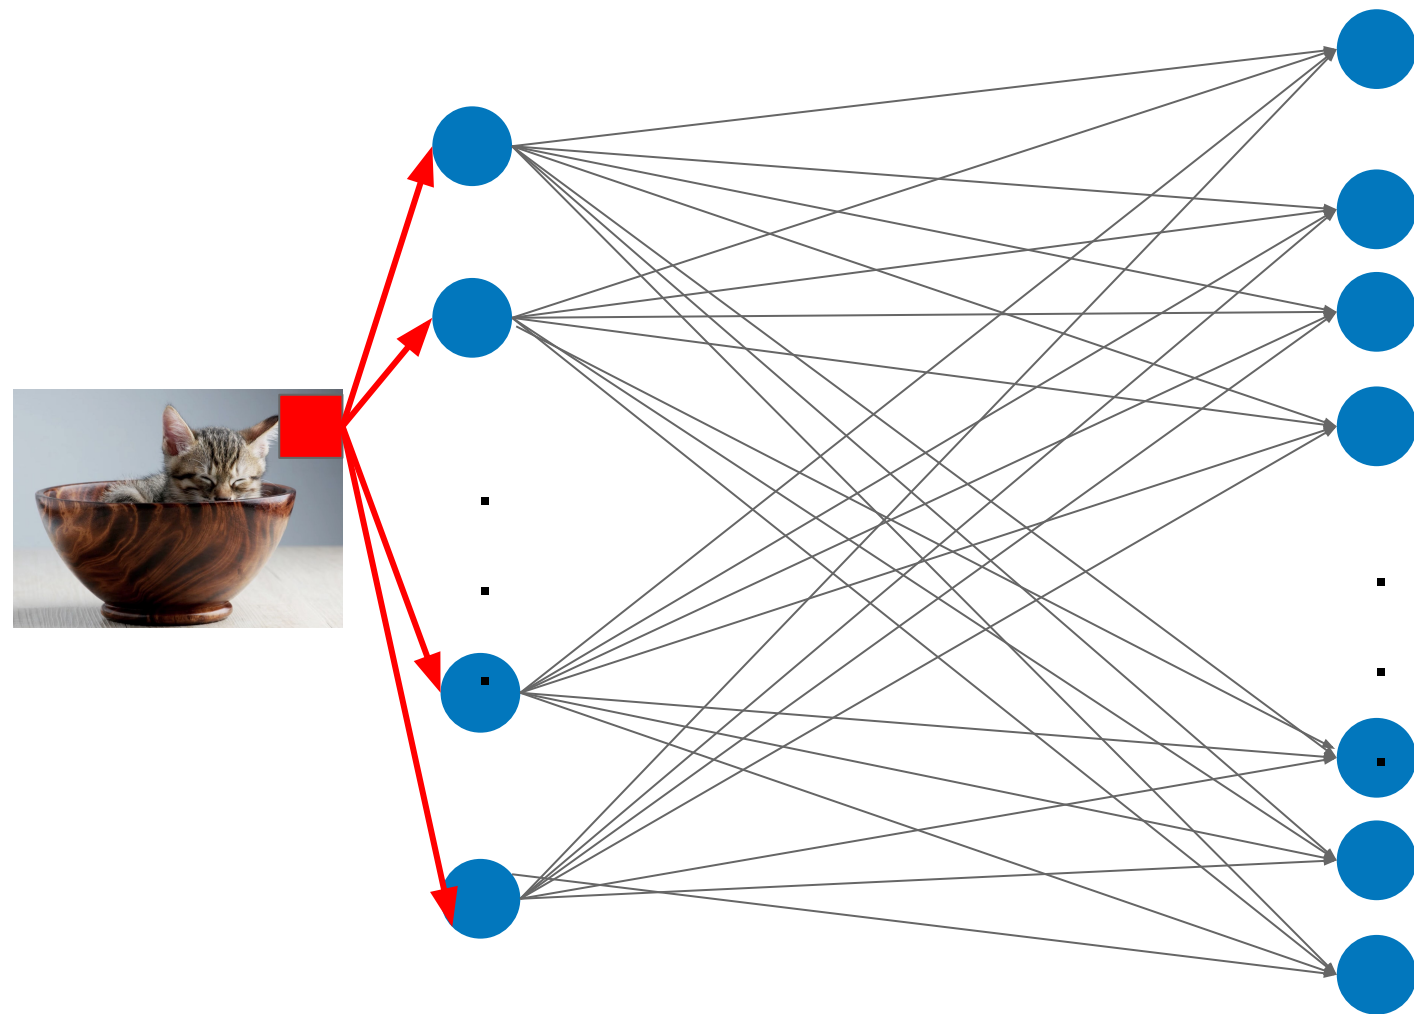
Using only a DNN for image classification
could have billions of weights




Using only a DNN for image classification
could have billions of weights



Using only a DNN for image classification
could have billions of weights

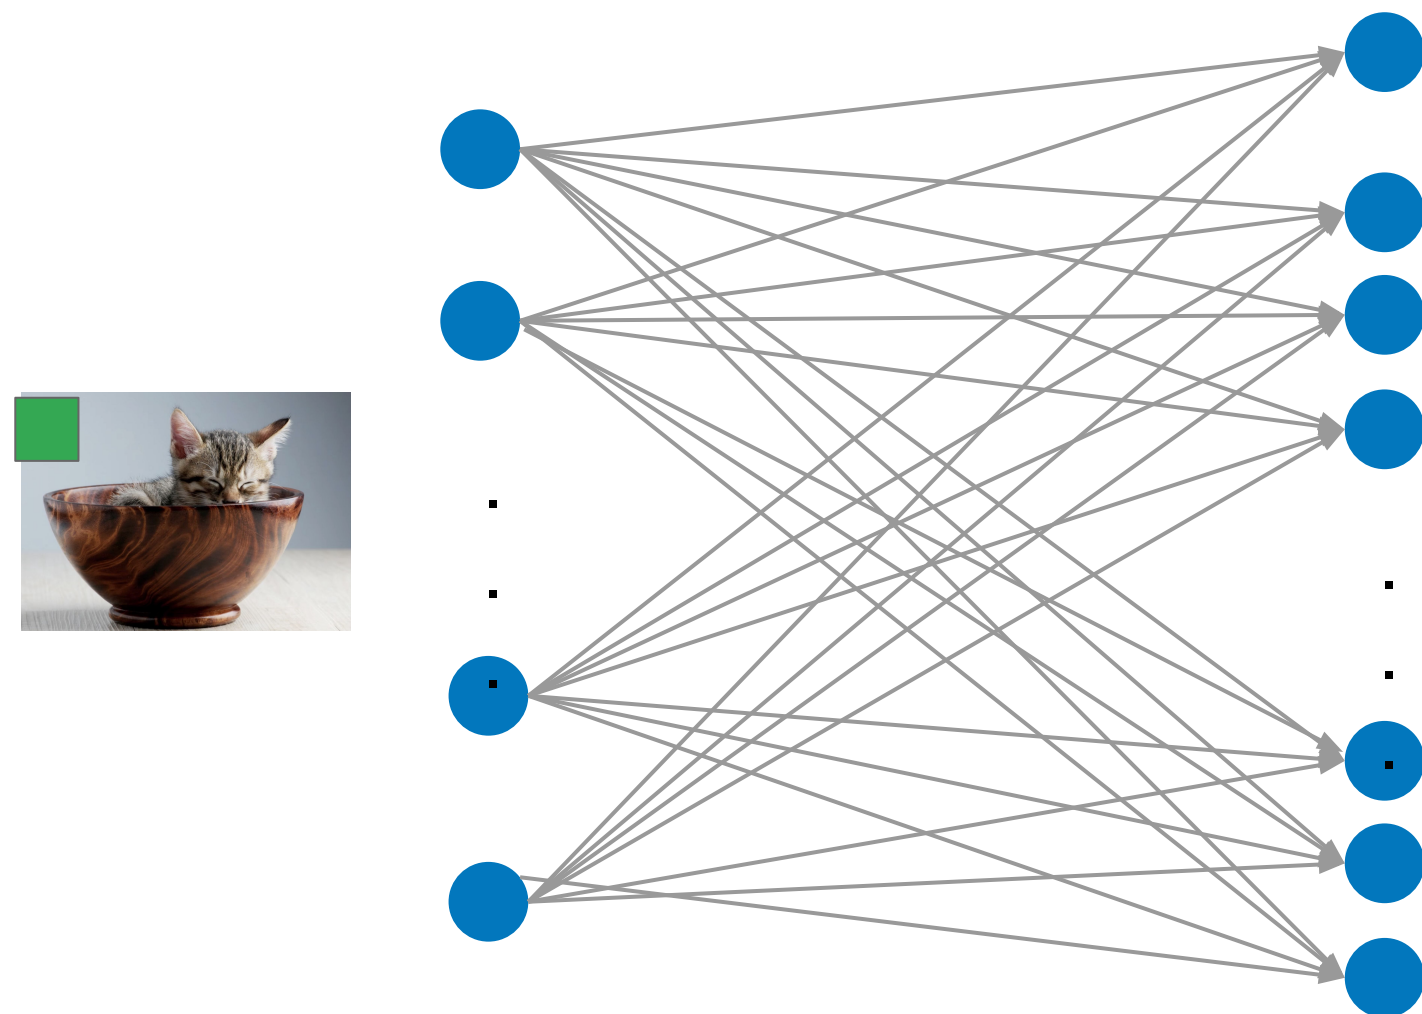


Using only a DNN for image classification
could have billions of weights




(Note: Many neurons not shown)

Using only a DNN for image classification
could have billions of weights



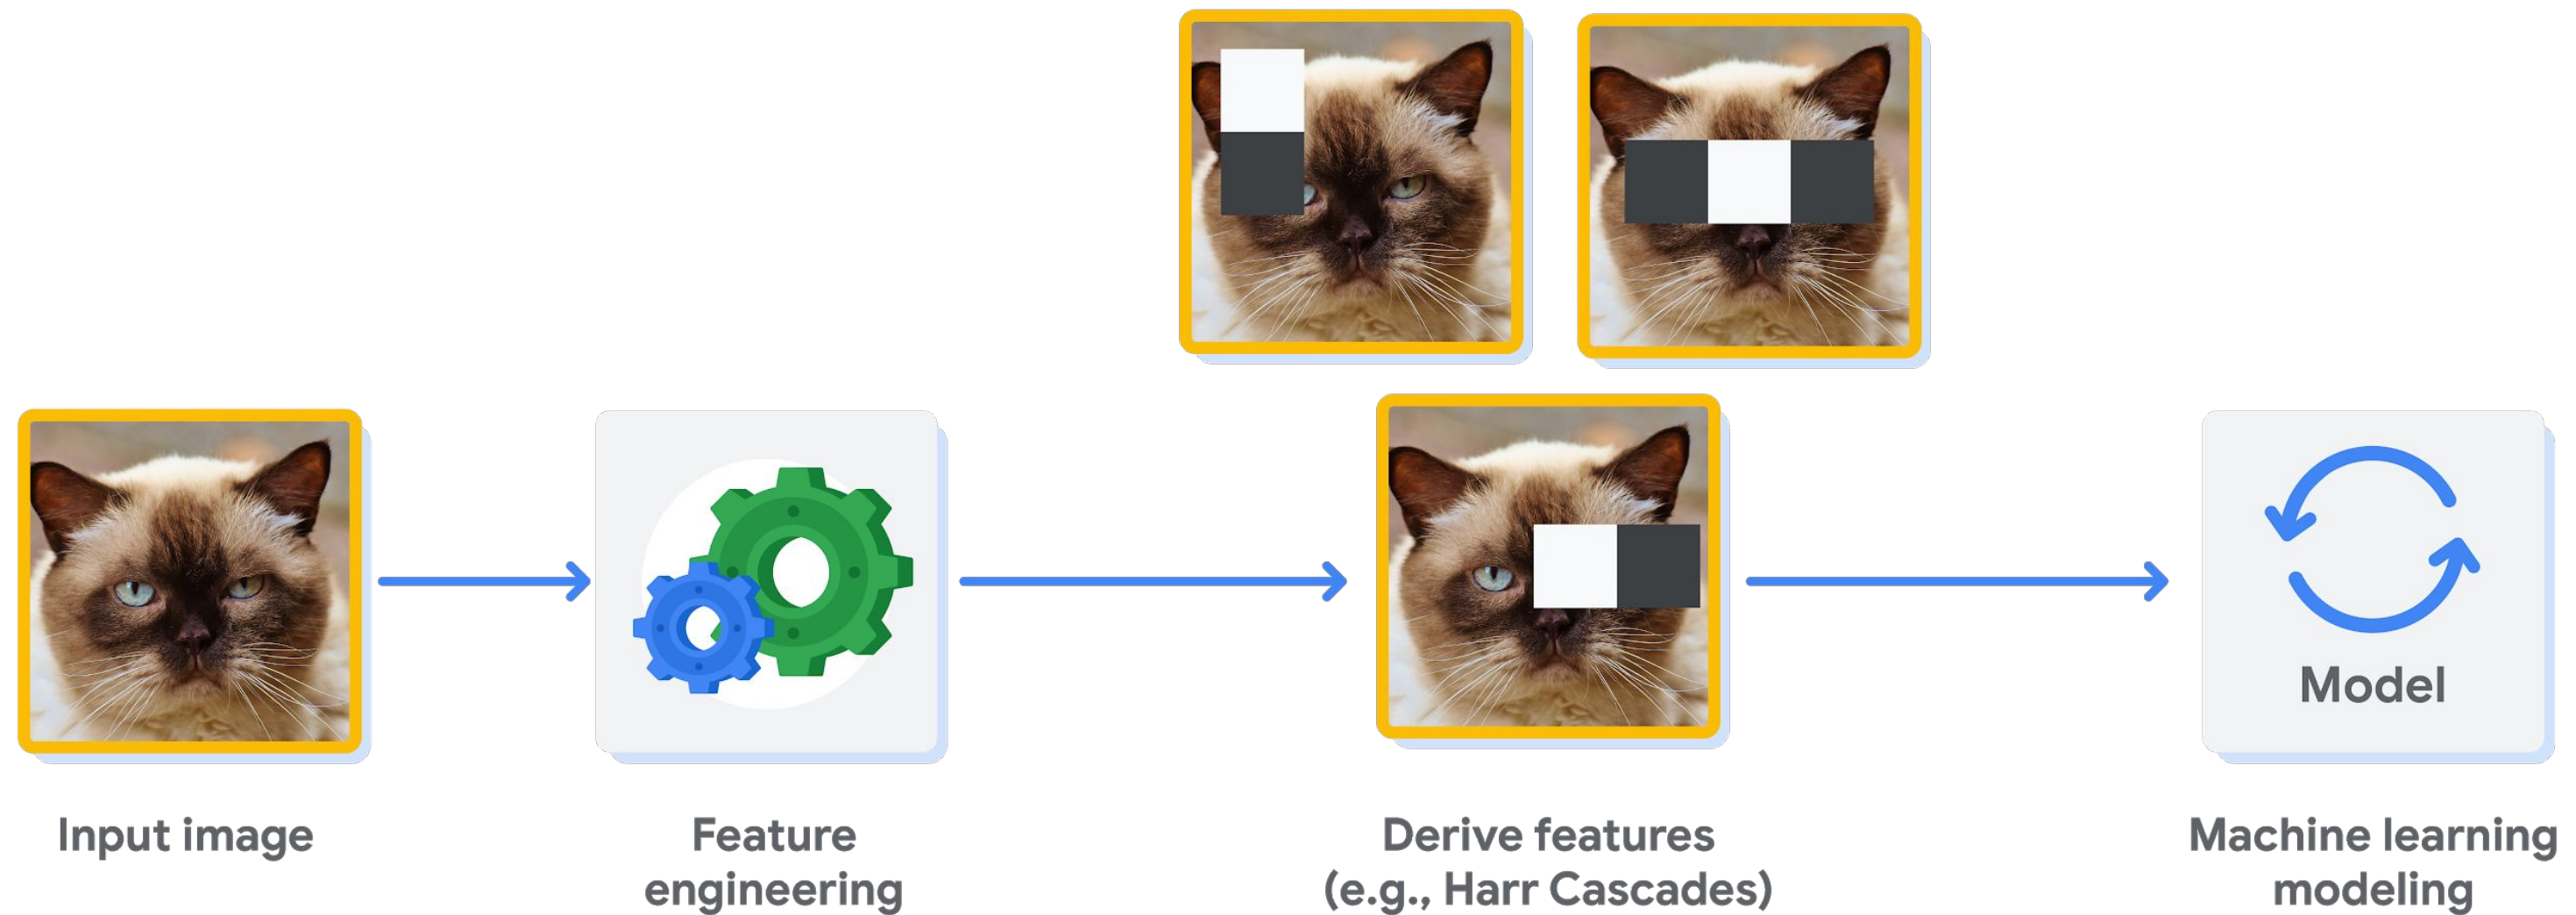
(Note: Many neurons not shown)

Using only a DNN for image classification
could have billions of weights

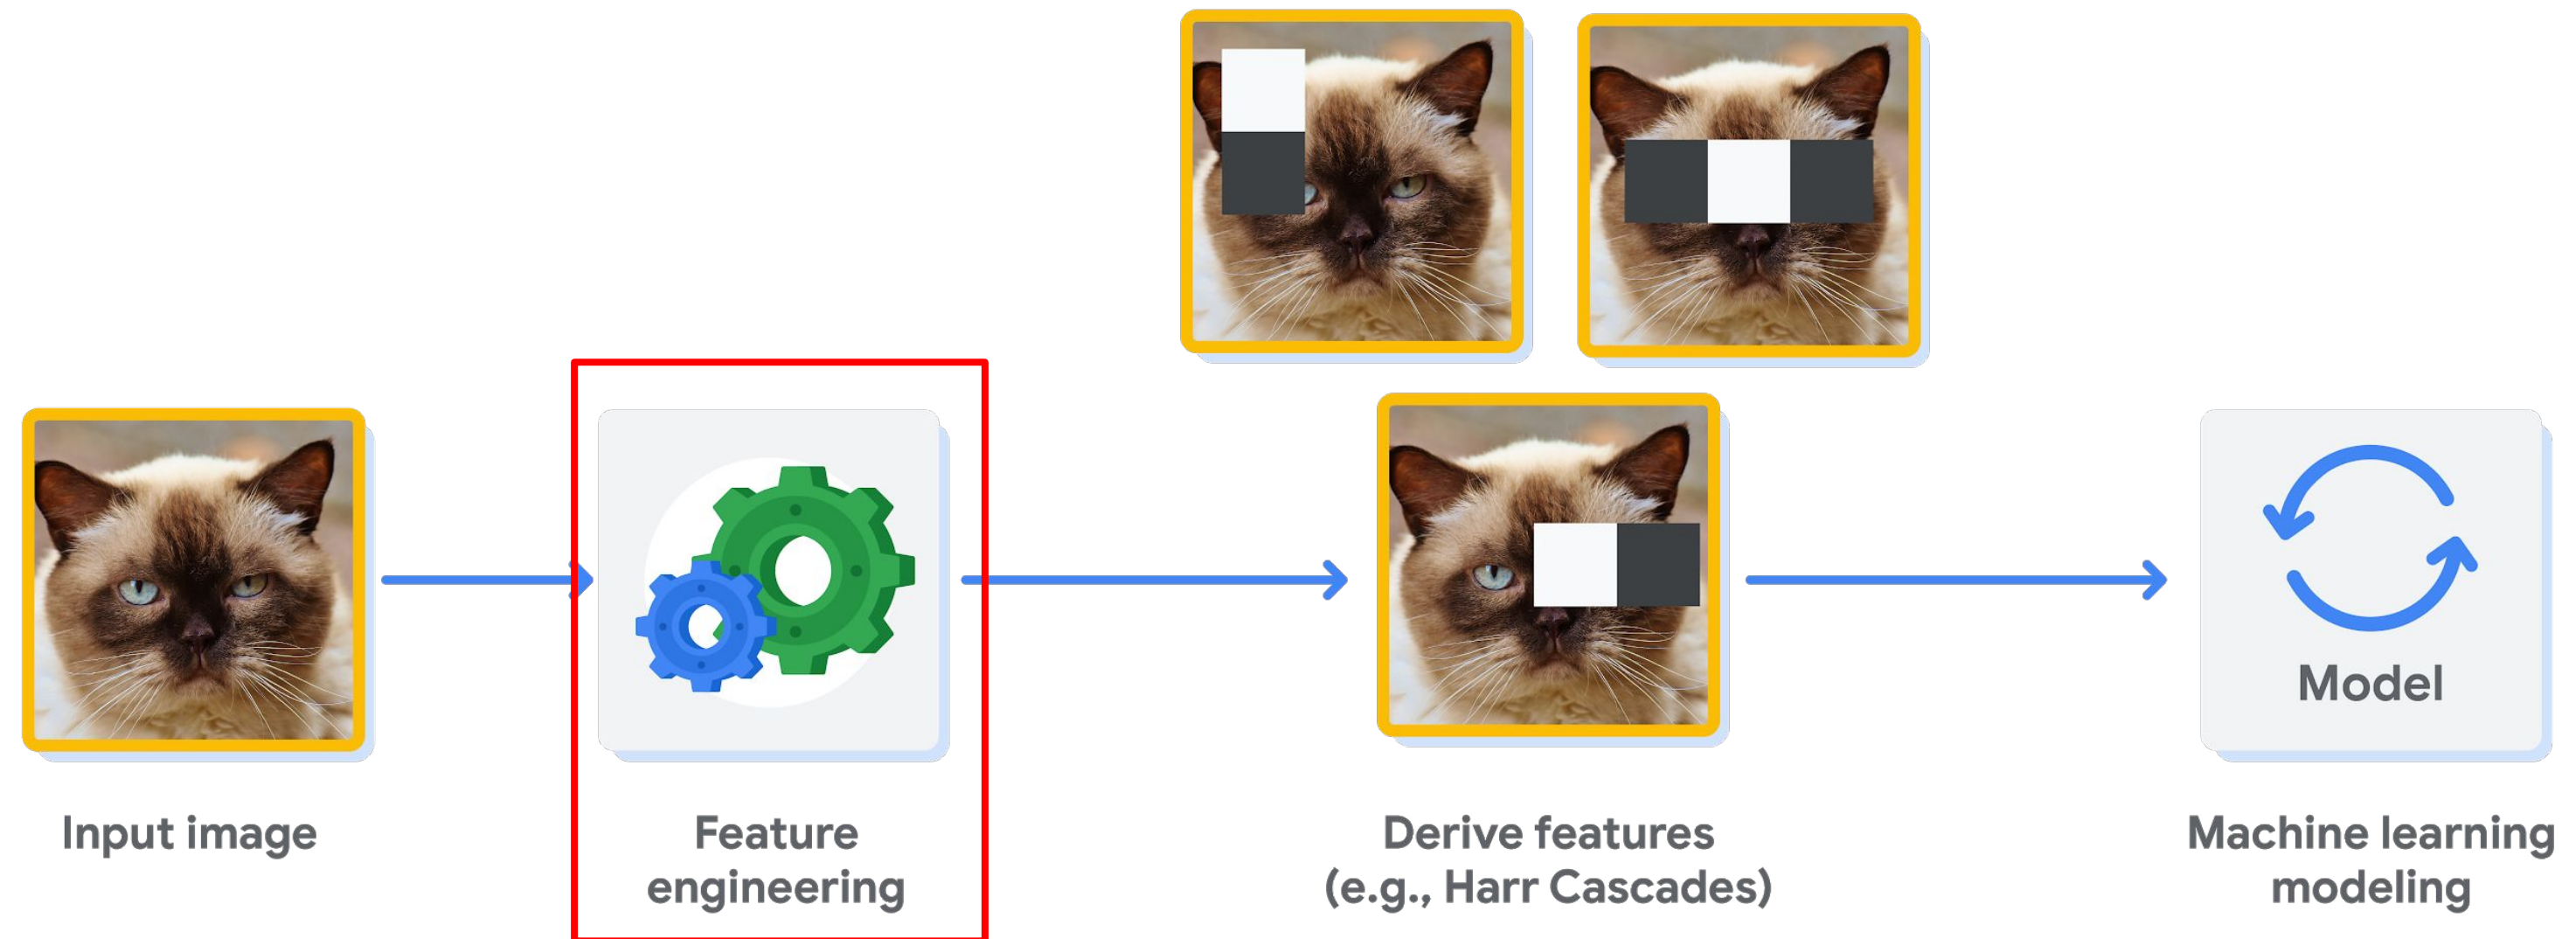


(Note: Many neurons not shown)

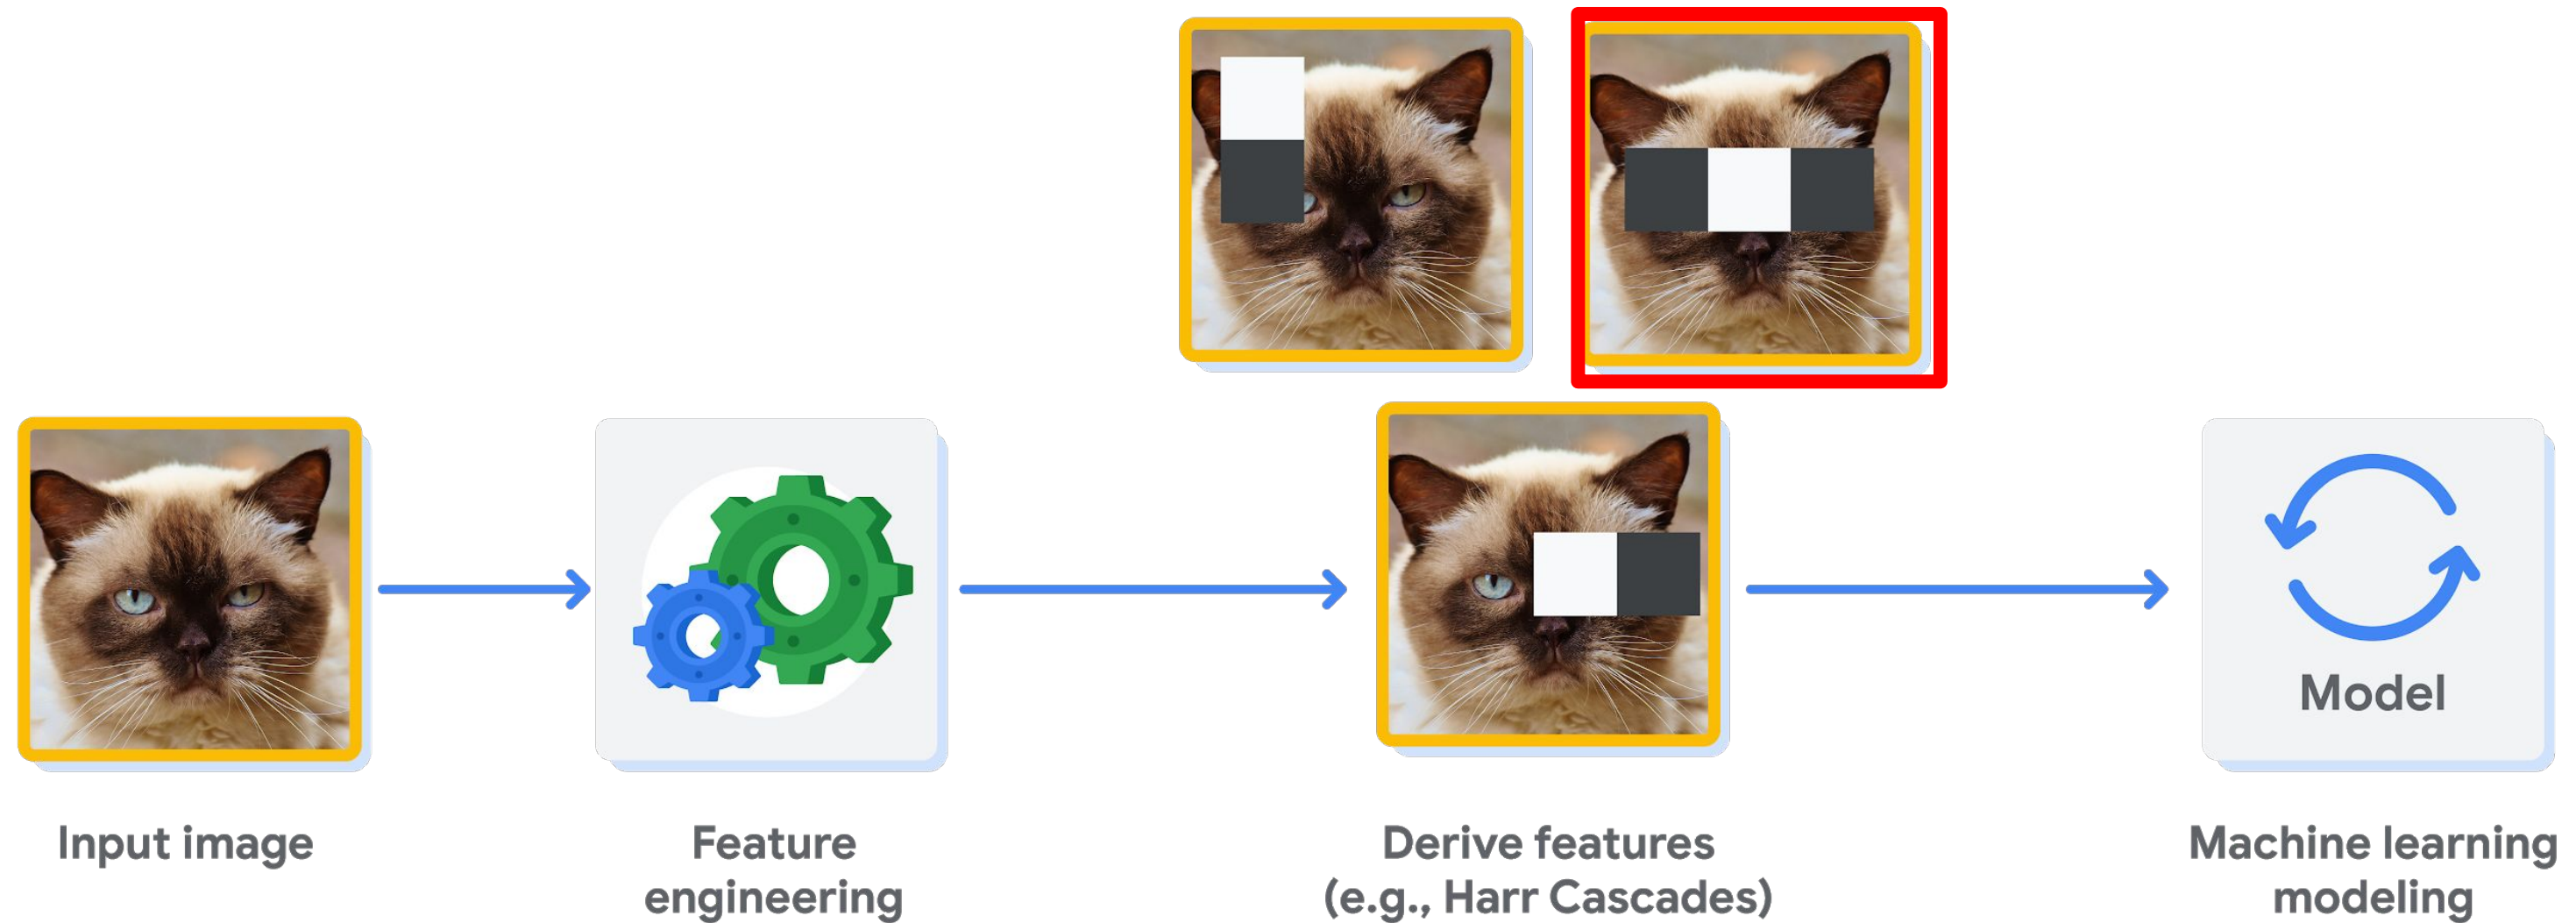
Traditionally image classification tasks utilize a feature engineering step



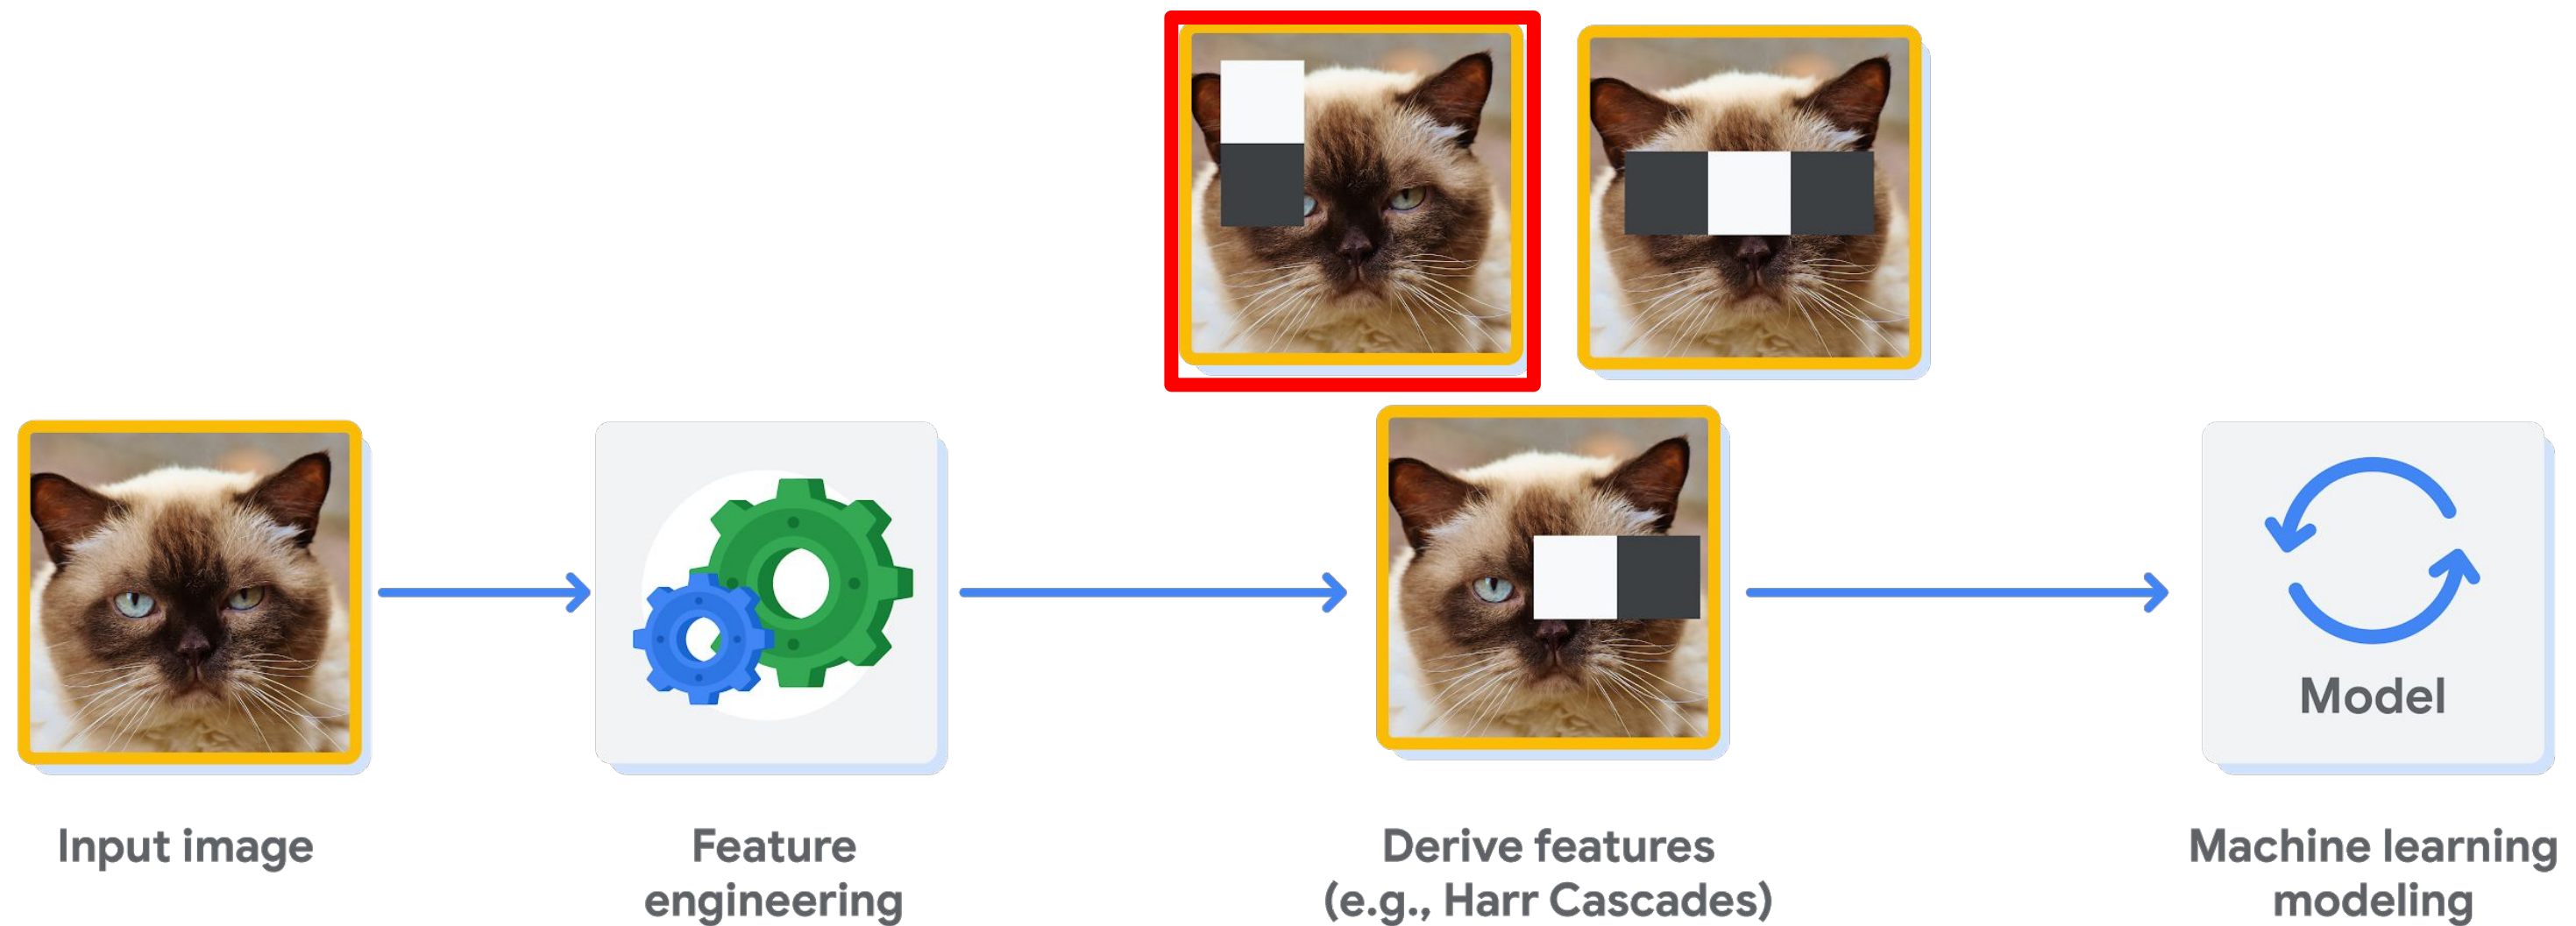
Traditionally image classification tasks utilize a feature engineering step



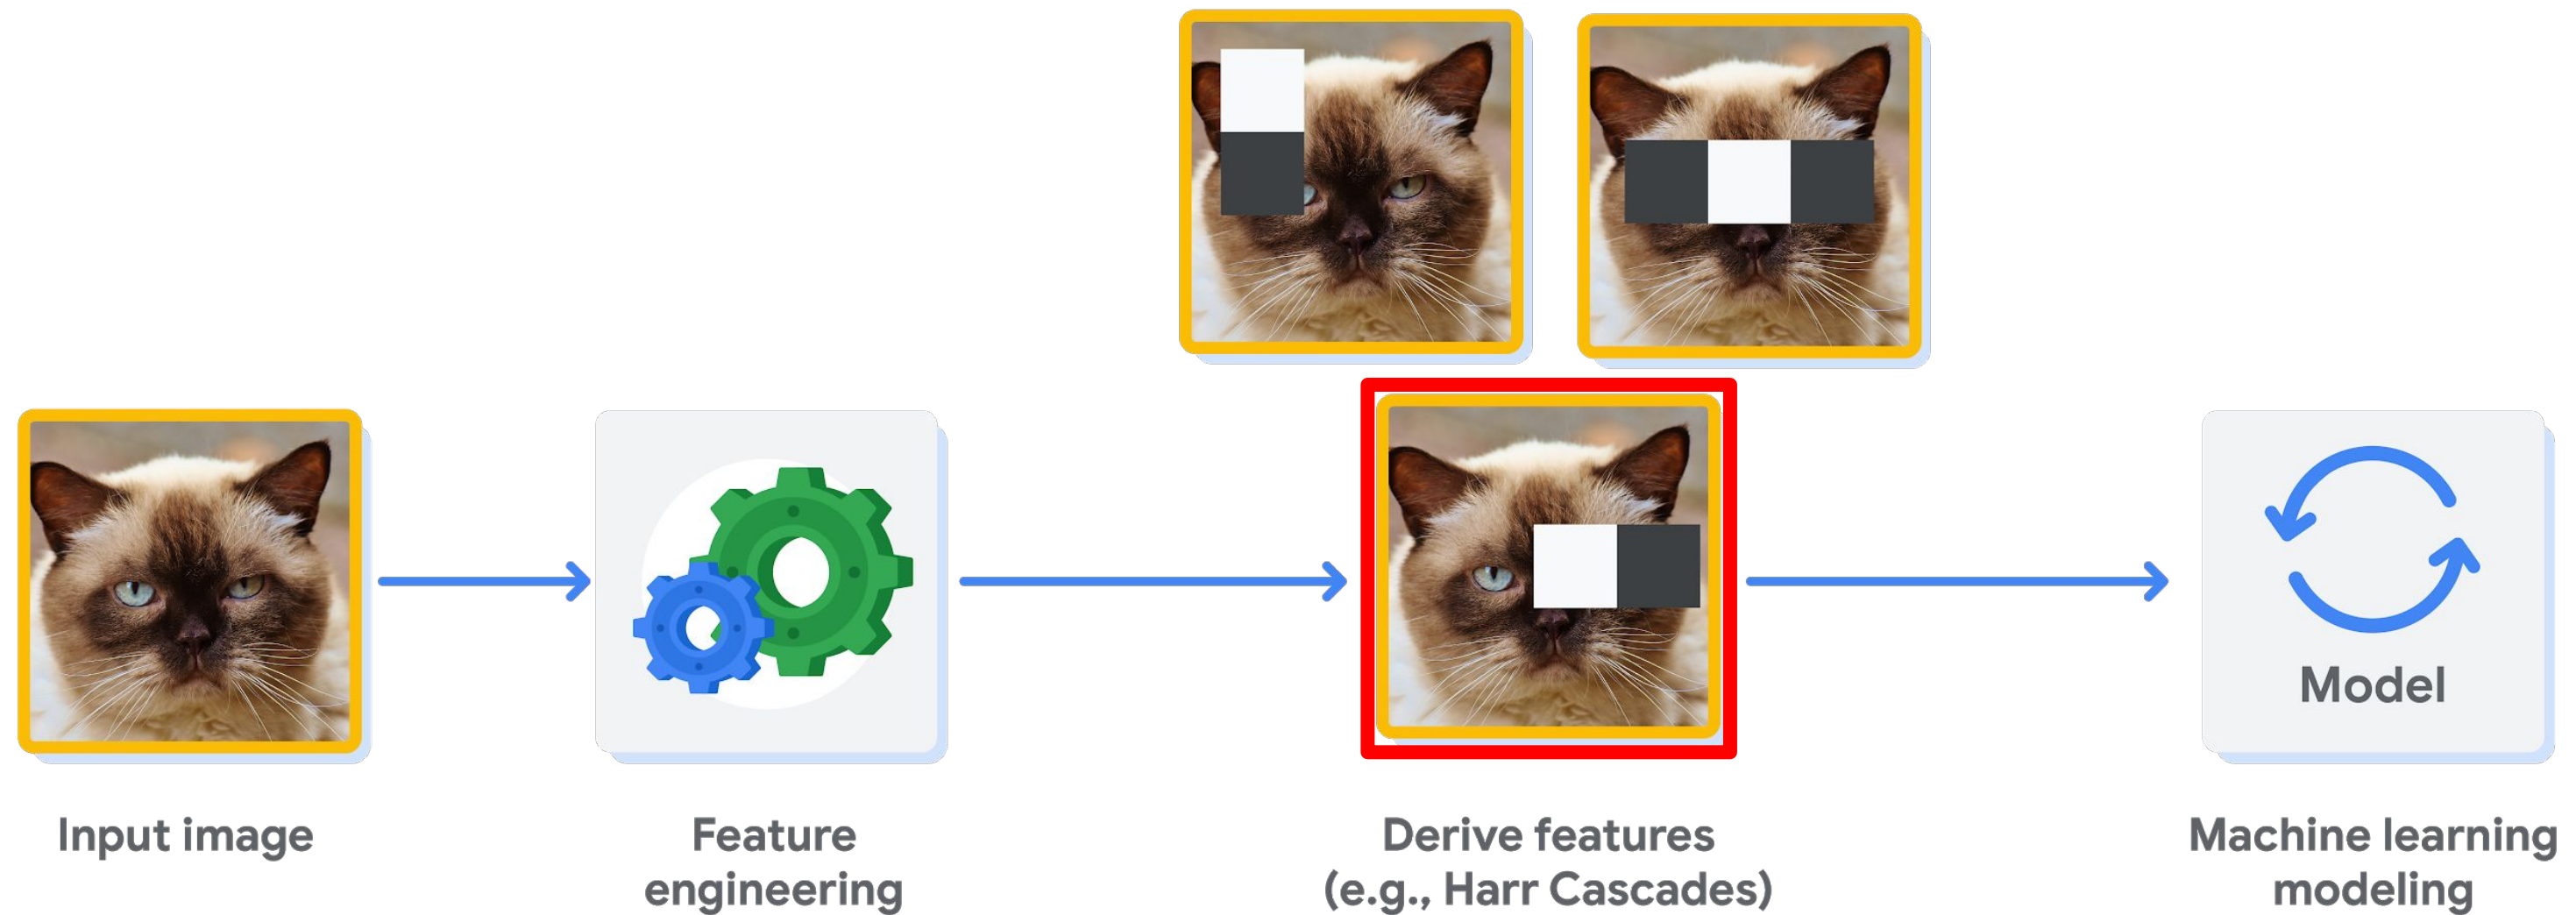
Pre-2012, ML practitioners tell the model
which features to identify



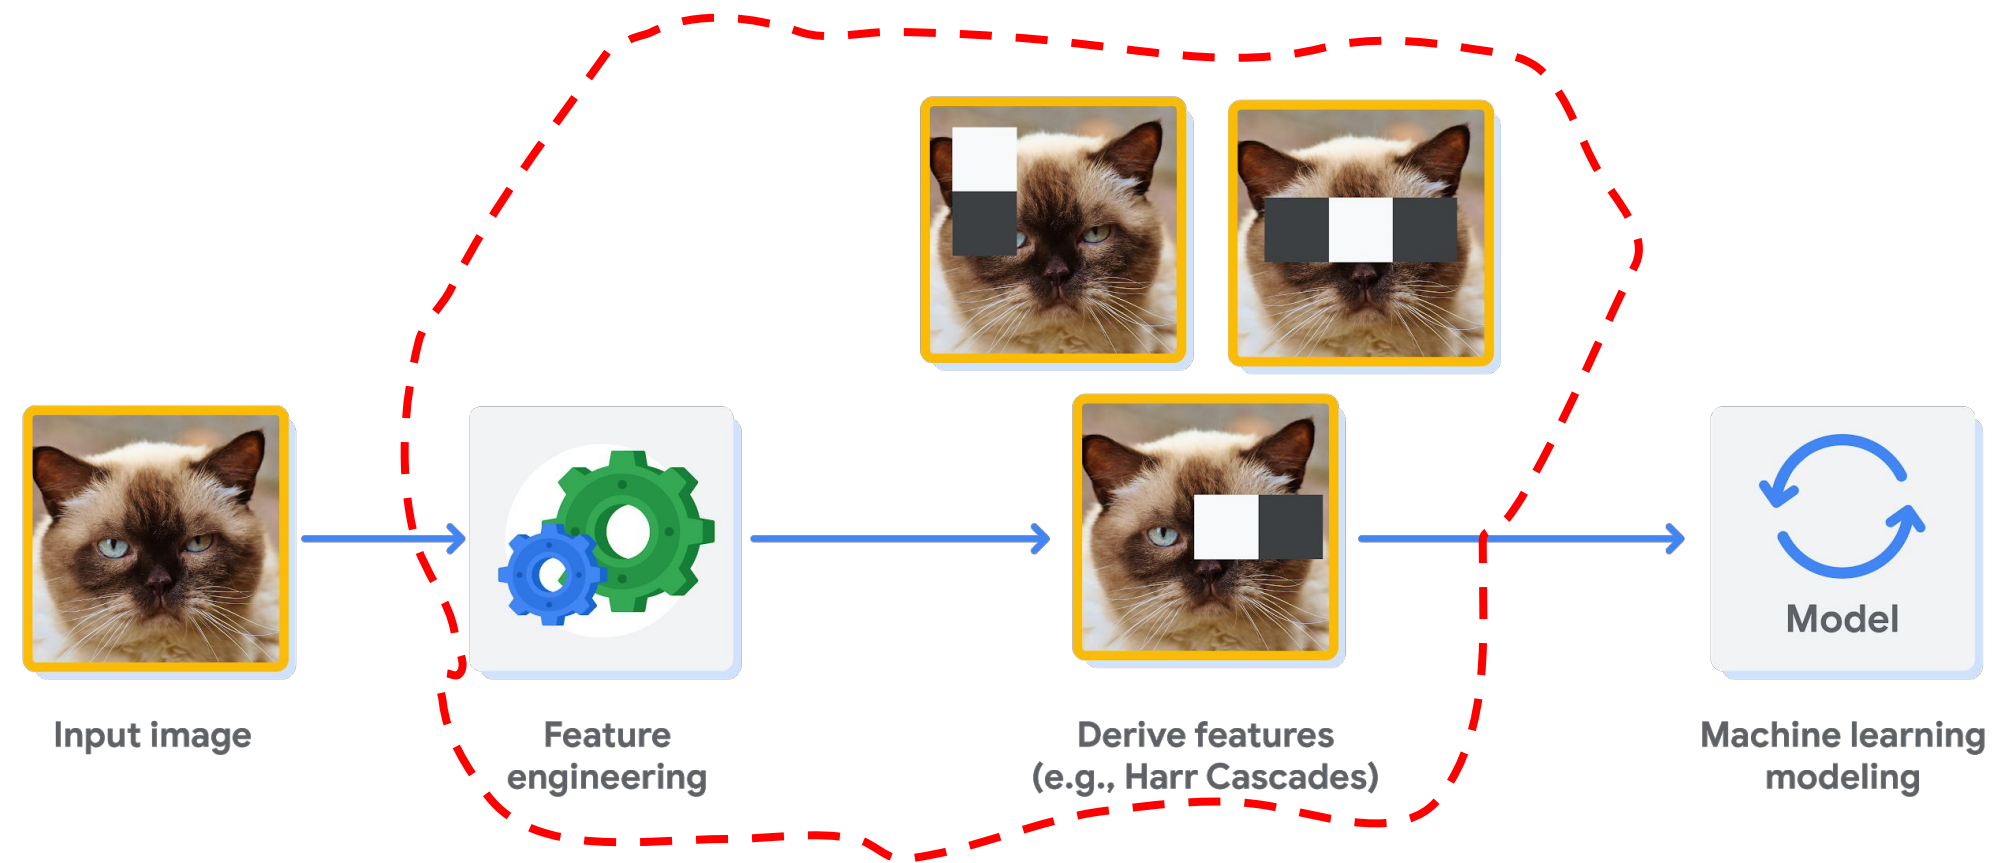
Pre-2012, ML practitioners tell the model
which features to identify



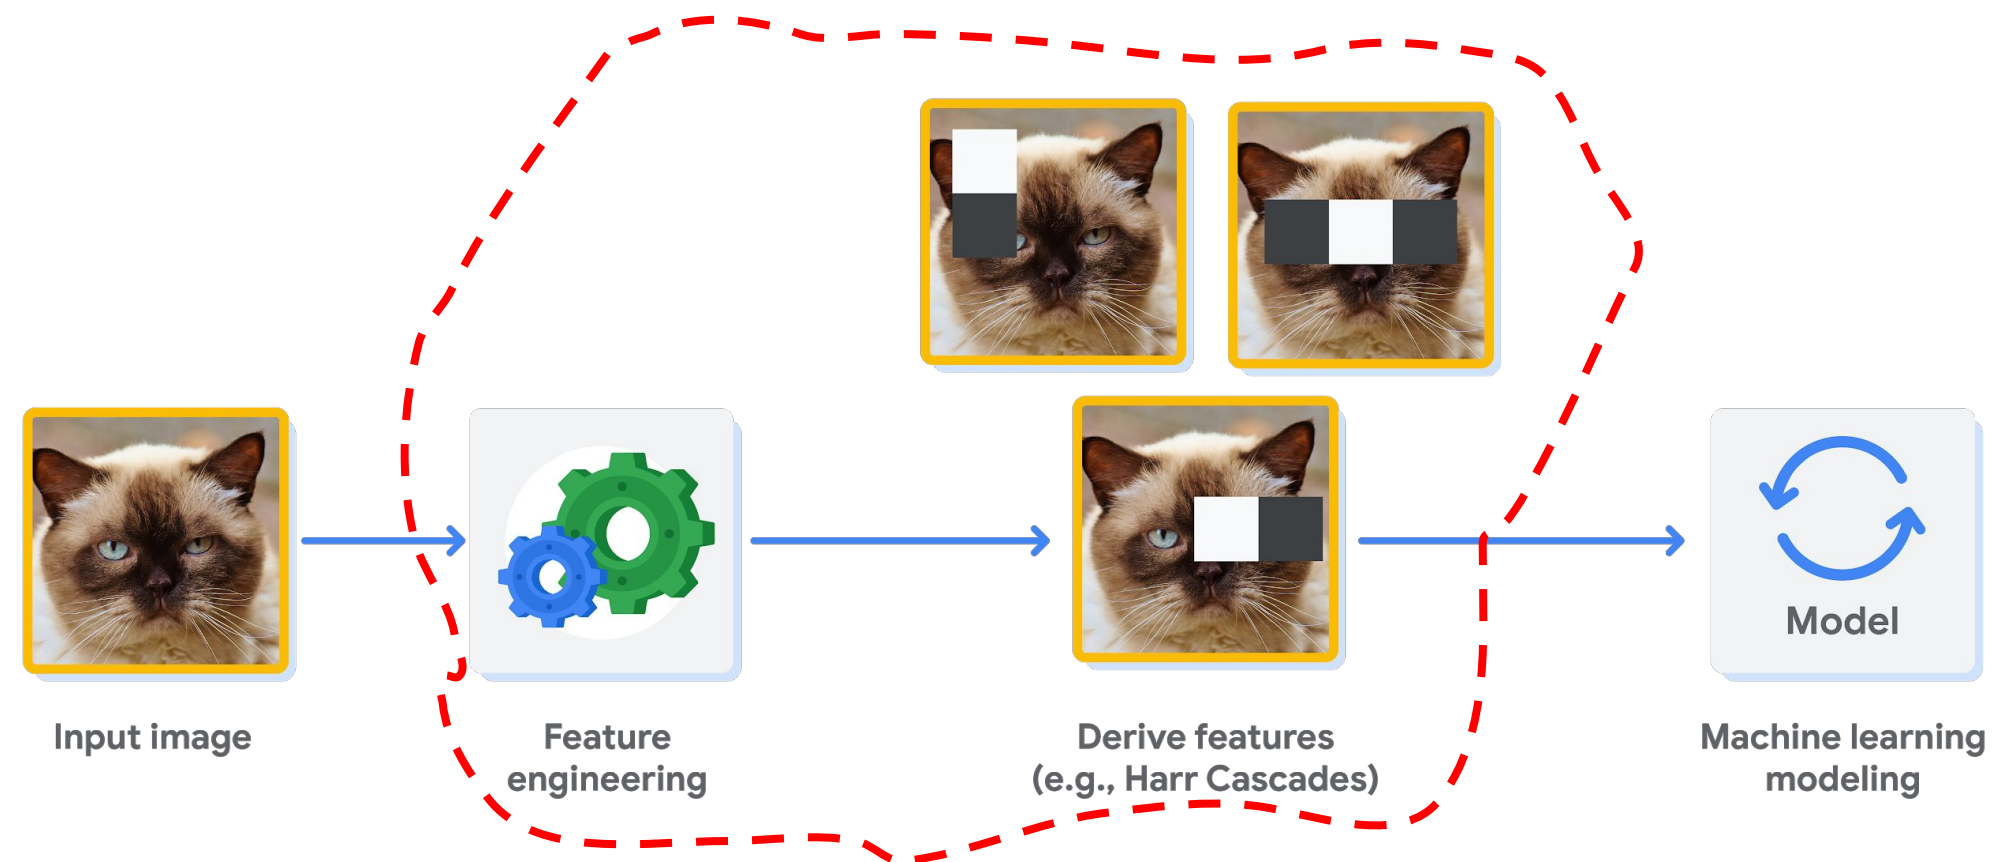
Pre-2012, ML practitioners tell the model
which features to identify



Pre-2012

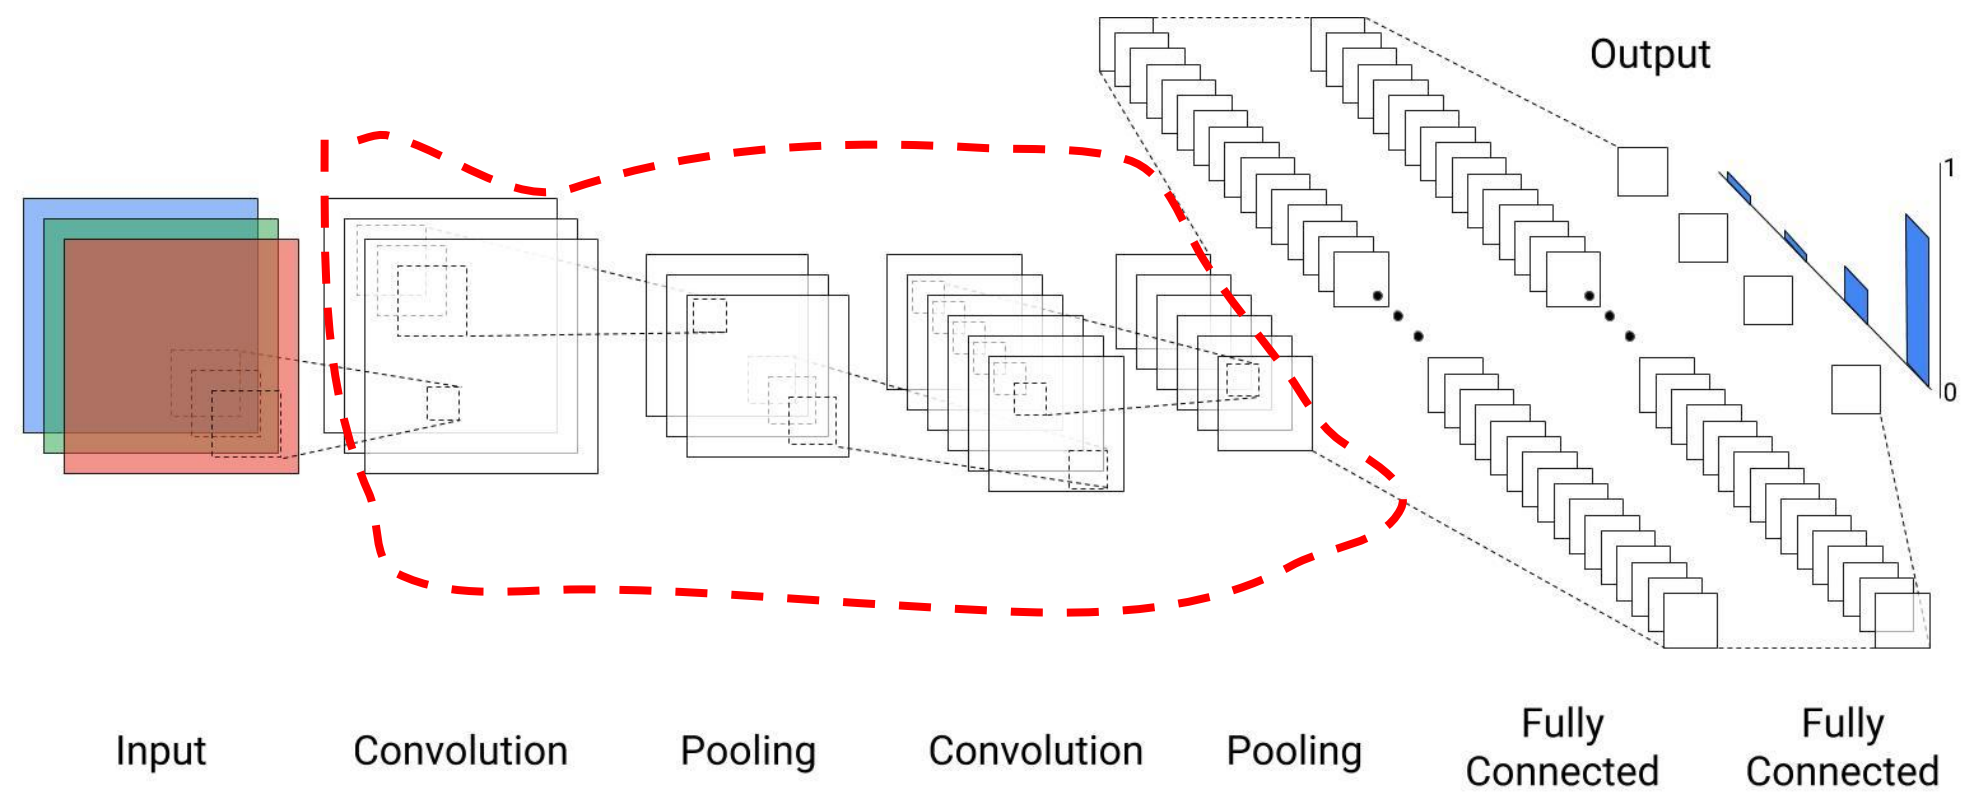


Pre-2012

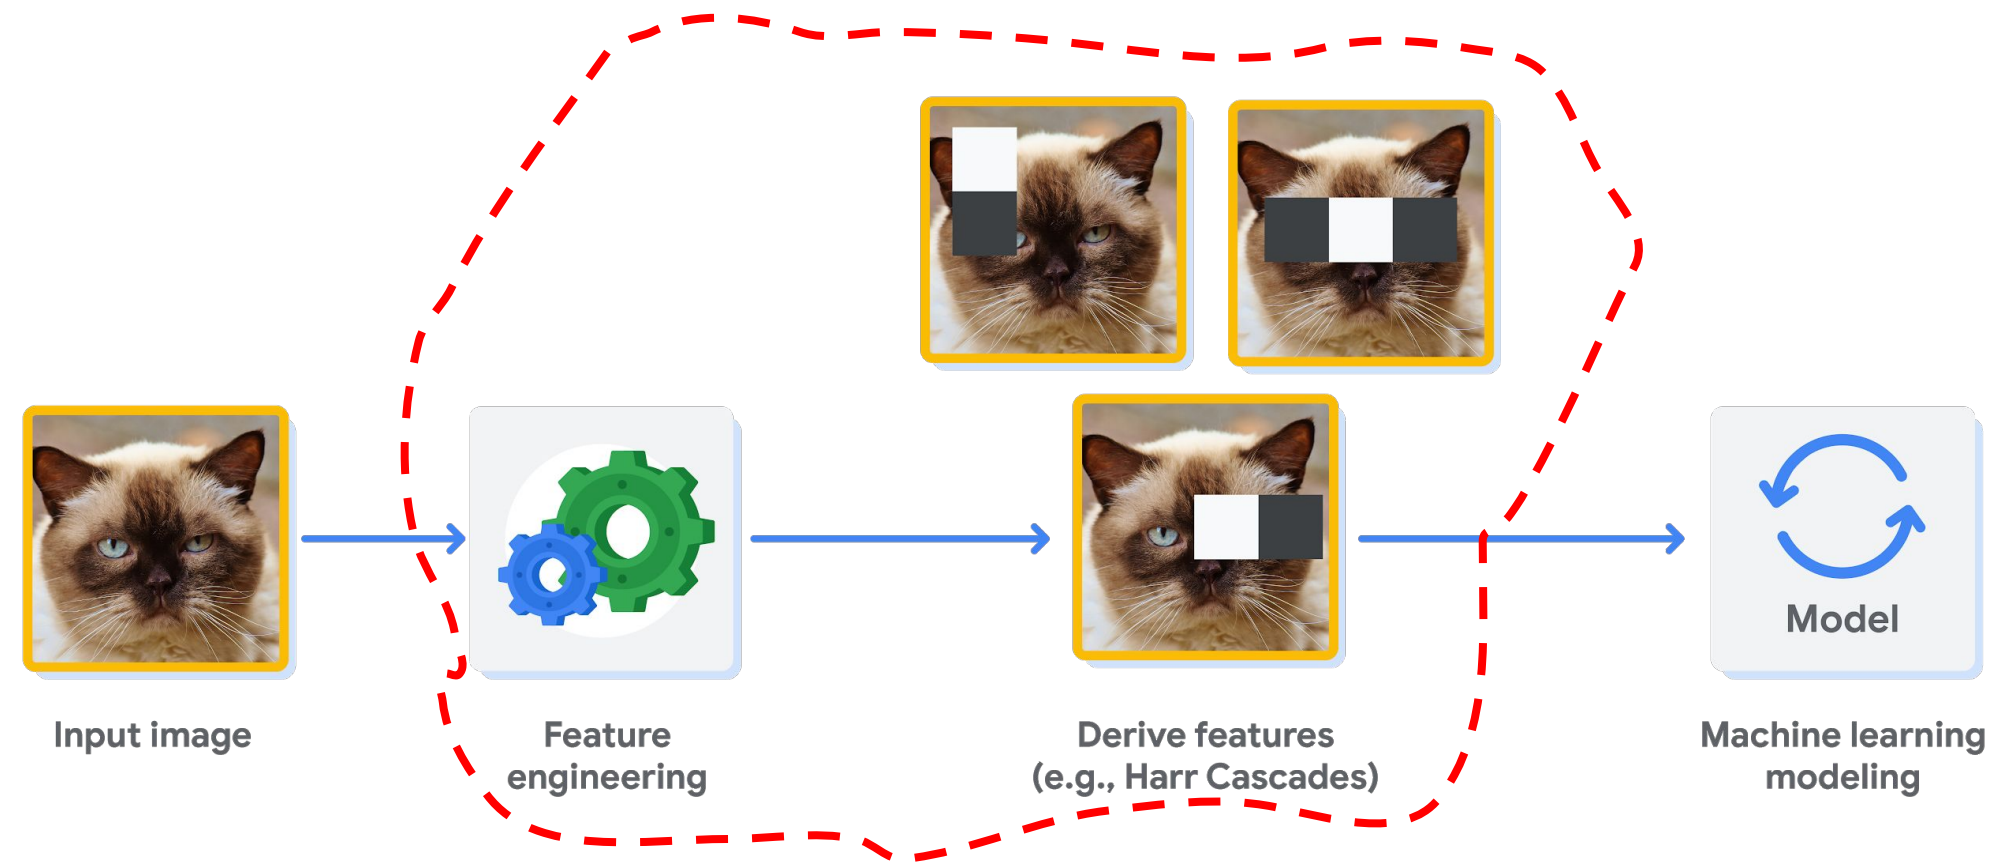


Post-2012

CNN model learns features itself during training

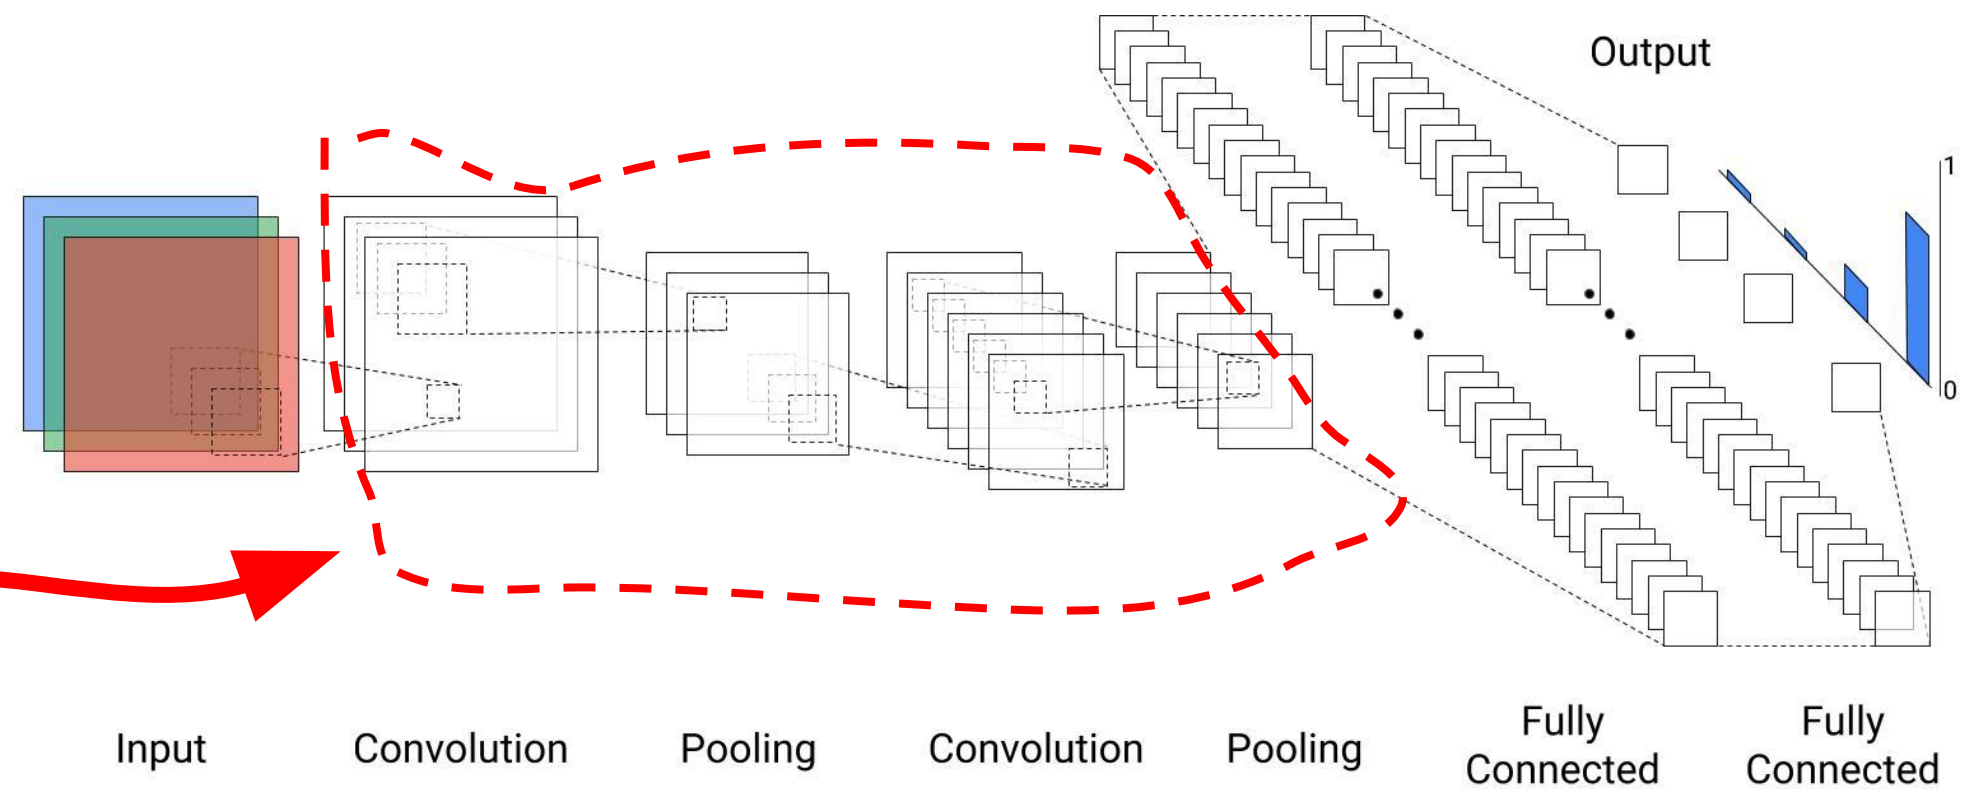


Pre-2012

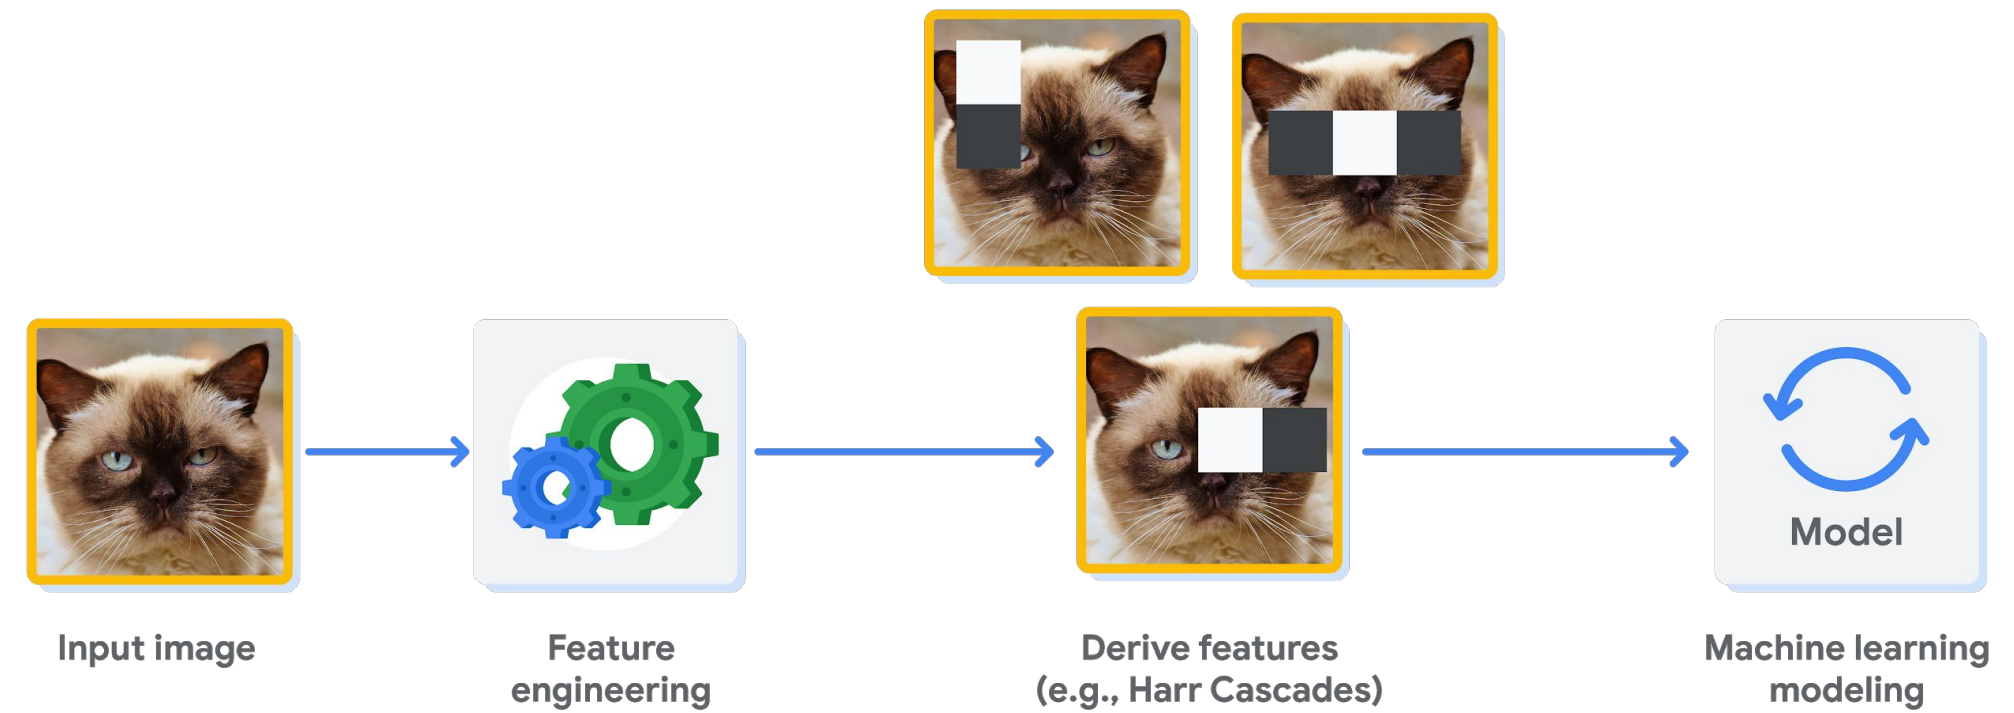


Post-2012

CNN model learns features itself during training




Pre-2012

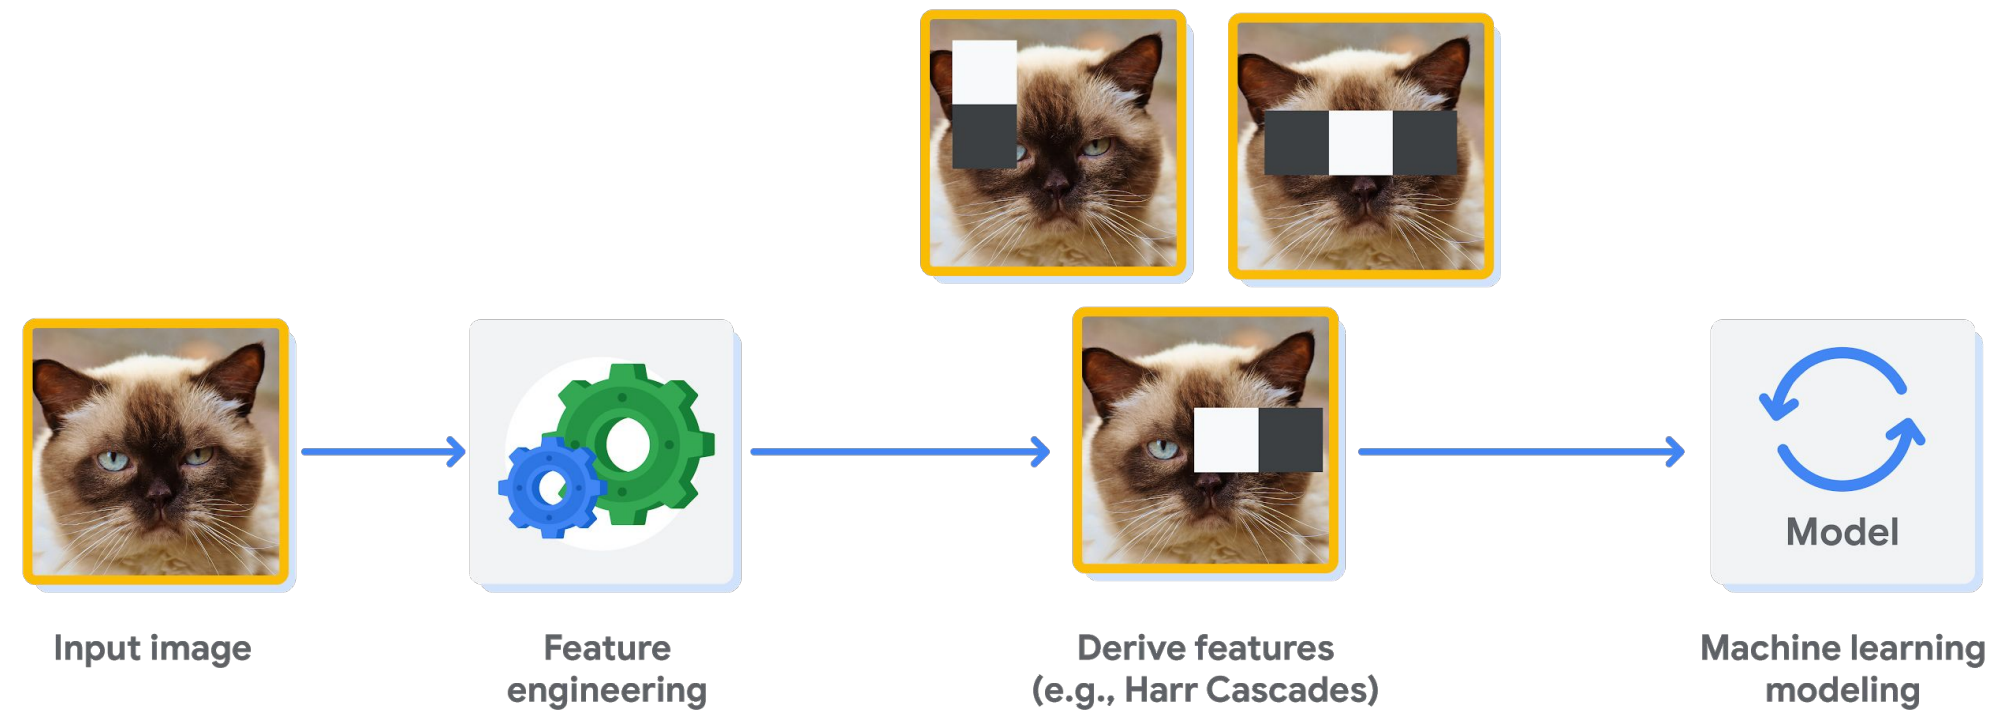


Post-2012

CNN model learns features itself during training

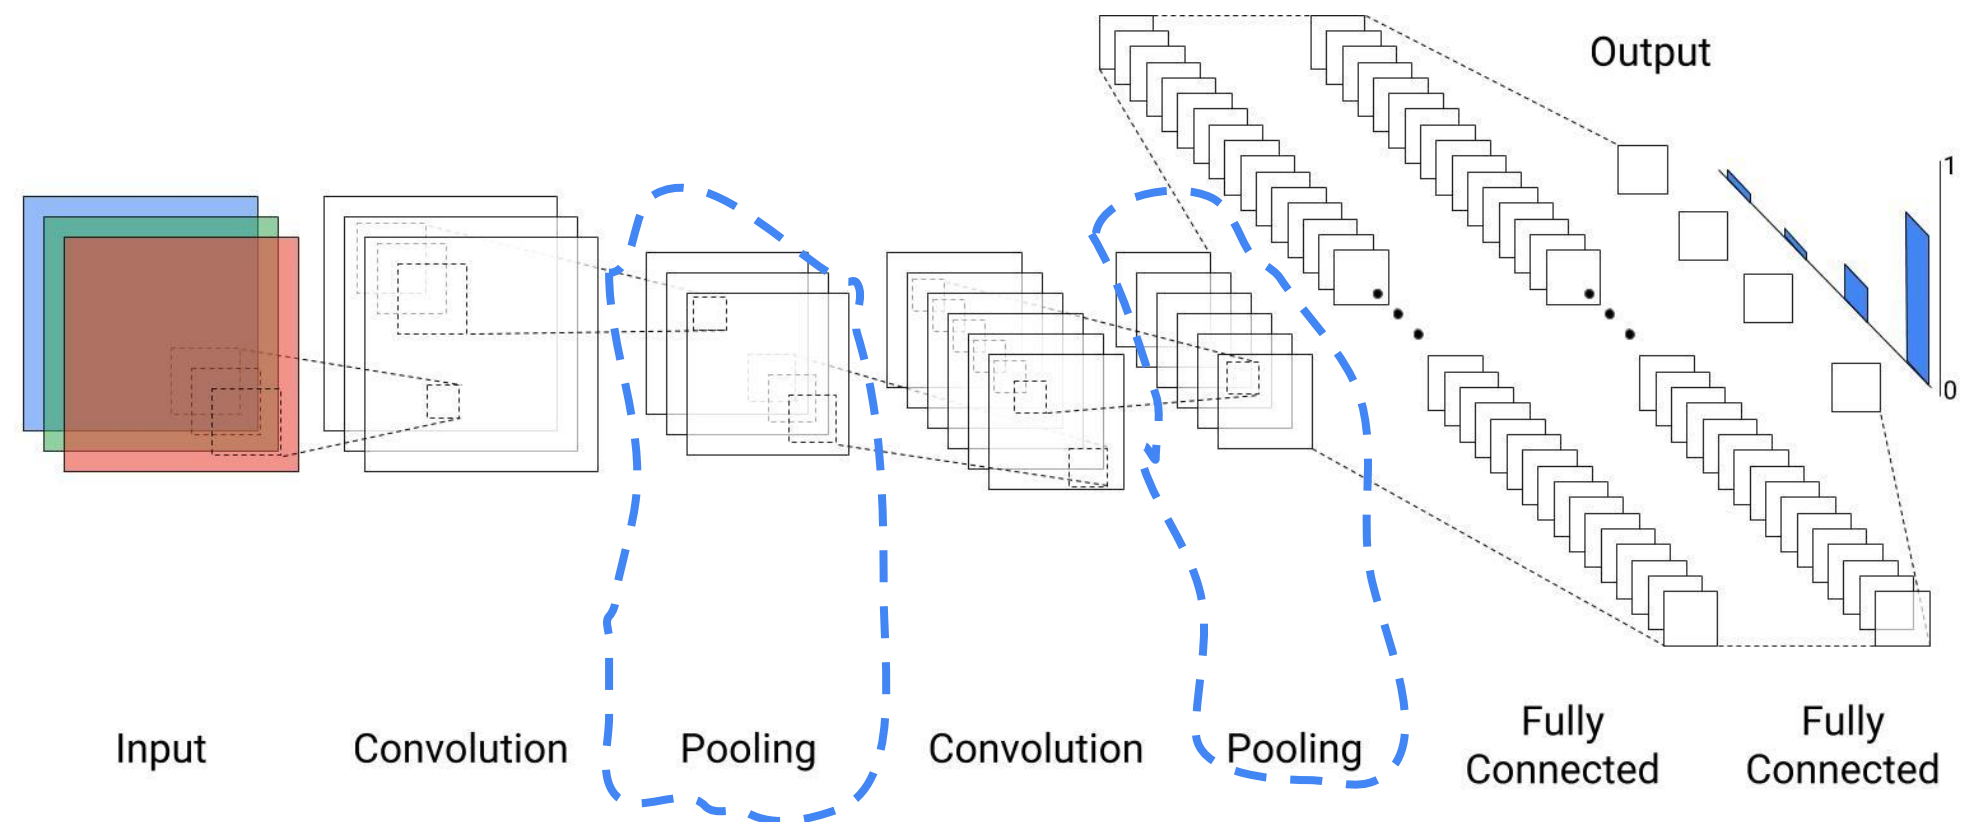


Pre-2012



Post-2012

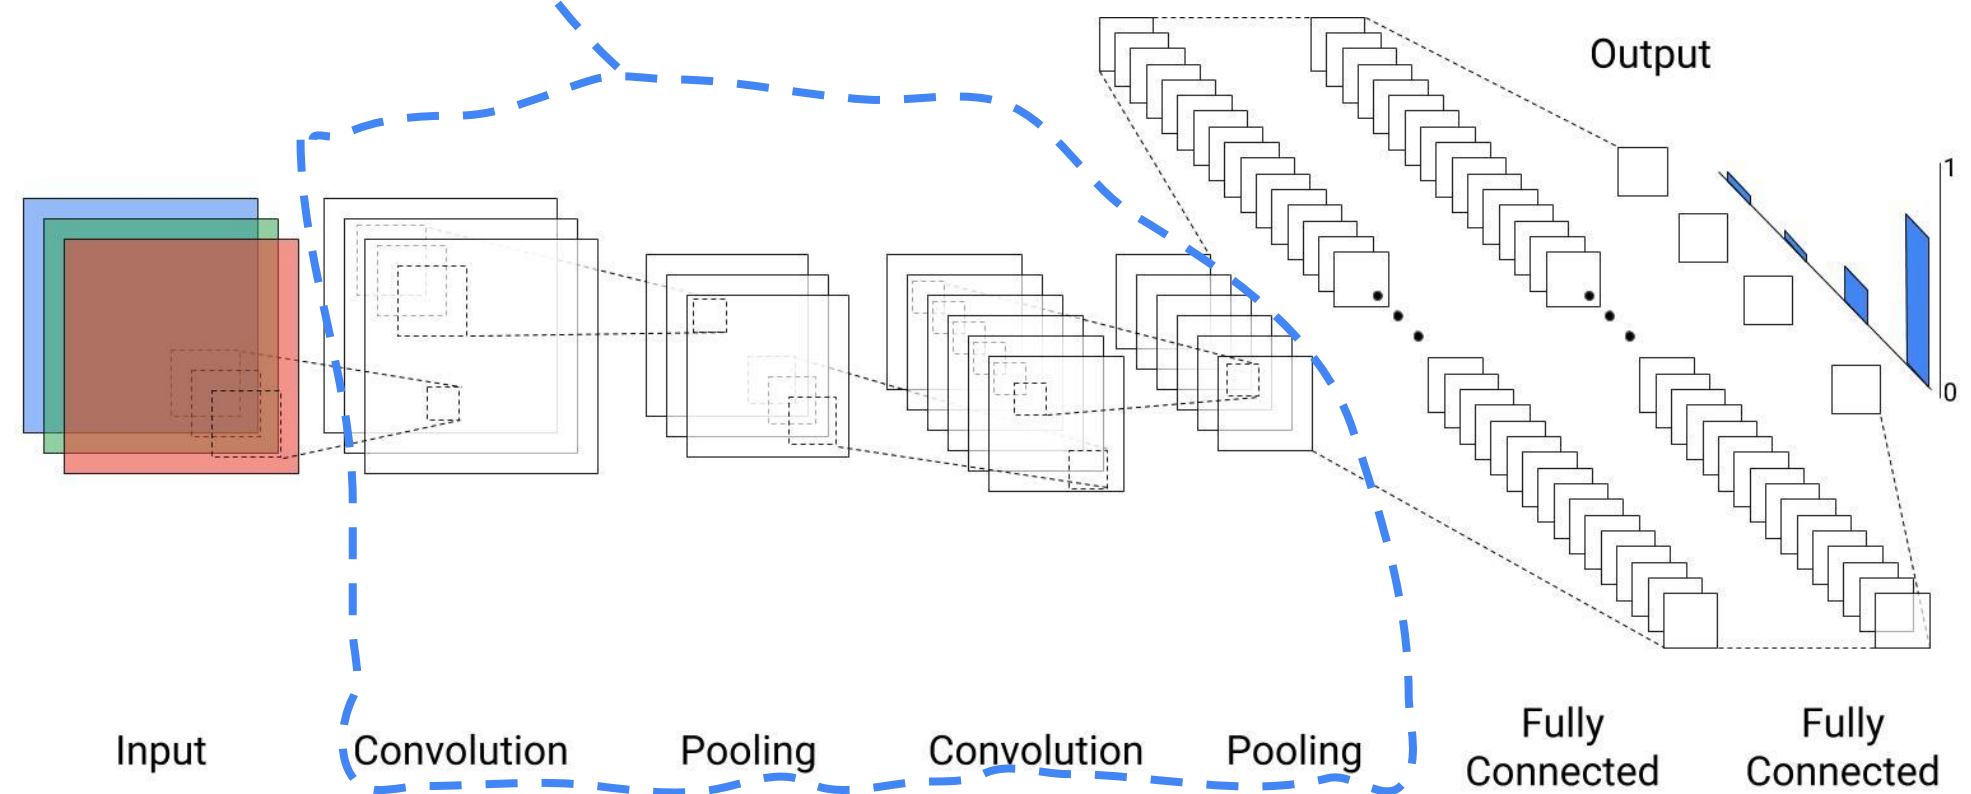
CNN model learns features itself during training



Post-2012

CNN model learns
features itself during
training

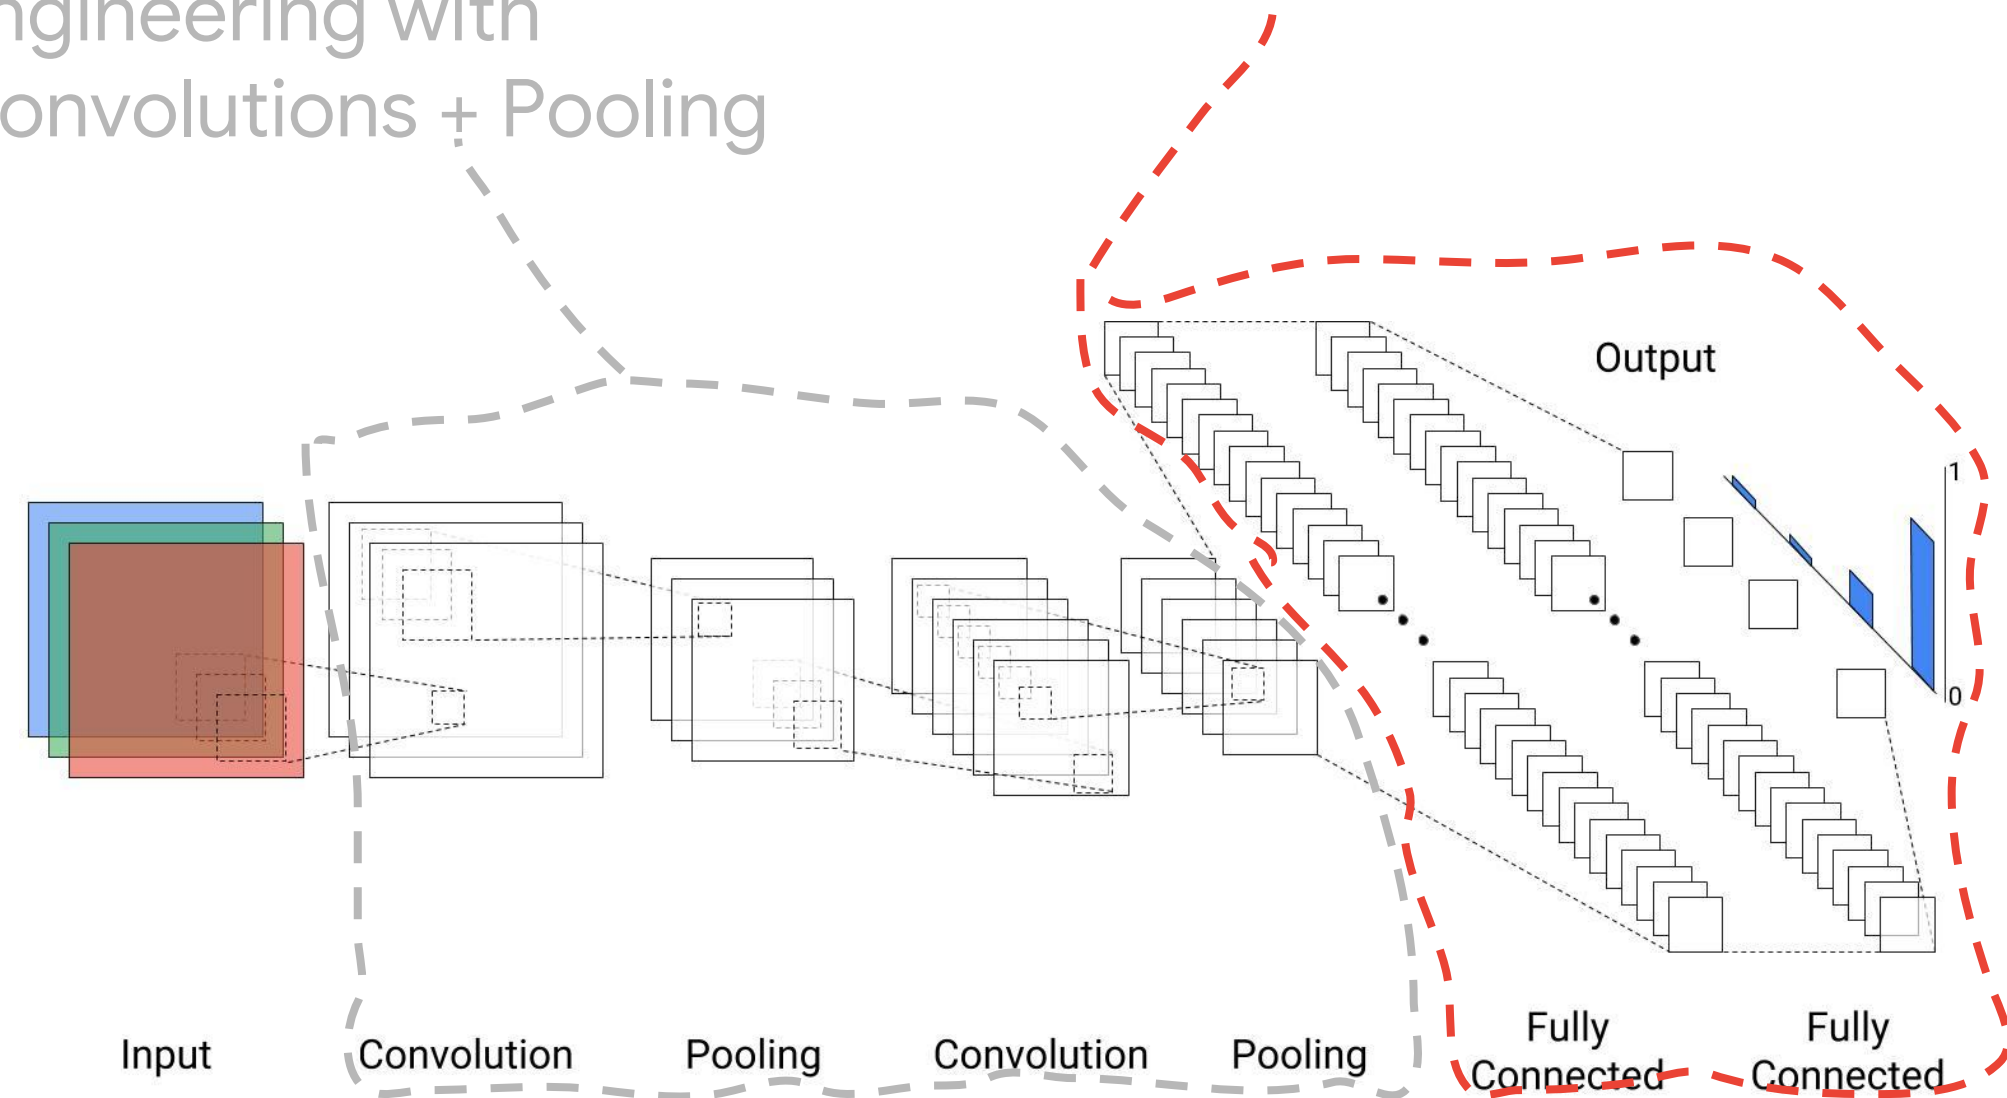
CNN Feature
Engineering with
Convolutions + Pooling



Post-2012
CNN model learns
features itself during
training

CNN Feature
Engineering with
Convolutions + Pooling

Traditional Deep Neural
Network for classification



File Name: T-ICML-O_2_I2_understanding_convolution

Format: Presenter in Studio

Presenter: Carl Osipov

A **convolution** is an operation that processes groups of nearby pixels

A **convolution** is an operation that processes groups of nearby pixels

| | | | | |
|-----------------|-----------------|-----------------|---|---|
| 1 _{x1} | 1 _{x0} | 1 _{x1} | 0 | 0 |
| 0 _{x0} | 1 _{x1} | 1 _{x0} | 1 | 0 |
| 0 _{x1} | 0 _{x0} | 1 _{x1} | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |

Convolved
Feature

| | | | | |
|-----------------|-----------------|-----------------|---|---|
| 1 _{x1} | 1 _{x0} | 1 _{x1} | 0 | 0 |
| 0 _{x0} | 1 _{x1} | 1 _{x0} | 1 | 0 |
| 0 _{x1} | 0 _{x0} | 1 _{x1} | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |

Convolved
Feature

| | | | | |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| | | |

Convolved
Feature

Weights

| | | | | |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| | | |

Convolved
Feature

Weights

Convolutional
Kernel

| | | | | |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| | | |

Convolved
Feature

Shared
Weights

Convolutional
Kernel

| | | | | |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

| | | |
|----|----|----|
| x1 | x0 | x1 |
| x0 | x1 | x0 |
| x1 | x0 | x1 |

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| | | |

Convolved
Feature

| | | | | |
|---|---|-----------------|-----------------|-----------------|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 _{x1} | 1 _{x0} | 0 _{x1} |
| 0 | 0 | 1 _{x0} | 1 _{x1} | 1 _{x0} |
| 0 | 0 | 1 _{x1} | 1 _{x0} | 0 _{x1} |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| | | |

Convolved
Feature

A convolved feature map is created by multiplying the kernel weights by the features

| | | | | |
|---|---|-----------------|-----------------|-----------------|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 _{x1} | 1 _{x0} | 0 _{x1} |
| 0 | 0 | 1 _{x0} | 1 _{x1} | 1 _{x0} |
| 0 | 0 | 1 _{x1} | 1 _{x0} | 0 _{x1} |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| | | |

Convolved
Feature

A convolved feature map is created by multiplying the kernel weights by the features

| | | | | |
|---|---|-----------------|-----------------|-----------------|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 _{x1} | 1 _{x0} | 0 _{x1} |
| 0 | 0 | 1 _{x0} | 1 _{x1} | 1 _{x0} |
| 0 | 0 | 1 _{x1} | 1 _{x0} | 0 _{x1} |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| | | |

Convolved
Feature

A convolved feature map is created by multiplying the kernel weights by the features

| | | | | |
|---|---|-----------------|-----------------|-----------------|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 _{x1} | | 0 _{x1} |
| 0 | 0 | | 1 _{x1} | |
| 0 | 0 | 1 _{x1} | | 0 _{x1} |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| | | |

Convolved
Feature

A convolved feature map is created by multiplying the kernel weights by the features

$$\text{SUM } (1 + 0 + 1 + 1 + 0)$$


| | | | | |
|---|---|-----------------|-----------------|-----------------|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 _{x1} | | 0 _{x1} |
| 0 | 0 | | 1 _{x1} | |
| 0 | 0 | 1 _{x1} | | 0 _{x1} |
| 0 | 1 | 1 | 0 | 0 |

Image



| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| | | |

Convolved
Feature

Convolutional Neural Networks commonly use multiple kernels




Convolutional Neural Networks
commonly use multiple kernels




TPU


Convolutional Neural Networks
commonly use multiple kernels




Detects
horizontal edges



Convolutional Neural Networks
commonly use multiple kernels




Detects
horizontal edges



Detects bright
spots


Different weights
detect different features



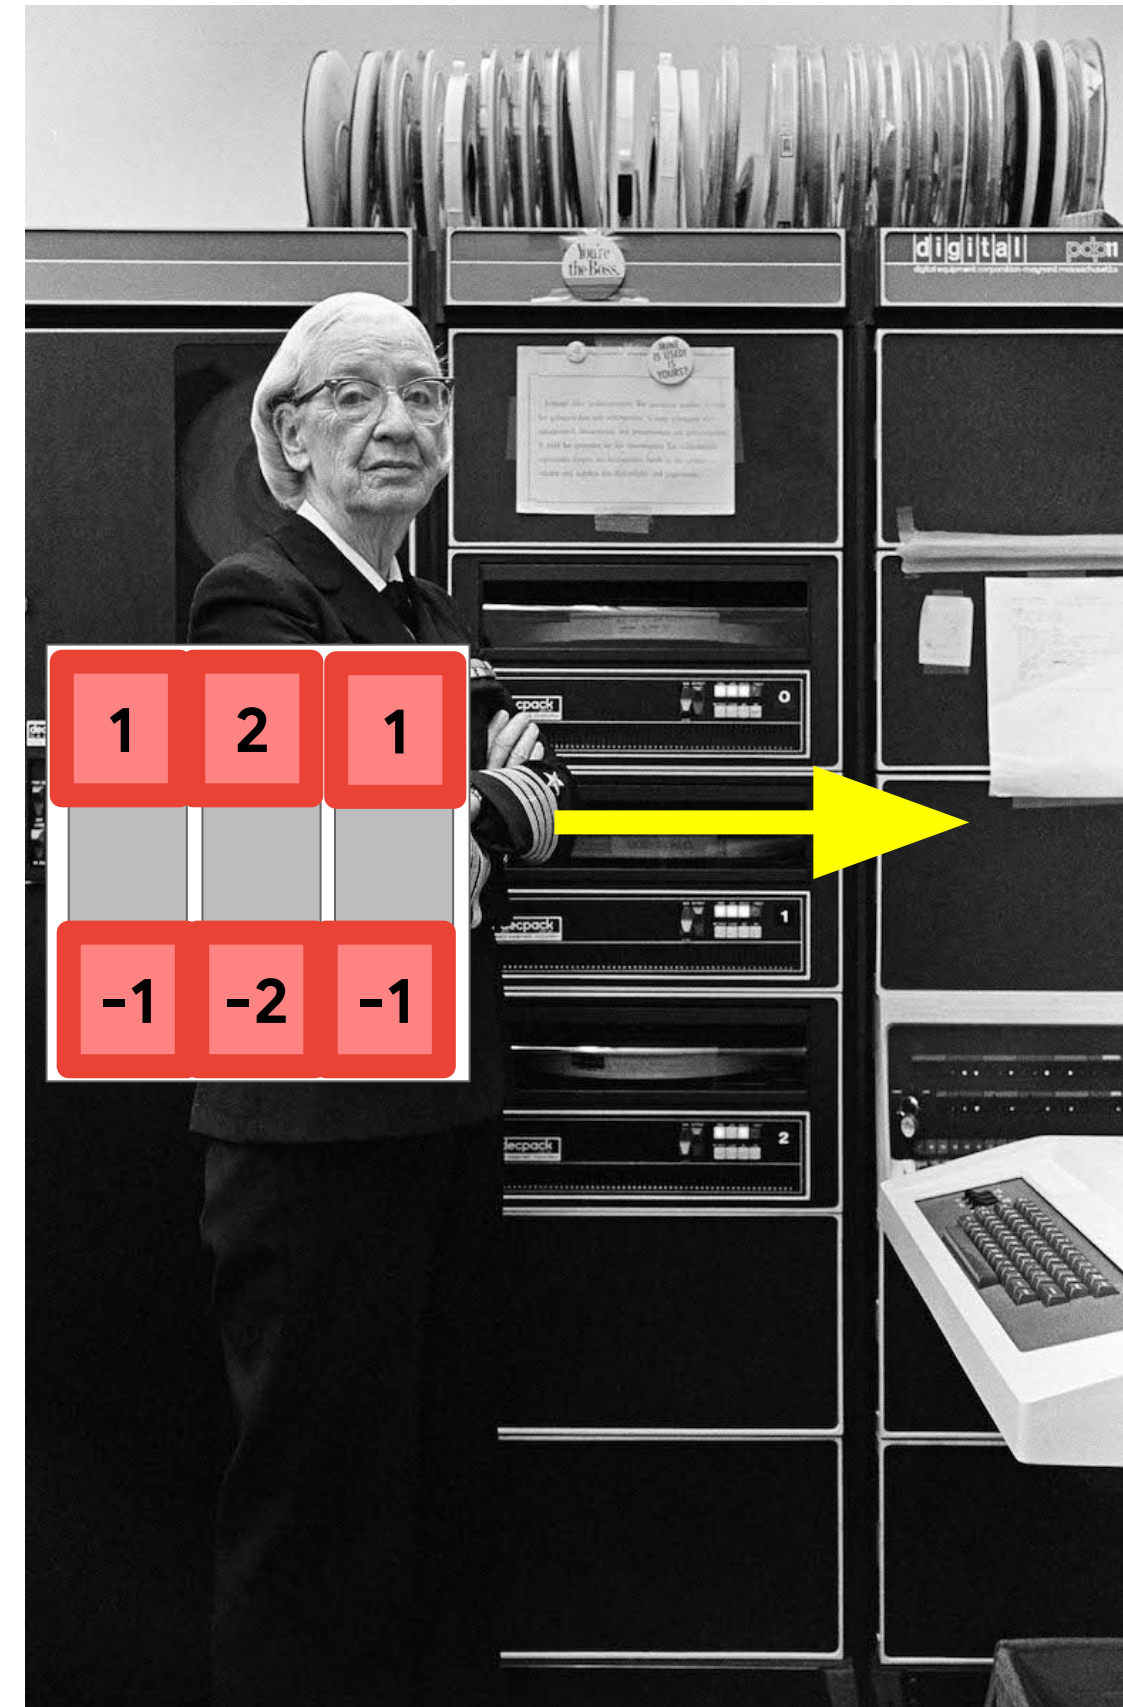
Detects
horizontal edges






Different weights
detect different features



Detects
horizontal edges





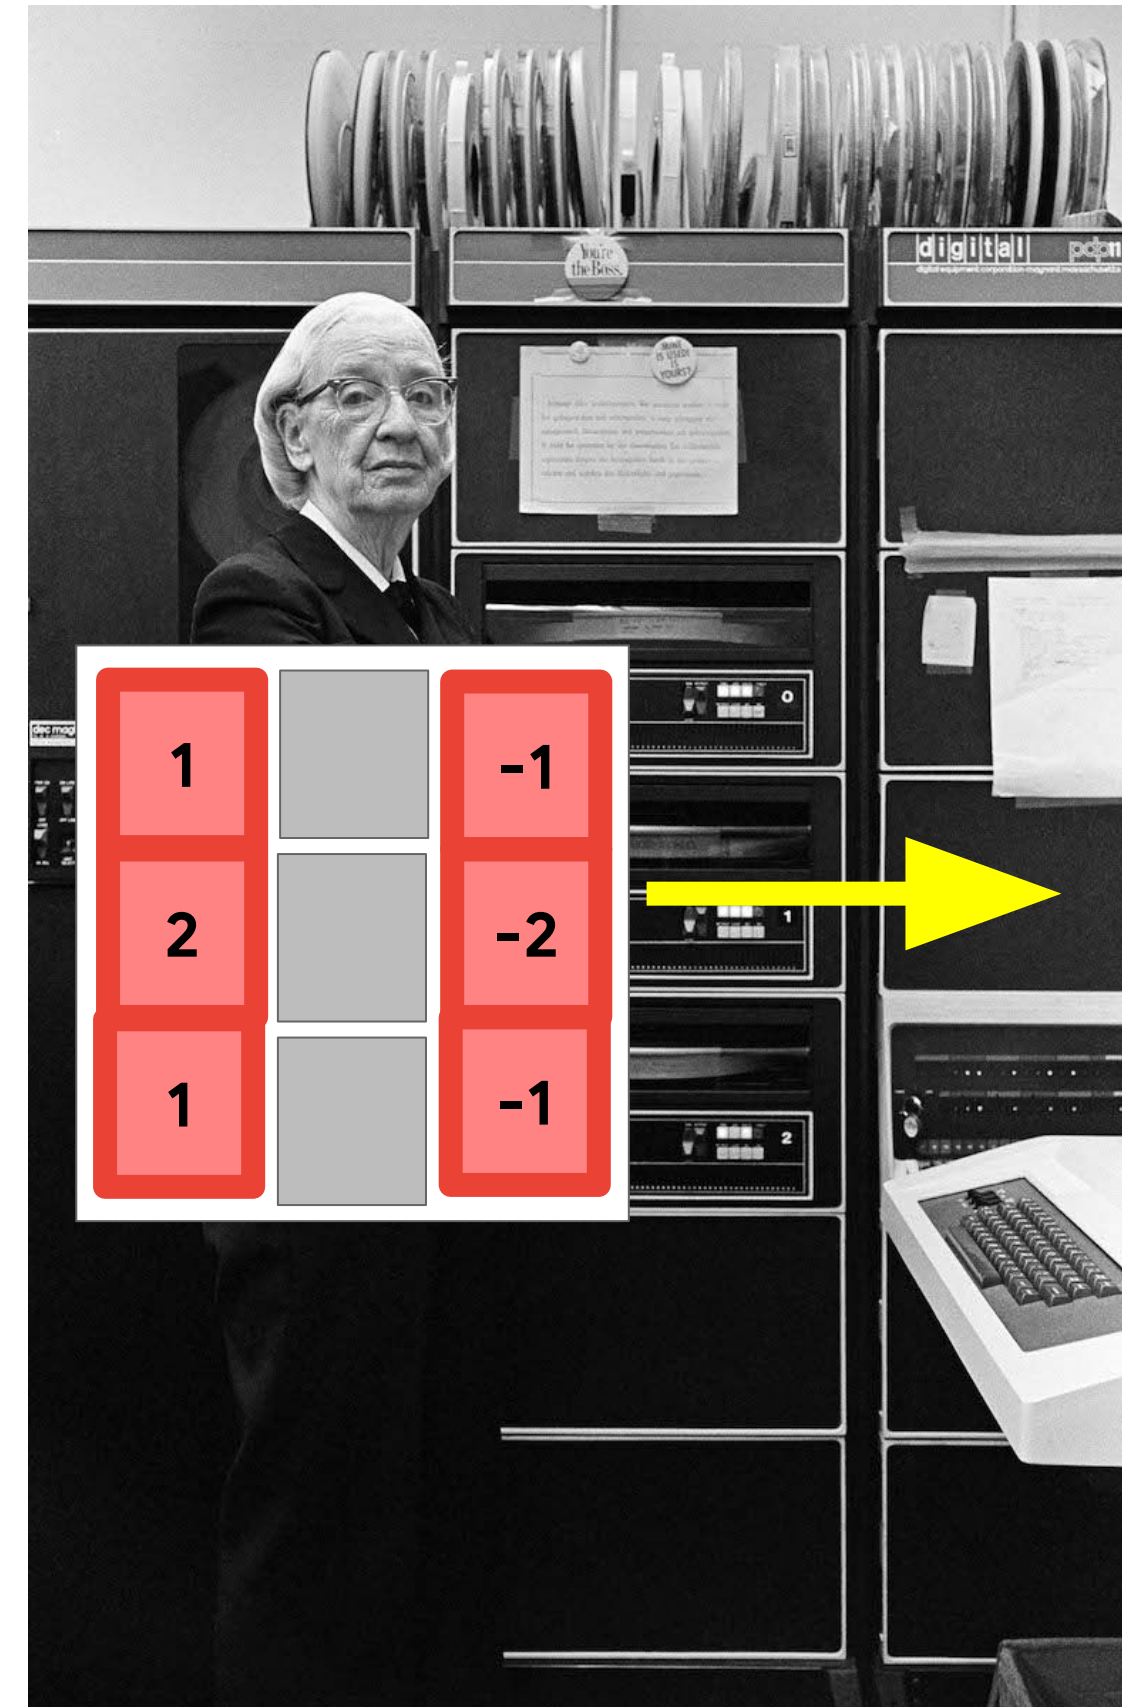
Different weights
detect different features


| | | |
|----|----|----|
| 1 | 2 | 1 |
| | | |
| -1 | -2 | -1 |

Detects
horizontal edges

| | | |
|---|--|----|
| 1 | | -1 |
| 2 | | -2 |
| 1 | | -1 |


Detects
vertical edges





| | | |
|----|----|----|
| 1 | 2 | 1 |
| | | |
| -1 | -2 | -1 |

| | | |
|---|--|----|
| 1 | | -1 |
| 2 | | -2 |
| 1 | | -1 |





Kernel 1

| | | |
|----|----|----|
| 1 | 2 | 1 |
| | | |
| -1 | -2 | -1 |

Kernel 2

| | | |
|---|--|----|
| 1 | | -1 |
| 2 | | -2 |
| 1 | | -1 |

Sum of the two
filter outputs



=





Kernel 1

| | | |
|----|----|----|
| 1 | 2 | 1 |
| | | |
| -1 | -2 | -1 |

Kernel 2

| | | |
|---|--|----|
| 1 | | -1 |
| 2 | | -2 |
| 1 | | -1 |

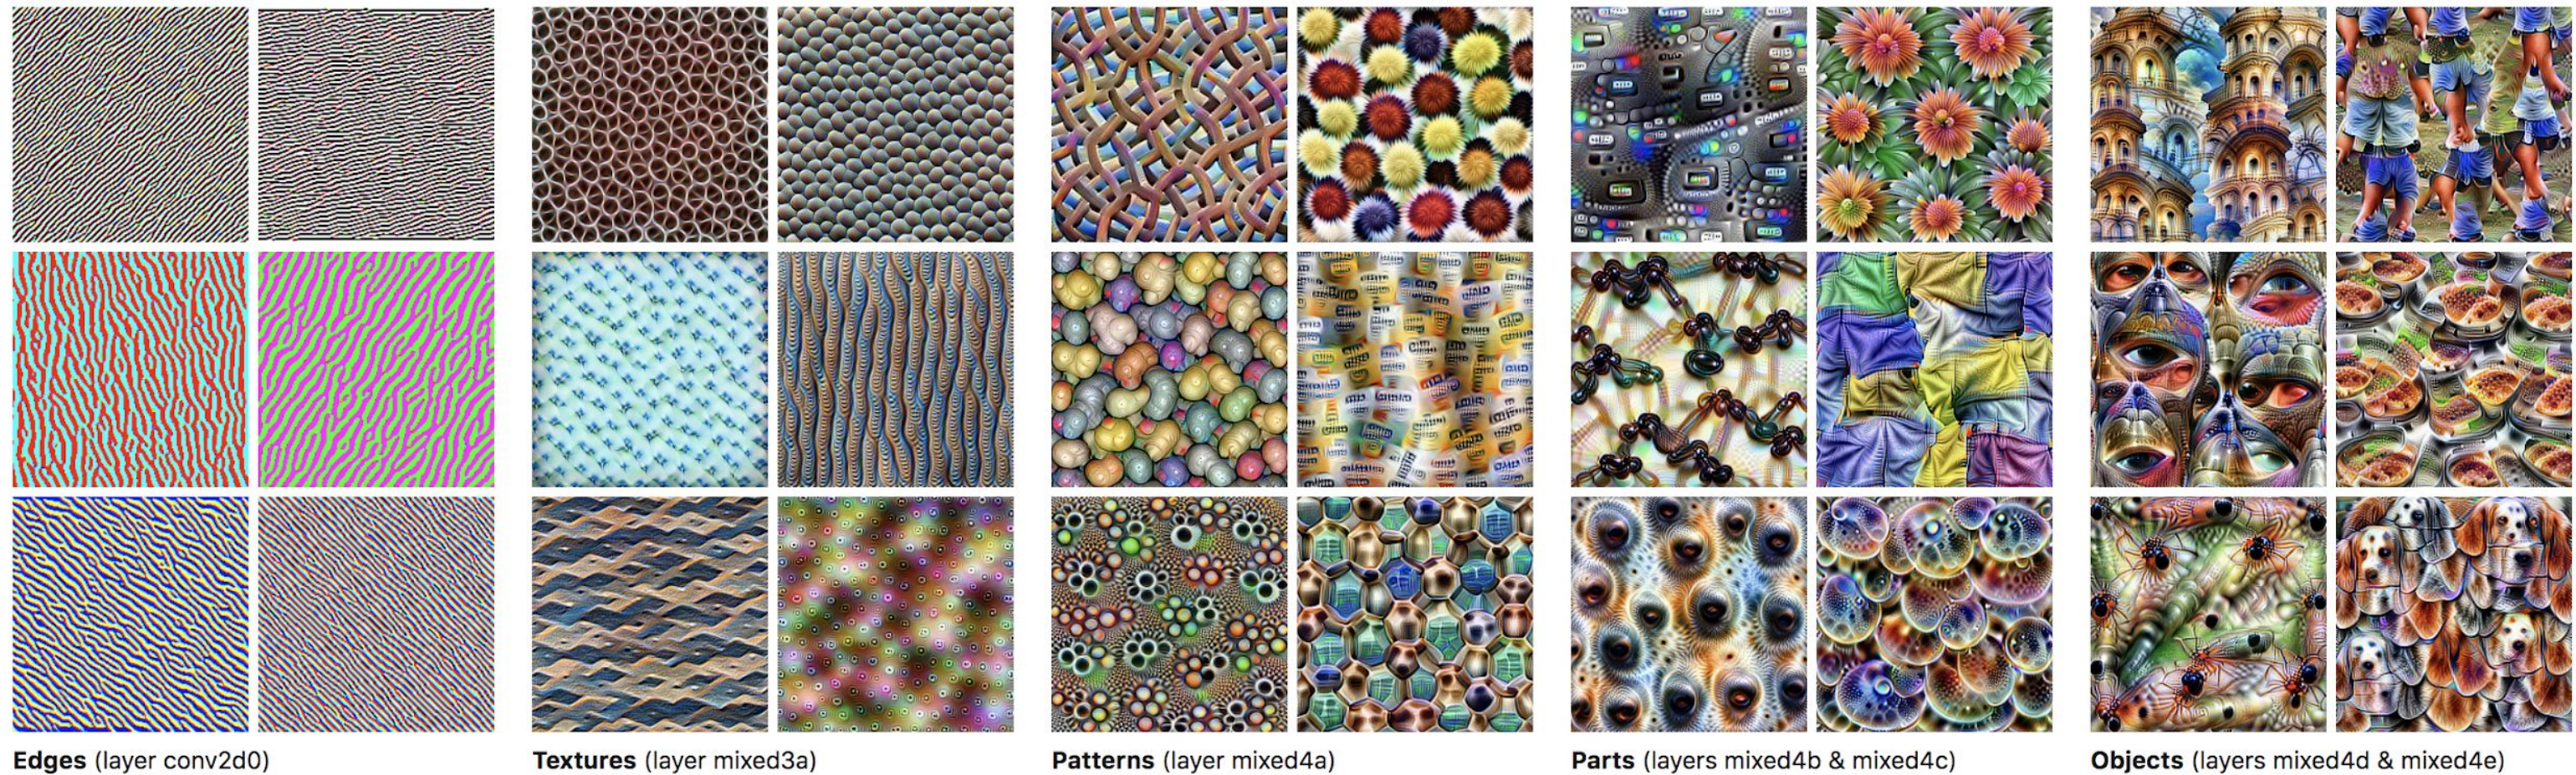
Sum of the two
filter outputs



=



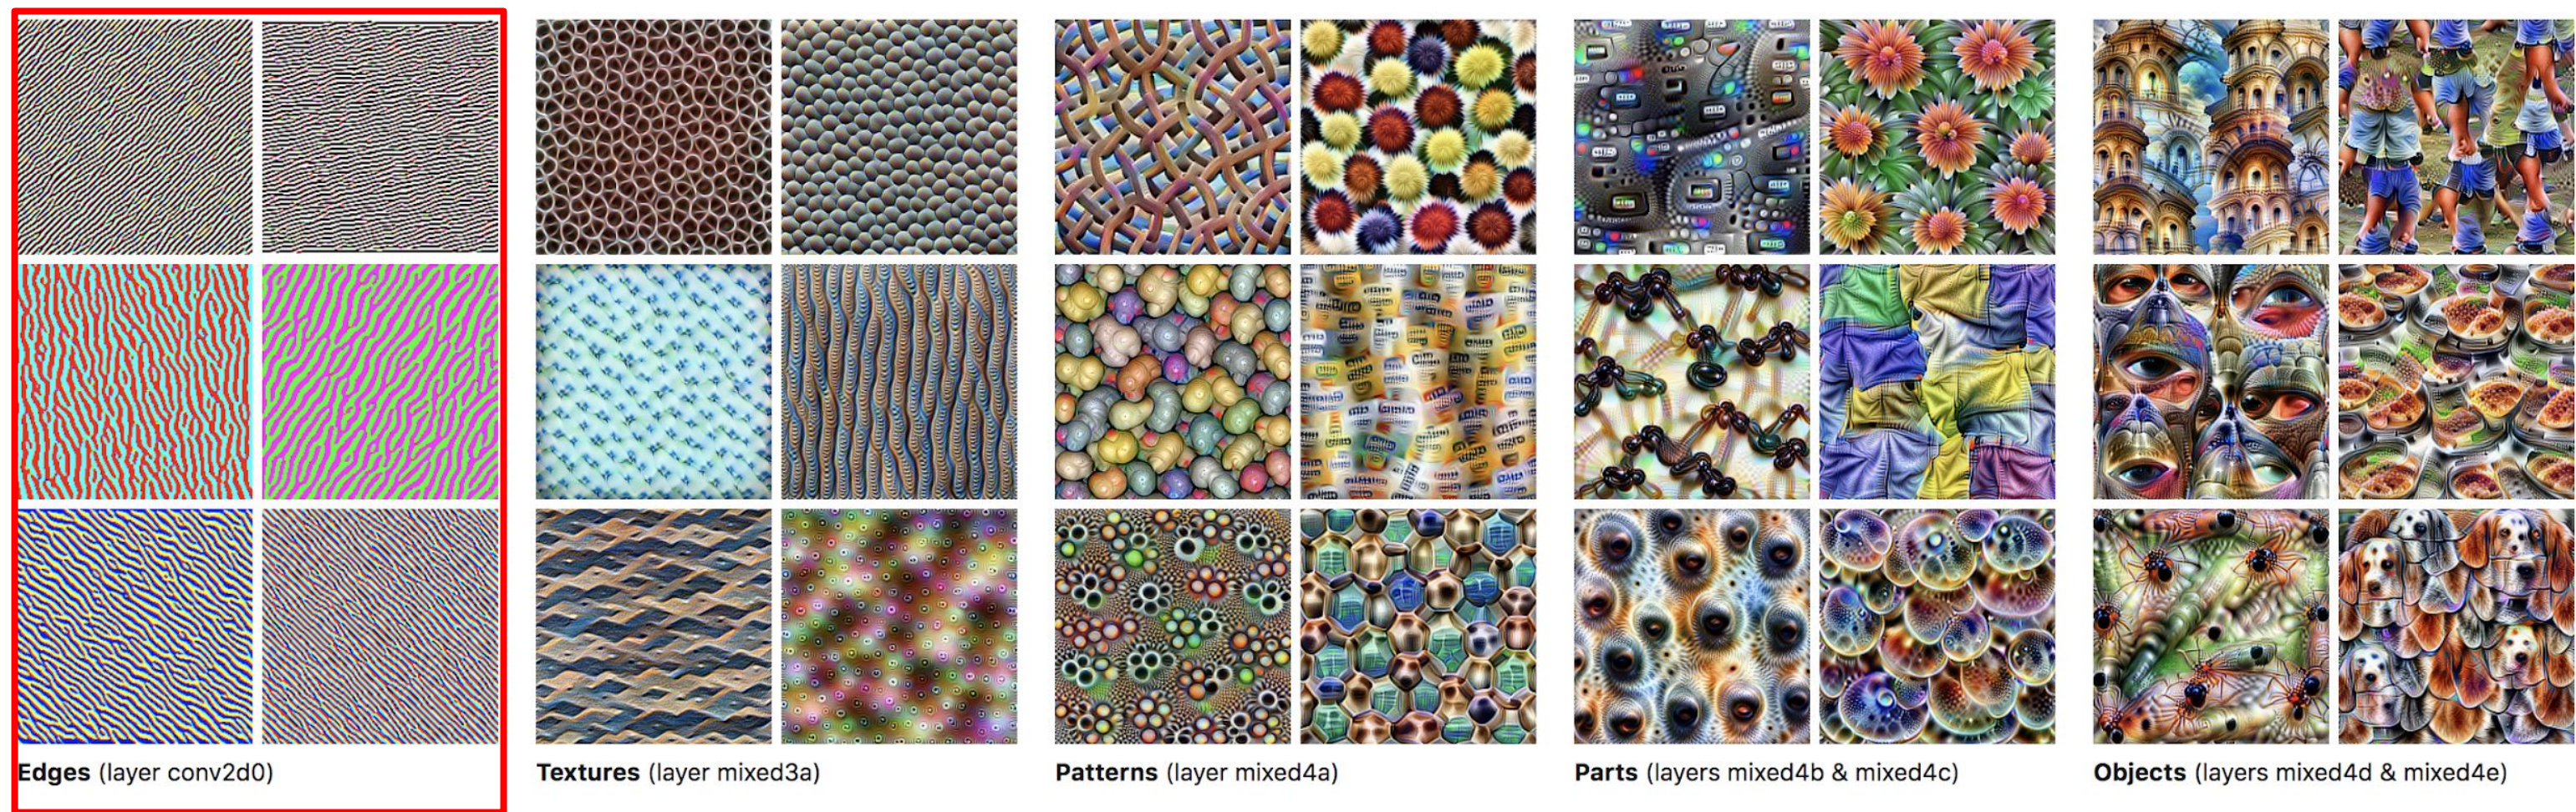
CNNs can learn a hierarchy of features



Flow of data in the network

<https://distill.pub/2017/feature-visualization/>


CNNs can learn a hierarchy of features



Flow of data in the network

<https://distill.pub/2017/feature-visualization/>

CNNs can learn a hierarchy of features



Flow of data in the network

<https://distill.pub/2017/feature-visualization/>

Quiz:

What is the process of "sliding" a kernel across an image called?

1. Convolution
2. Confusion matrix
3. Contortion
4. Curved detection

Quiz:

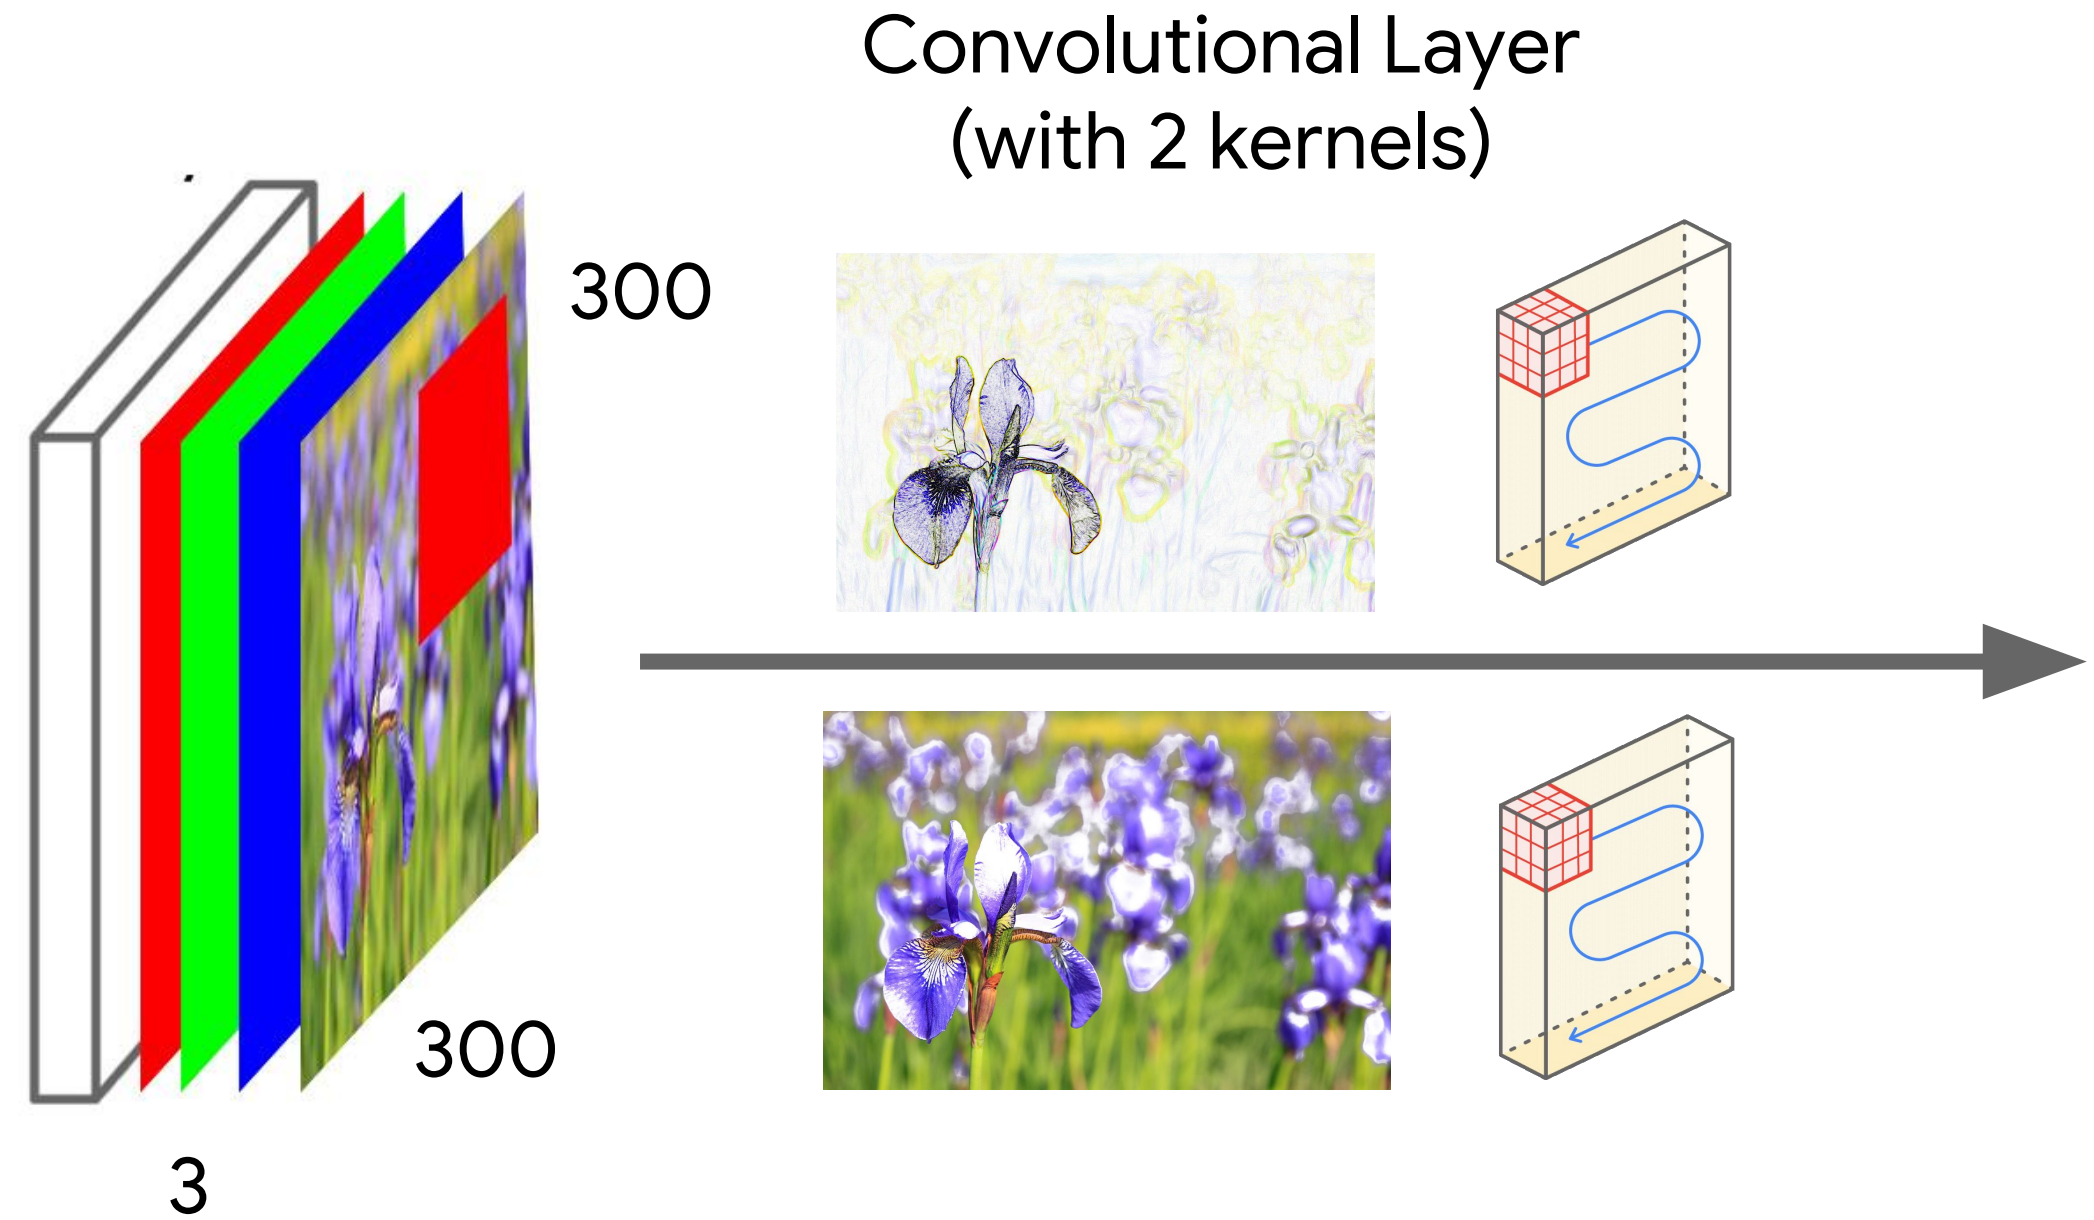
What is the process of "sliding" a kernel across an image called?

1. **Convolution**
2. Confusion matrix
3. Contortion
4. Curved detection

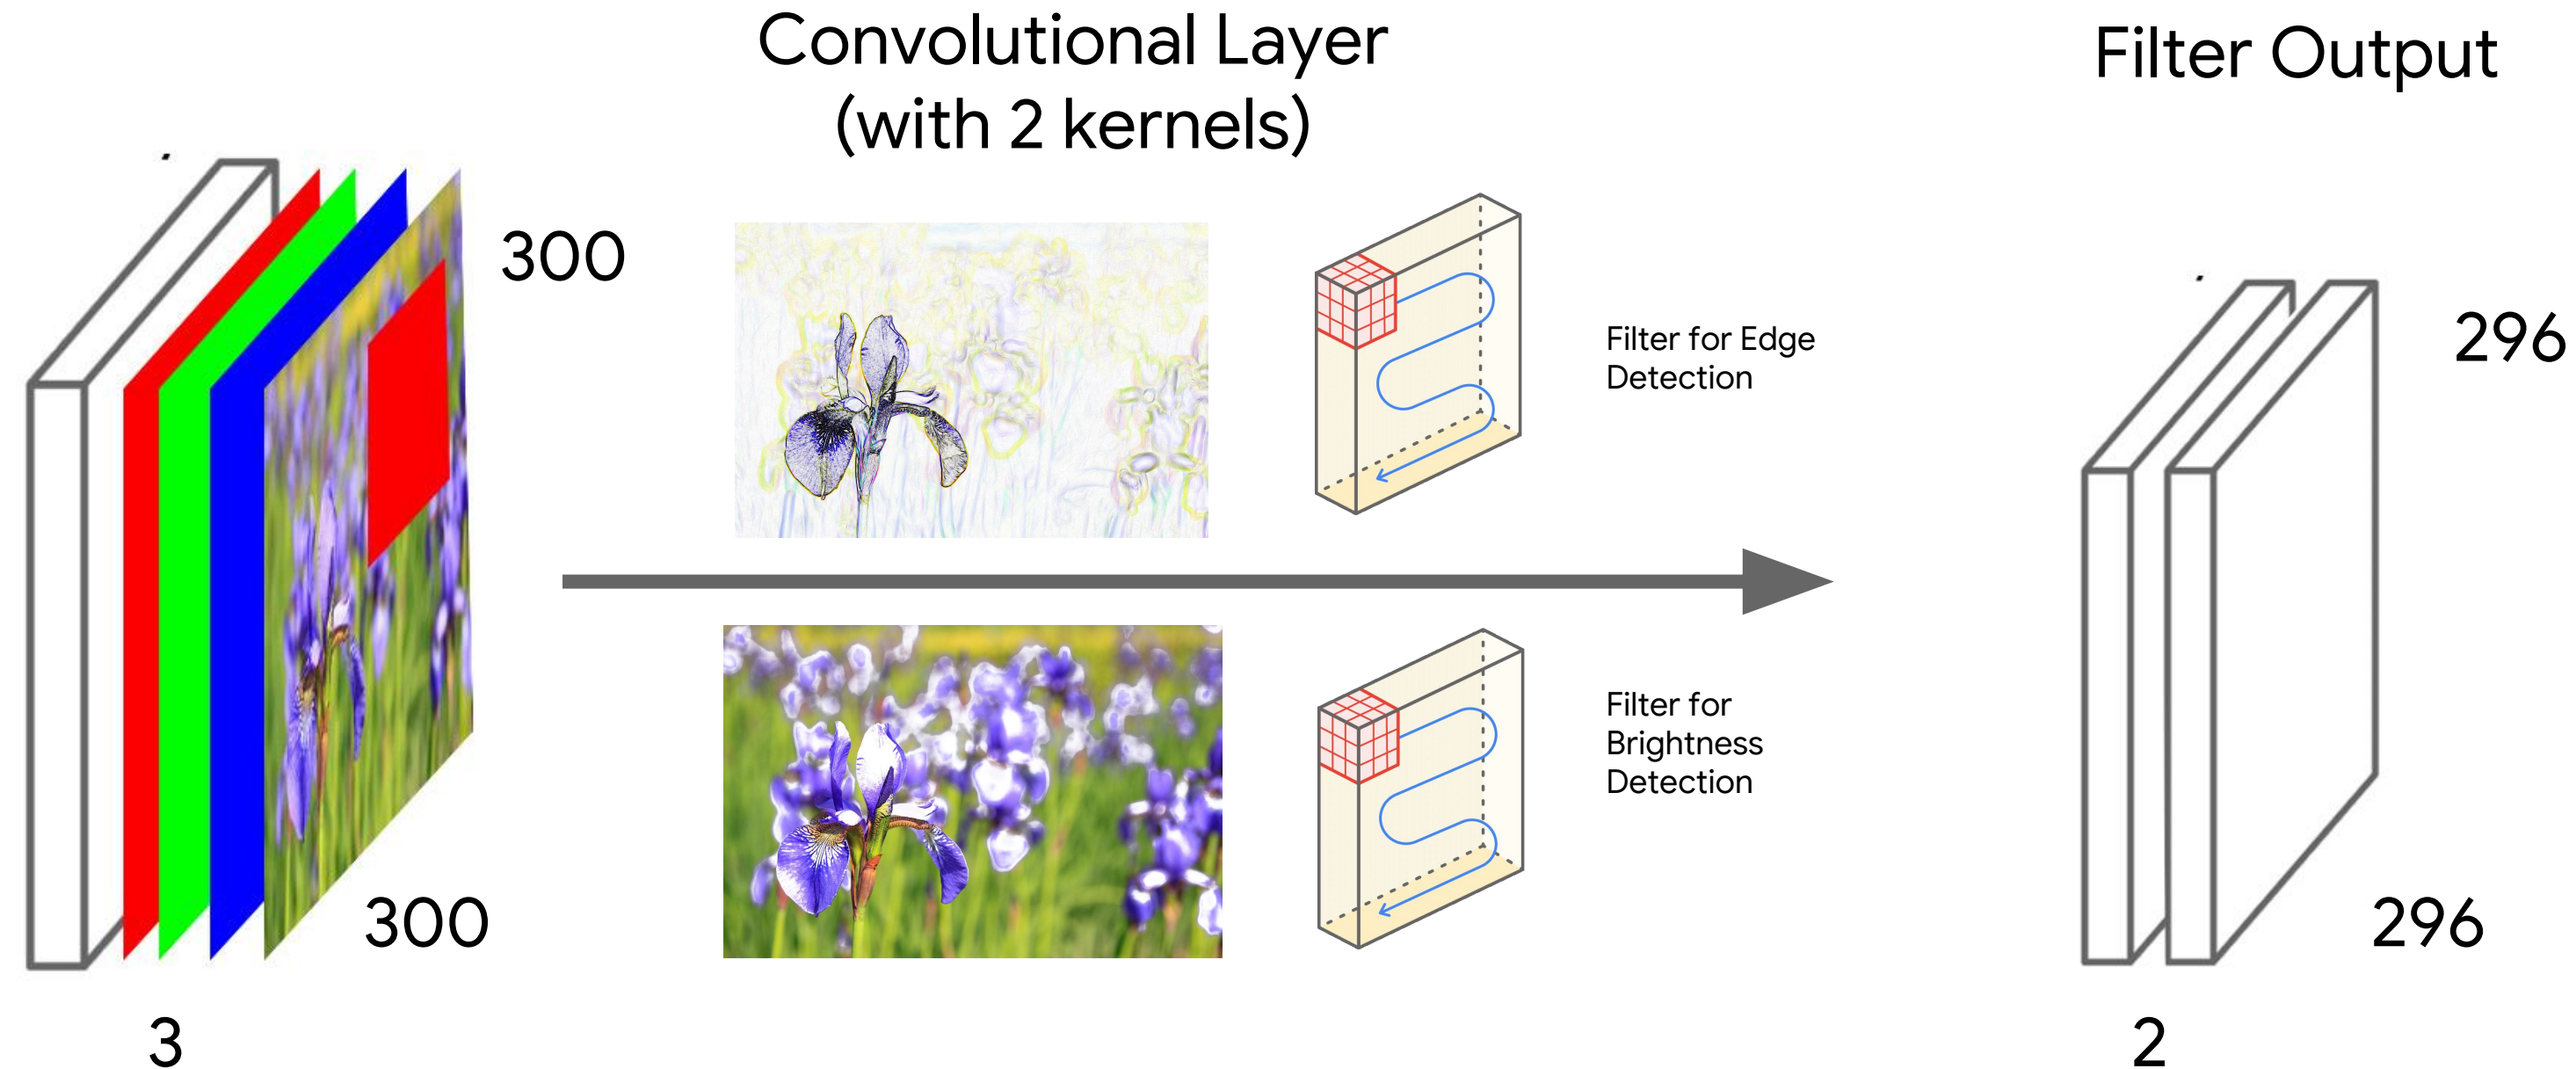
Convolutional Layers are collections of filters



Convolutional Layers are collections of filters



Convolutional Layers are collections of filters

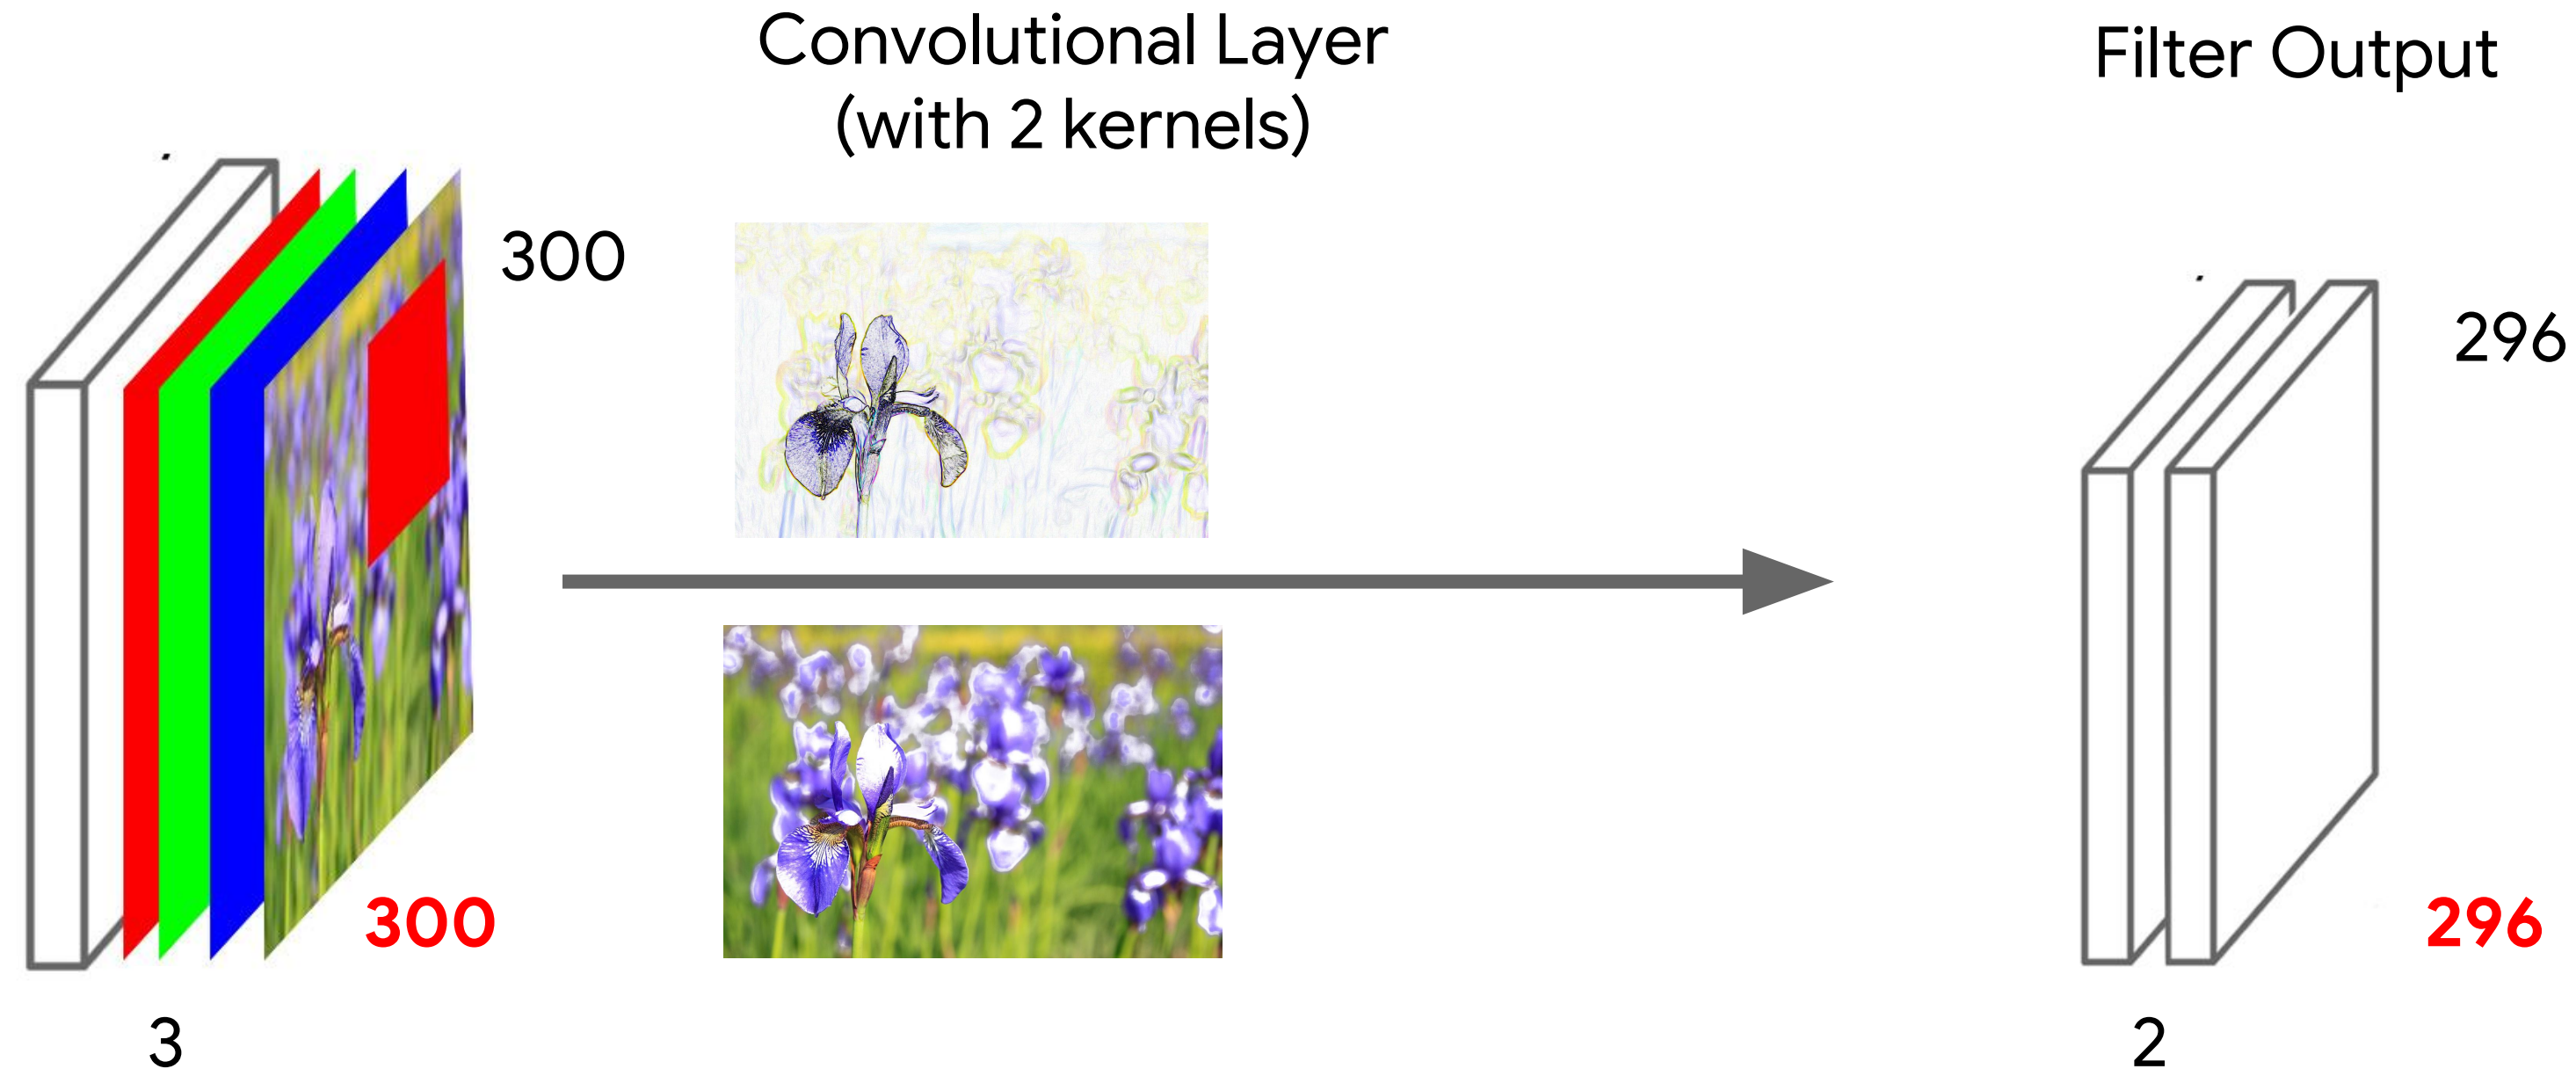


File Name: T-ICML-O_2_I3_cnn_parameters

Format: Presenter in Studio

Presenter: Carl Osipov

Convolutional Layers are collections of filters



Padding preserves the shape of the input after the convolution

| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 | 0 | 0 |
| 0 | 0 | 1 | 1 | 1 | 0 | 0 |
| 0 | 0 | 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 0 | 1 | 1 | 0 | 0 |
| 0 | 0 | 1 | 1 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

"Same" padding for 3x3 kernel size

TensorFlow provides a high level API to set up a convolutional layer

```
tf.layers.conv2d(  
    inputs=...,          # [batch_size, filter_height, filter_width, in_channels]  
    filters=...,         # number of filters, i.e. out_channels  
    kernel_size=3,       # size of the kernel, e.g. 3 for a 3x3 kernel  
    padding='same'       # maintain the same shape across the input and output  
    ...  
)
```

TensorFlow provides a high level API to set up a convolutional layer

```
tf.layers.conv2d(  
    inputs=...,          # [batch_size, filter_height, filter_width, in_channels]  
    filters=...,         # number of filters, i.e. out_channels  
    kernel_size=3,       # size of the kernel, e.g. 3 for a 3x3 kernel  
    padding='same'       # maintain the same shape across the input and output  
    ...  
)
```

TensorFlow provides a high level API to set up a convolutional layer

```
tf.layers.conv2d(  
    inputs=...,          # [batch_size, filter_height, filter_width, in_channels]  
    filters=...,         # number of filters, i.e. out_channels  
    kernel_size=3,       # size of the kernel, e.g. 3 for a 3x3 kernel  
    padding='same'       # maintain the same shape across the input and output  
    ...  
)
```


TensorFlow provides a high level API to set up a convolutional layer

```
tf.layers.conv2d(  
    inputs=...,          # [batch_size, filter_height, filter_width, in_channels]  
    filters=...,         # number of filters, i.e. out_channels  
    kernel_size=3,       # size of the kernel, e.g. 3 for a 3x3 kernel  
    padding='same'       # maintain the same shape across the input and output  
    ...  
)
```


TensorFlow provides a high level API to set up a convolutional layer

```
tf.layers.conv2d(  
    inputs=...,          # [batch_size, filter_height, filter_width, in_channels]  
    filters=...,         # number of filters, i.e. out_channels  
    kernel_size=3,       # size of the kernel, e.g. 3 for a 3x3 kernel  
    padding='same'       # maintain the same shape across the input and output  
    ...  
)
```

TensorFlow provides a high level API to set up a convolutional layer



```
tf.layers.conv2d(  
    inputs=...,          # [batch_size, filter_height, filter_width, in_channels]  
    filters=...,         # number of filters, i.e. out_channels  
    kernel_size=3,       # size of the kernel, e.g. 3 for a 3x3 kernel  
    padding='same'       # maintain the same shape across the input and output  
    ...  
)
```

Dimensionality reduction decreases the total # of parameters and computation time





Dimensionality reduction decreases the total # of parameters and computation time

stride=2





Dimensionality reduction decreases the total # of parameters and computation time

stride=2




stride=3



Quiz:

What does the green part of the below diagram represent?


1. The kernel
2. The filter output
3. The original image



Quiz:

What does the green part of the below diagram represent?


1. The kernel
2. The filter output
- 3. The original image**



Quiz:

What does the yellow part of the below diagram represent?


1. The kernel convolving
2. The filter output
3. The original image



Quiz:

What does the yellow part of the below diagram represent?

1. The kernel convolving
2. The filter output
3. The original image



Quiz:

What is the size of the stride step in the below diagram?

1. Stride of 3
2. Stride of 2
3. Stride of 1
4. Unknown

| | | | | |
|------------------------|------------------------|------------------------|---|---|
| 1 <small>x1</small> | 1 <small>x0</small> | 1 <small>x1</small> | 0 | 0 |
| 0 <small>x0</small> | 1 <small>x1</small> | 1 <small>x0</small> | 1 | 0 |
| 0 <small>x1</small> | 0 <small>x0</small> | 1 <small>x1</small> | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |

Convolved
Feature

Quiz:

What is the size of the stride step in the below diagram?

1. Stride of 3
2. Stride of 2
- 3. Stride of 1**
4. Unknown

| | | | | |
|------------------------|------------------------|------------------------|---|---|
| 1 <small>x1</small> | 1 <small>x0</small> | 1 <small>x1</small> | 0 | 0 |
| 0 <small>x0</small> | 1 <small>x1</small> | 1 <small>x0</small> | 1 | 0 |
| 0 <small>x1</small> | 0 <small>x0</small> | 1 <small>x1</small> | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |


Convolved
Feature

File Name: T-ICML-O_2_I4_pooling_layers

Format: Presenter in Studio


Presenter: Carl Osipov

Maxpooling also reduces dimensionality



Rectified Feature Map

Maxpooling example



Brighter values correspond to larger pixel values



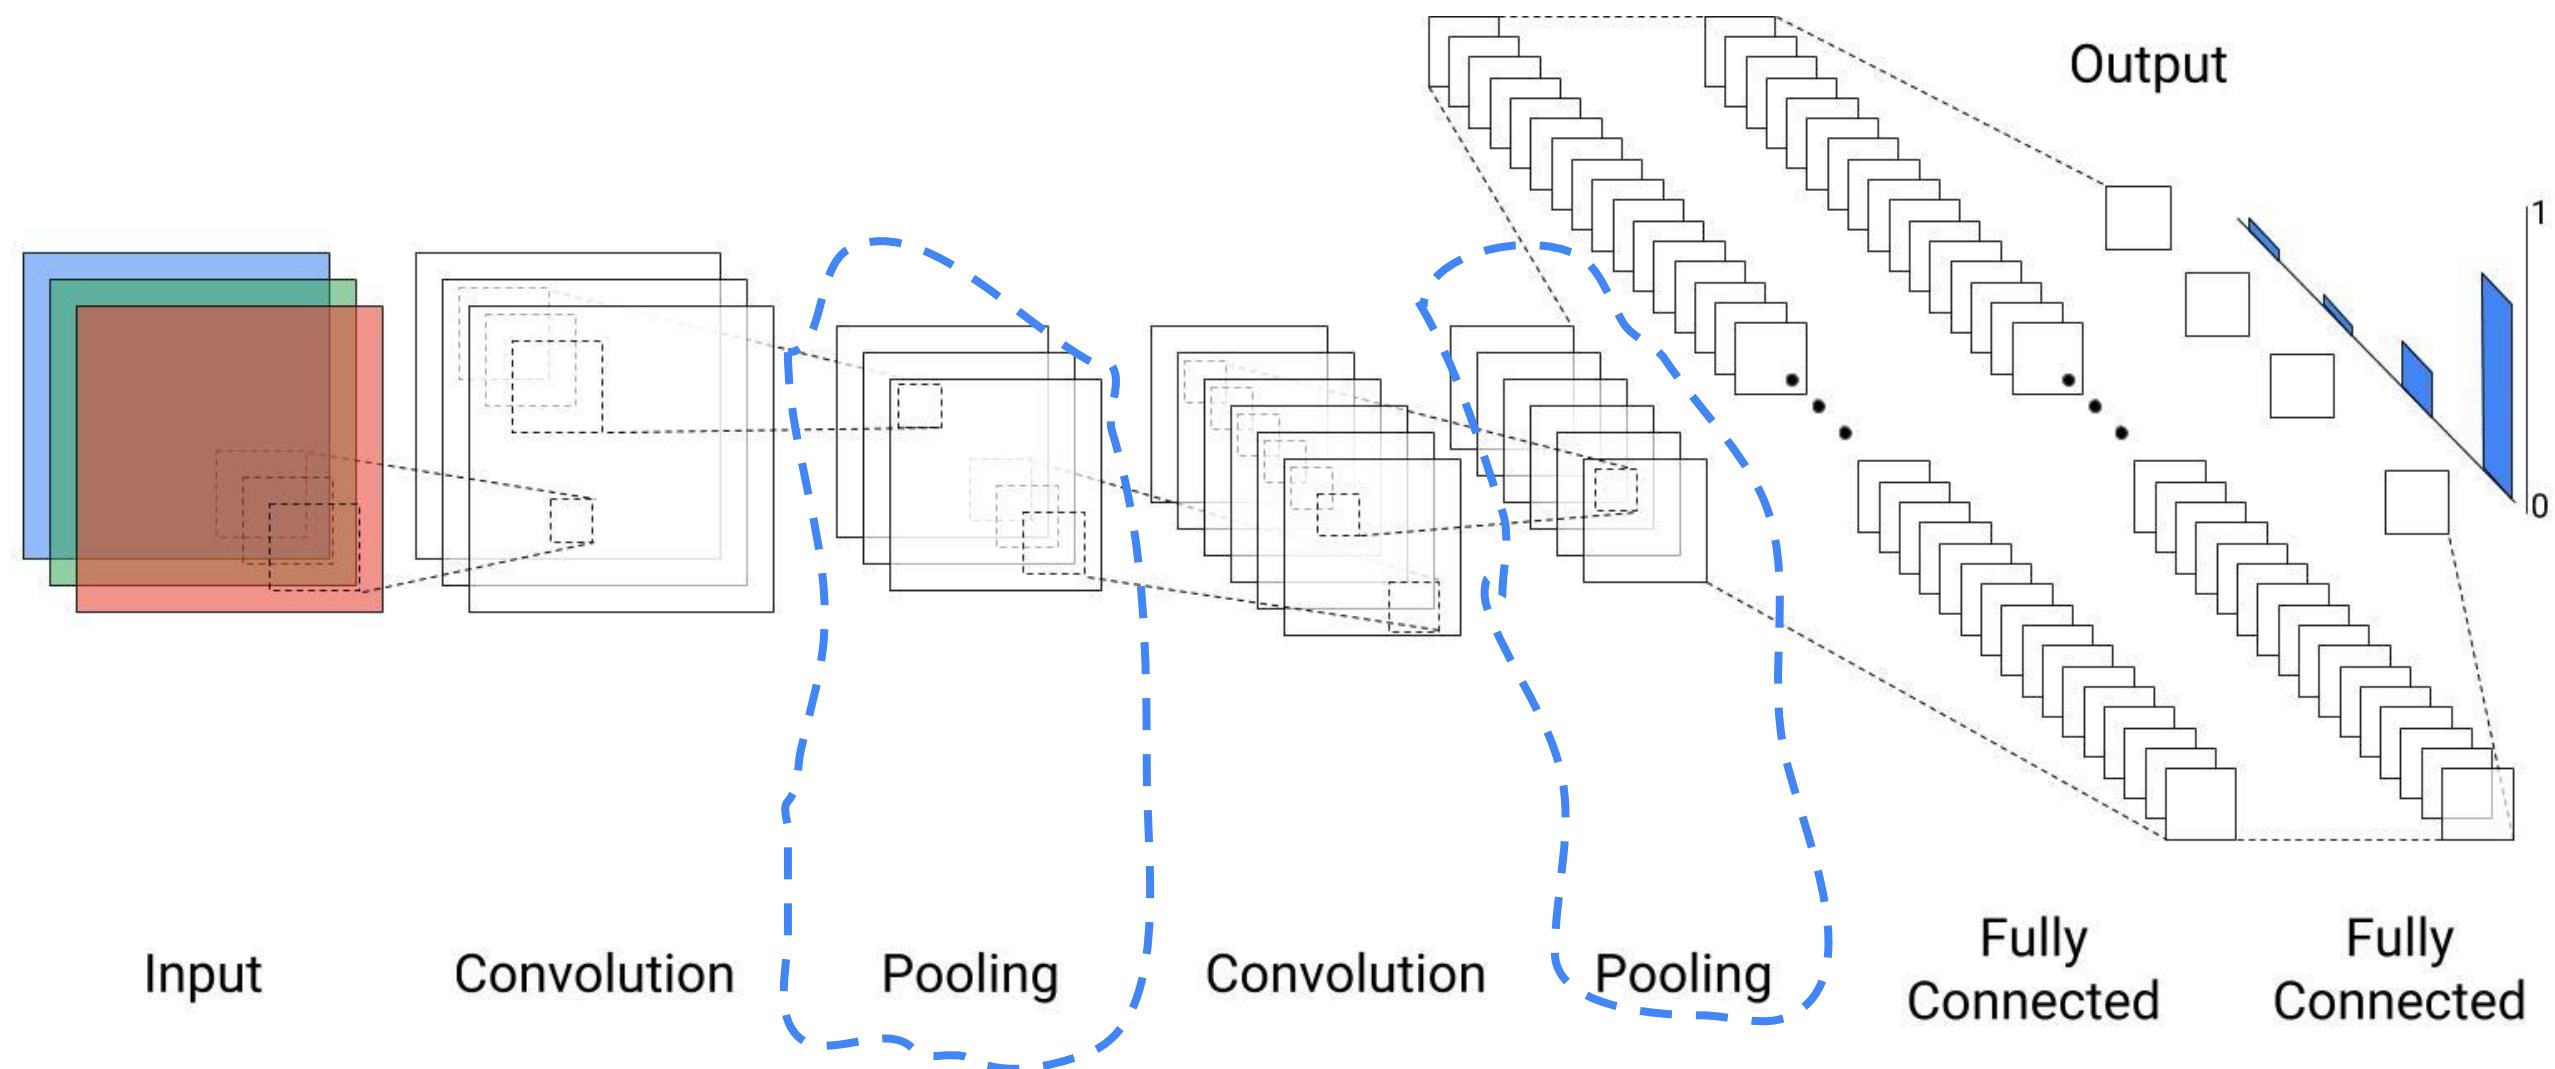
Maxpool

Kernel: 2x2
Stride: 2



Dimensionality is decreased by a factor of 2 both horizontally and vertically

Maxpooling operations are just additional layers in our network



TensorFlow provides the tf.layers.max_pooling2d API

```
tf.layers.max_pooling2d(  
    inputs,          # [batch_size, filter_height, filter_width, in_channels]  
    pool_size=2,     # size of the maxpooling kernel  
    strides=2        # size of the stride step  
)
```

Quiz:

What do the smaller red numbers represent in this image you've seen before?

1. The weights of the particular kernel we are applying (i.e. what feature we are detecting)
2. The intensity of the pixel for that area of the original image
3. The channel depth of the image

| | | |
|---------|---------|---------|
| 1 x1 | 0 x0 | 0 x1 |
| 1 x0 | 1 x1 | 0 x0 |
| 1 x1 | 1 x0 | 1 x1 |

Quiz:

What do the smaller red numbers represent in this image you've seen before?

1. The weights of the particular kernel we are applying (i.e. what feature we are detecting)
2. The intensity of the pixel for that area of the original image
3. The channel depth of the image

| | | |
|---------|---------|---------|
| 1 x1 | 0 x0 | 0 x1 |
| 1 x0 | 1 x1 | 0 x0 |
| 1 x1 | 1 x0 | 1 x1 |

Quiz:

Given the above, what the value of the output that the kernel will generate for this part of the image?

- 0
- 1
- 2
- 3
- 4
- 5

| | | |
|---------|---------|---------|
| 1 x1 | 0 x0 | 0 x1 |
| 1 x0 | 1 x1 | 0 x0 |
| 1 x1 | 1 x0 | 1 x1 |

Quiz:

Given the above, what the value of the output that the kernel will generate for this part of the image?

- 0
- 1
- 2
- 3
- 4
- 5

| | | |
|-----------------|-----------------|-----------------|
| 1 _{x1} | 0 _{x0} | 0 _{x1} |
| 1 _{x0} | 1 _{x1} | 0 _{x0} |
| 1 _{x1} | 1 _{x0} | 1 _{x1} |


File Name: T-ICML-O_2_I5_implementing_cnns

Format: Presenter in Studio

Presenter: Carl Osipov


Successive convolution layers apply filters to increasing scales

28x28x1




convolutional layer, 4 channels deep
W1[5, 5, 1, 4] stride 1

28x28x4




convolutional layer, 8 channels deep
W2[3, 3, 4, 8] stride 2

14x14x8



convolutional layer, 12 channels deep
W3[3, 3, 8, 12] stride 2

7x7x12

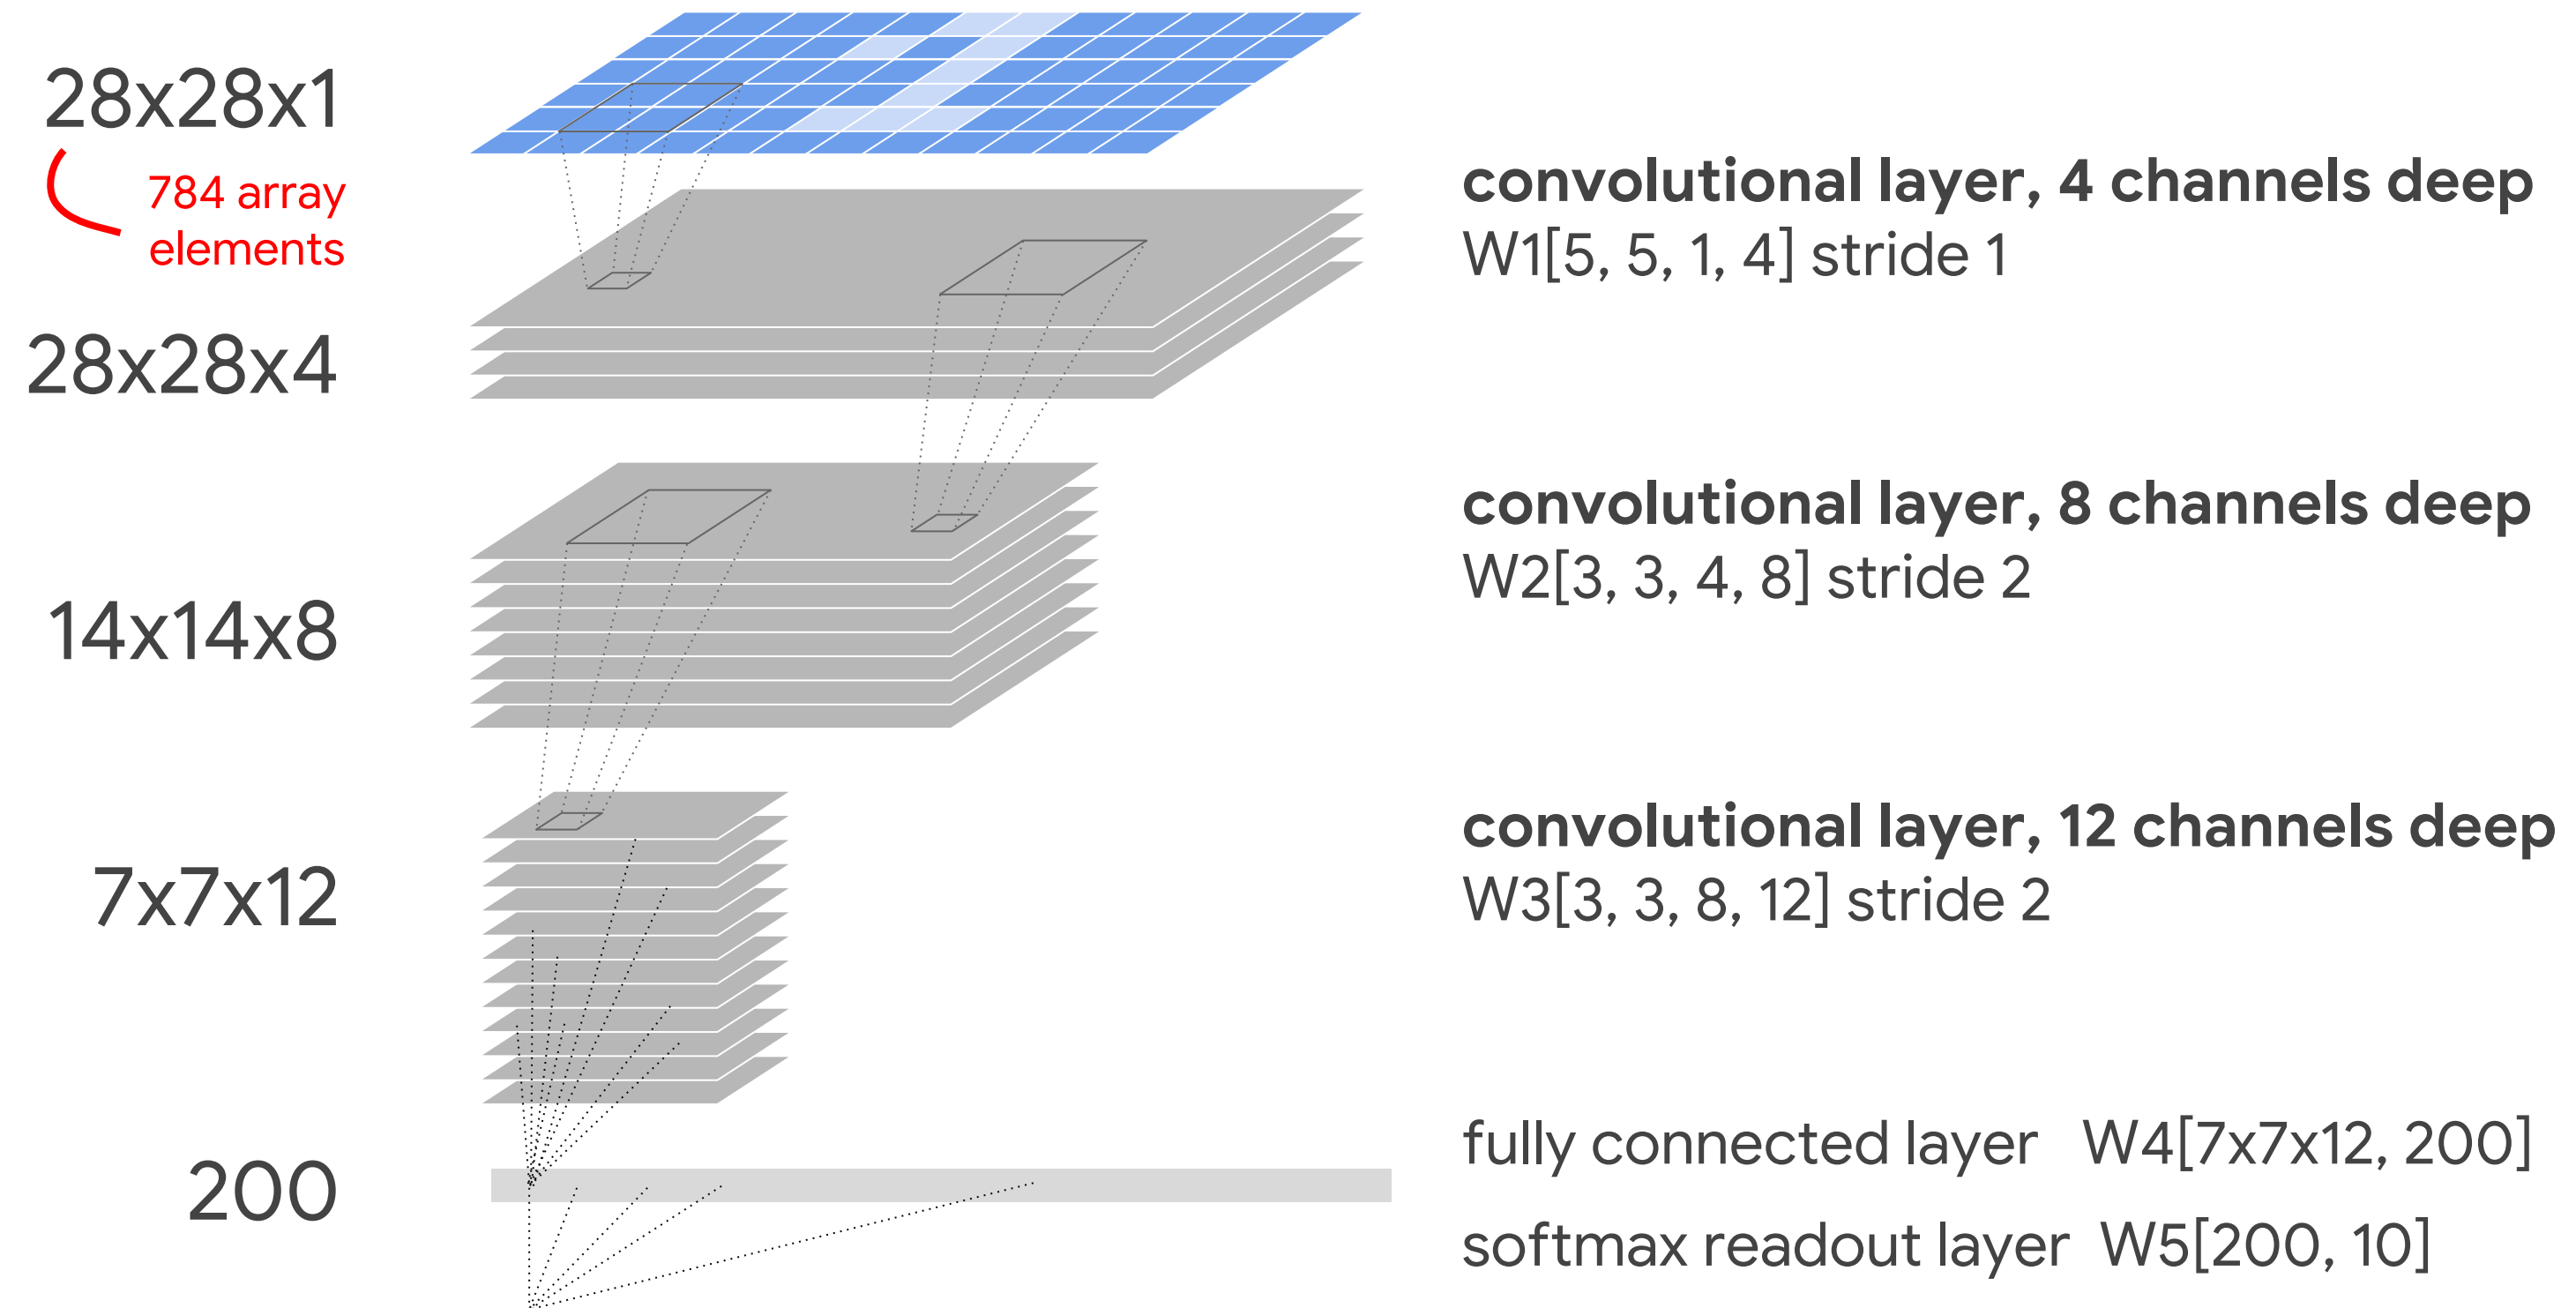


200




fully connected layer W4[7x7x12, 200]
softmax readout layer W5[200, 10]

Successive convolution layers apply filters to increasing scales



CNN model architecture for MNIST image



Setting up a model with Convolutional layers

```
def cnn_model(img, mode):
    X = tf.reshape(img, [-1, HEIGHT, WIDTH, 1]) # as a 2D image with one grayscale
    channel


    c1 = tf.layers.conv2d(X, filters=4, kernel_size=5, strides=1,
                           padding='same', activation=tf.nn.relu) # ?x28x28x4

    c2 = tf.layers.conv2d(c1, filters=8, kernel_size=3, strides=2,
                           padding='same', activation=tf.nn.relu)# ?x14x14x8

    c3 = tf.layers.conv2d(c2, filters=12, kernel_size=3, strides=2,
                           padding='same', activation=tf.nn.relu)# ?x7x7x12

    c3flat = tf.reshape(c3, [-1, 7*7*12]) # flattened
    h3 = tf.layers.dense(c3flat, 200, activation=tf.nn.relu)
    ylogits = tf.layers.dense(h3, NCLASSES, activation=None)
    return ylogits, NCLASSES
```

Using only a DNN for image classification
could have billions of weights



(Note: Many neurons not shown)



Kernels detect patterns and output filters

Kernel 1

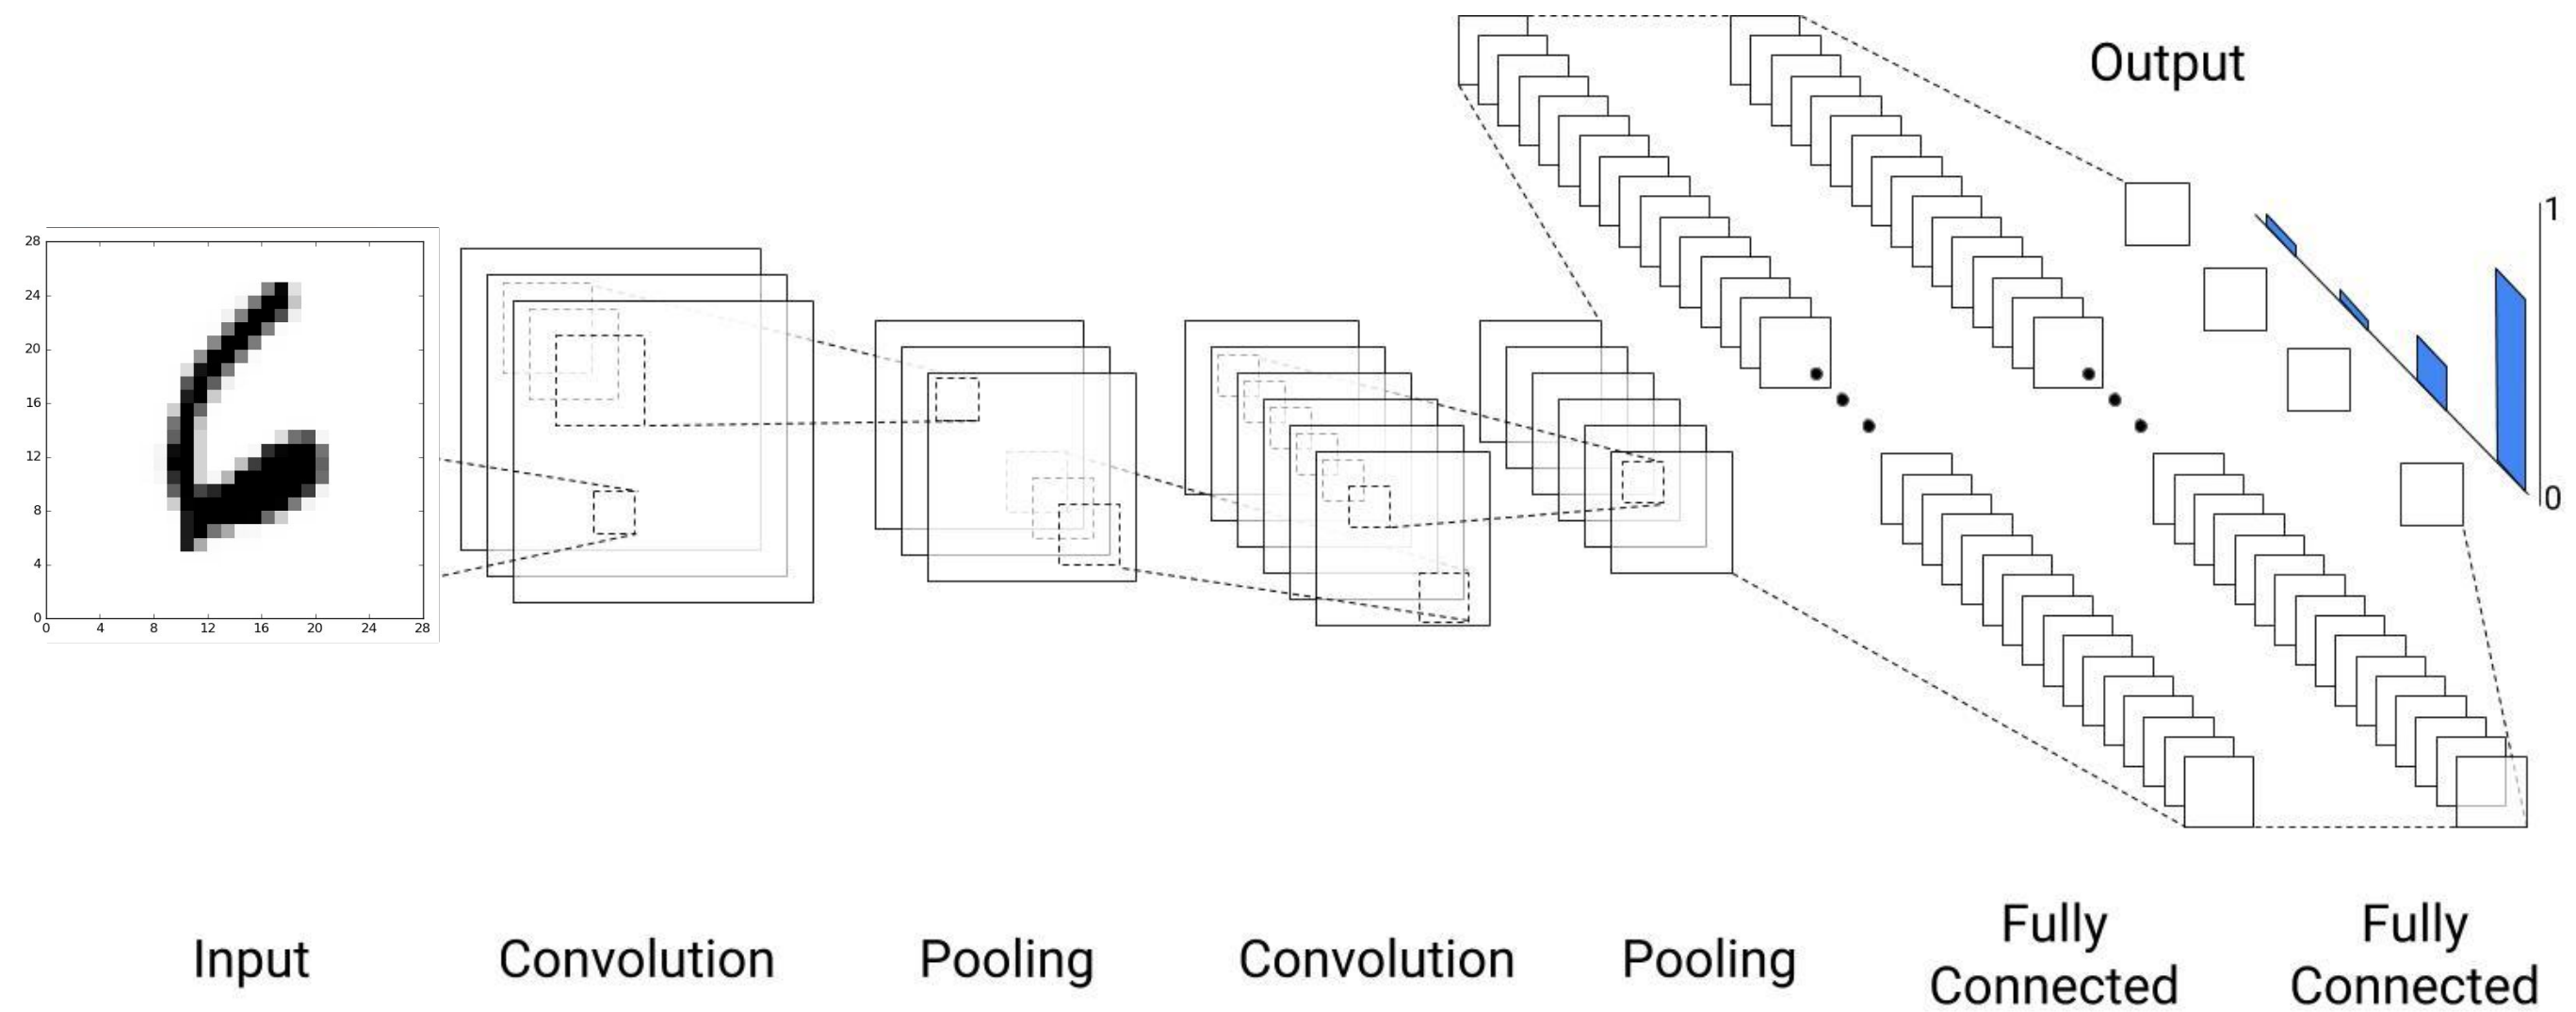
| | | |
|----|----|----|
| 1 | 2 | 1 |
| | | |
| -1 | -2 | -1 |

Kernel 2

| | | |
|---|--|----|
| 1 | | -1 |
| 2 | | -2 |
| 1 | | -1 |



Convolutional neural networks for image classification



Quiz:

Choose the advantage(s) of processing a small patch of the image at a time using a convolution kernel instead of using a dense layer for the entire image

1. Fewer weights
2. Faster processing
3. Weights are shared across the input
4. All of the above

Quiz:

Choose the advantage(s) of processing a small patch of the image at a time using a convolution kernel instead of using a dense layer for the entire image

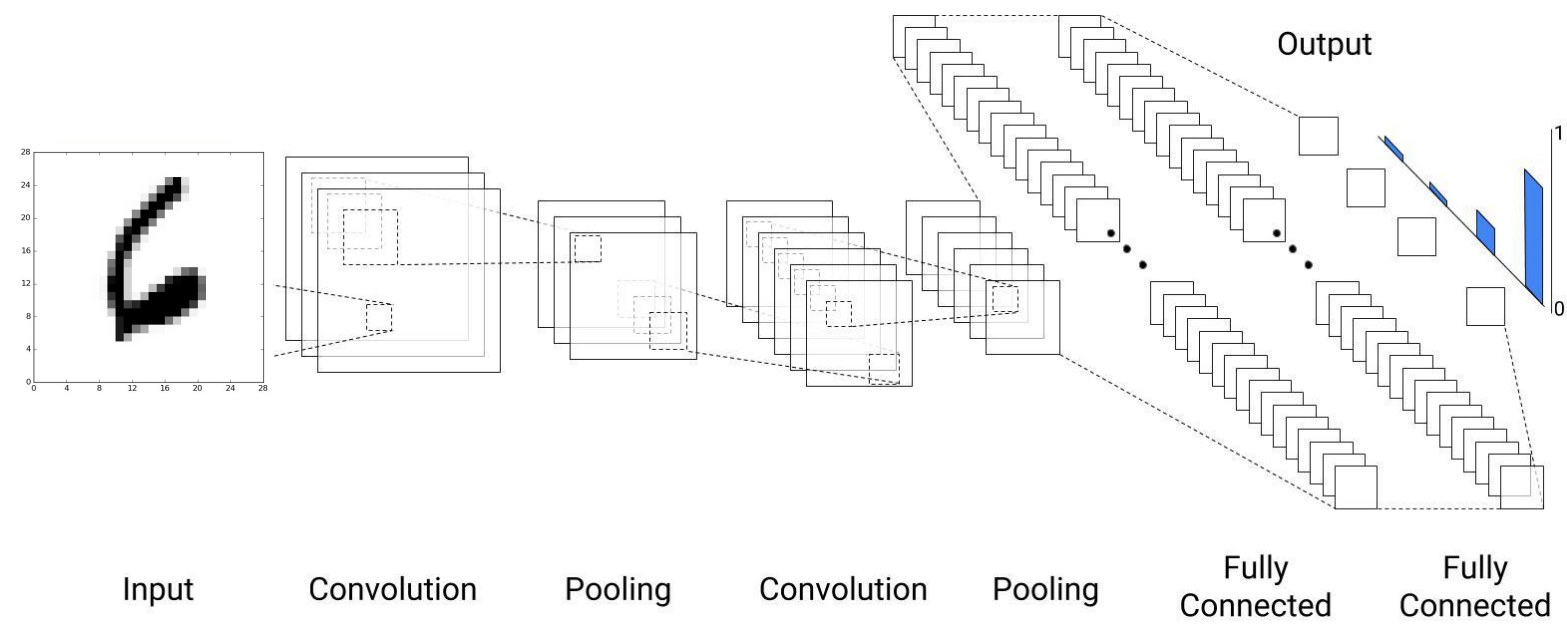
1. Fewer weights
2. Faster processing
3. Weights are shared across the input
- 4. All of the above**

File Name: T-ICML-O_2_I6_lab_intro_cnn

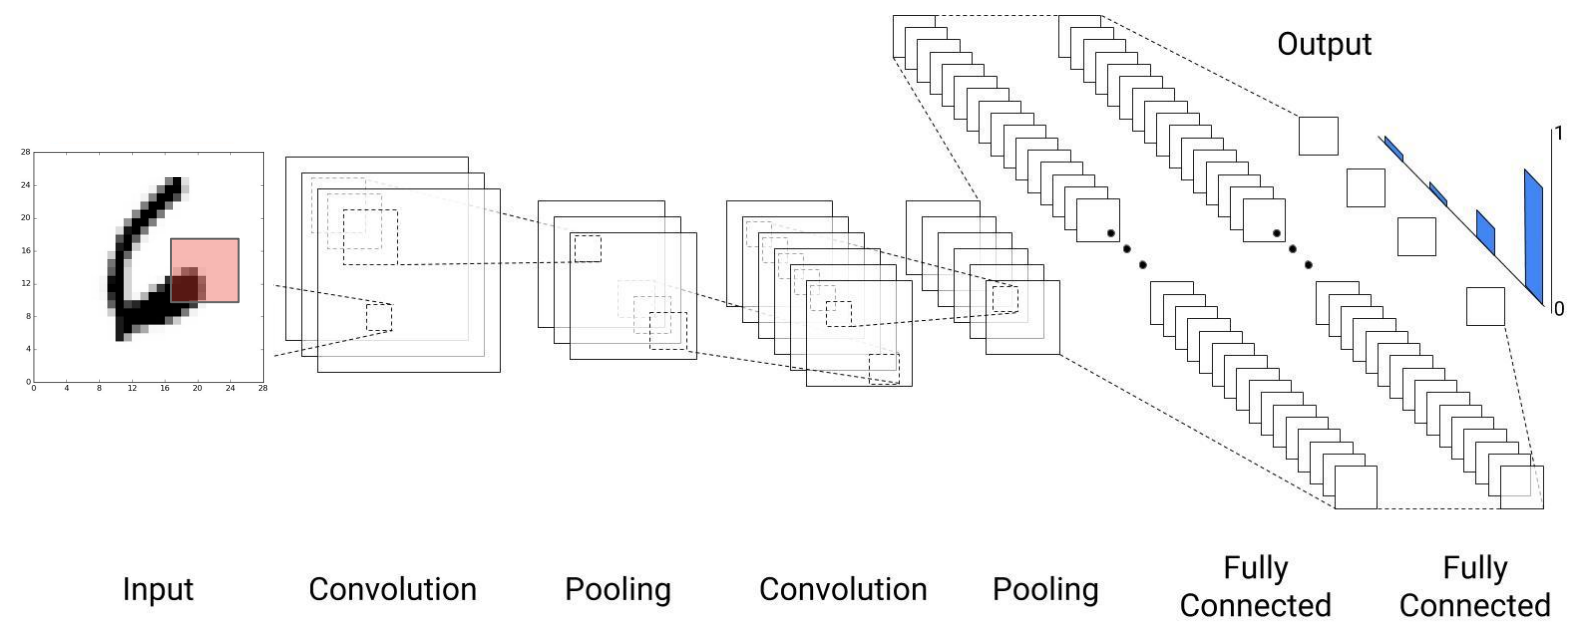
Format: Presenter in Studio

Presenter: Carl Osipov

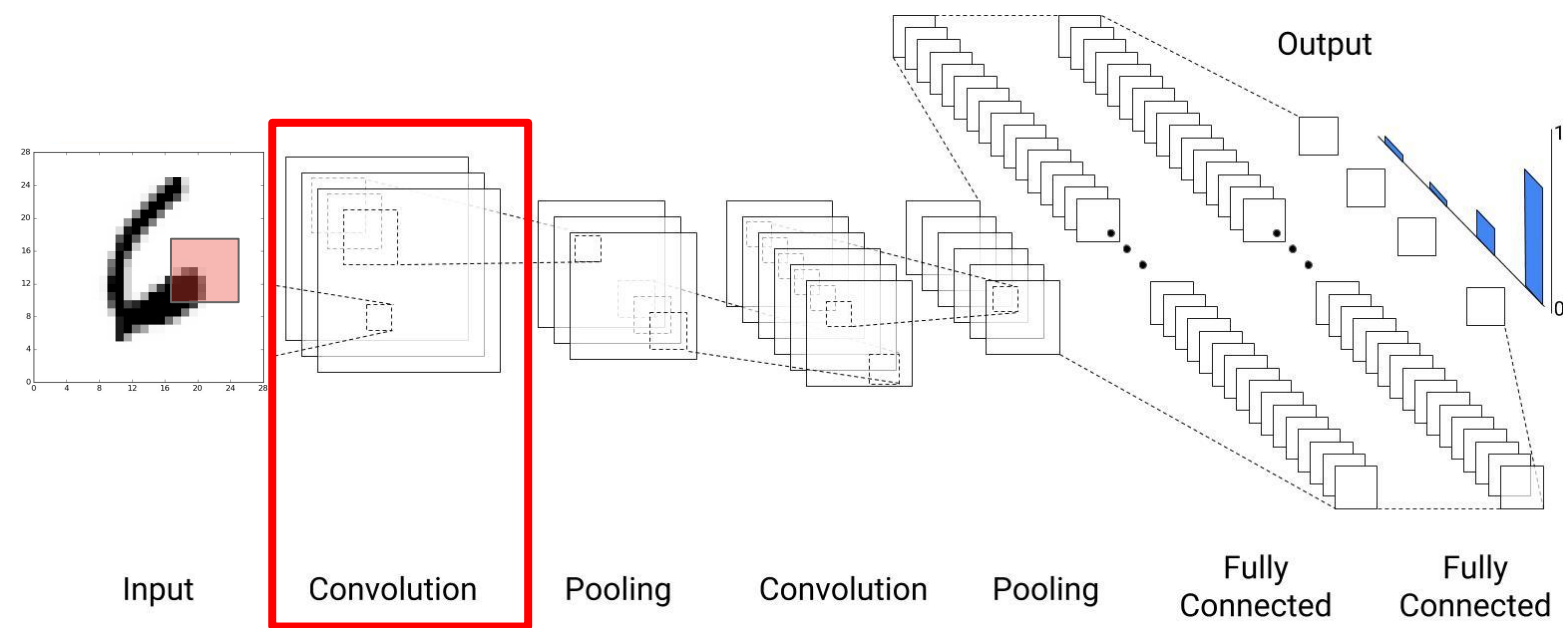
CNN Architecture Review



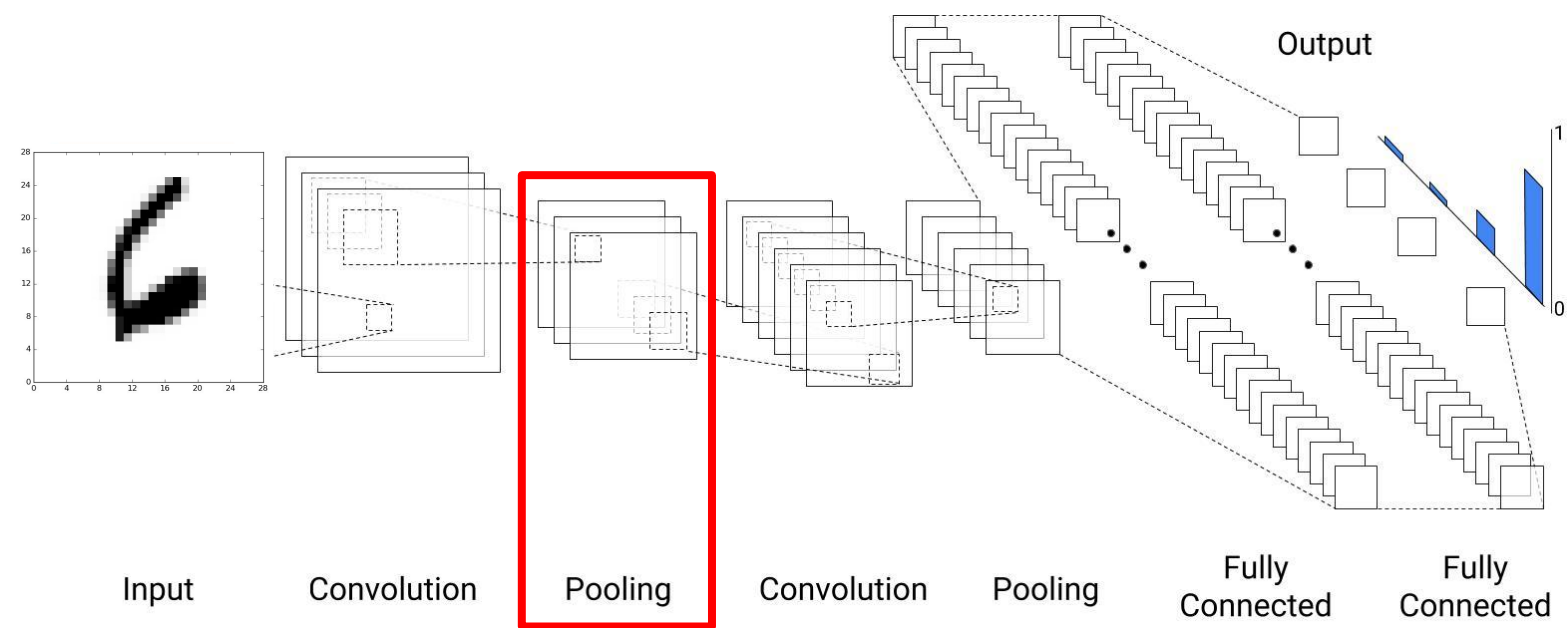
CNN Architecture Review



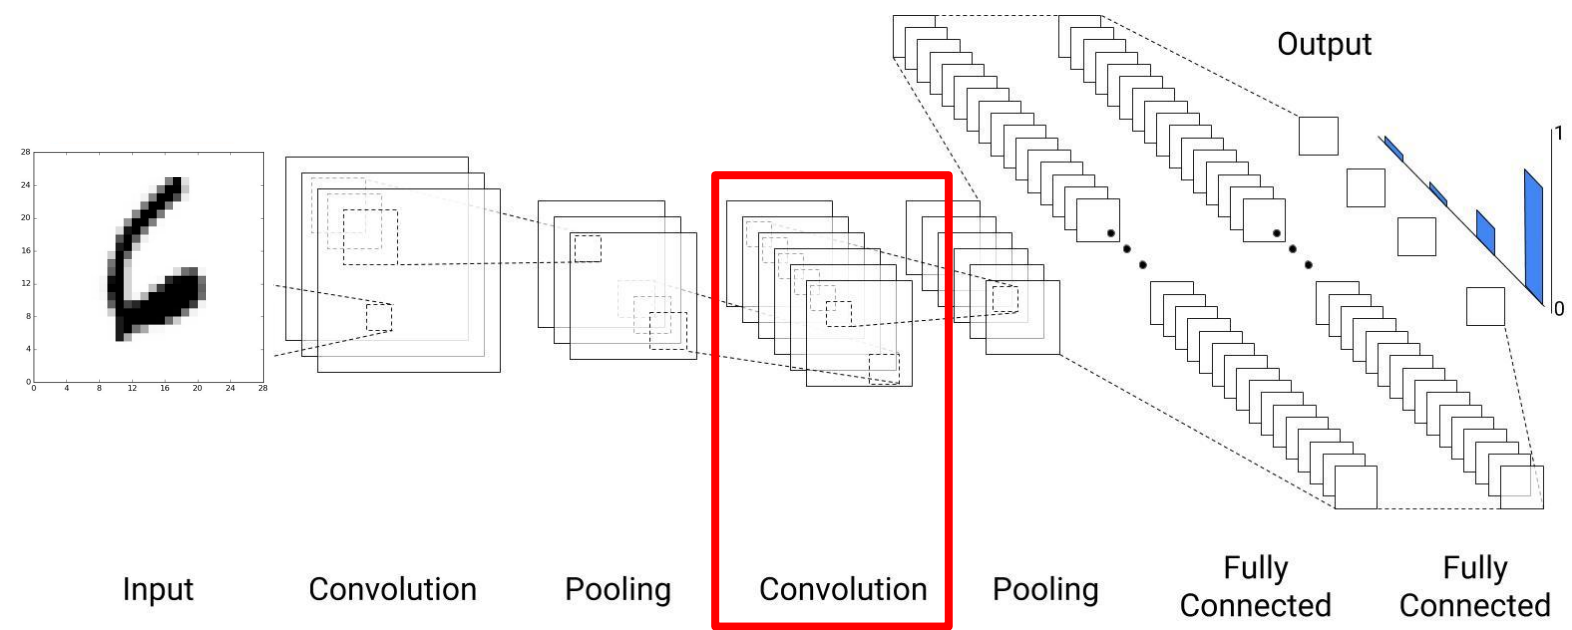
CNN Architecture Review




CNN Architecture Review



CNN Architecture Review



CNN Architecture Review



Fully-Connected
DNN for
classification

Lab

MNIST Image Classification
with TensorFlow on CMLE
(CNN)

Carl Osipov

Lab Steps

- Import the training dataset of MNIST handwritten images
- Reshape and preprocess the image data
- Setup your CNN with 10 classes
- Create convolutional and pooling layers + softmax function
- Define and create your EstimatorSpec in tensorflow to create your custom estimator
- Define and run your train_and_evaluate

File Name: T-ICML-O_2_I7_lab_solution_cnn

Format: Presenter in Studio

Presenter: Carl Osipov

